

capstone

April 23, 2025

0.1 Data import and Cleaning

```
[1]: import pandas as pd
from collections import Counter
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import Bio
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, accuracy_score, \
    confusion_matrix, roc_auc_score, roc_curve
from Bio.SeqUtils.ProtParam import ProteinAnalysis
import requests
from io import StringIO
from Bio import SeqIO

epitopes = pd.read_csv(r'/Users/tariq/Documents/capstone/data/
    ↳epitope_table_export_1740279588.csv')
assays = pd.read_csv(r'/Users/tariq/Documents/capstone/data/
    ↳tcell_table_export_1740279970.csv')

def fetch_full_sequence(url):
    if pd.notna(url): # Check if the URL is not NaN
        url = f'{url}.fasta'
        try:
            response = requests.get(url)
            if response.status_code == 200:
                fasta_io = StringIO(response.text)
                records = list(SeqIO.parse(fasta_io, "fasta"))
                if records: # Check if there are any records
                    return str(records[0].seq)
                else:
                    print("No records found in the FASTA file.")
            except requests.exceptions.RequestException as e:
                print(f"Request failed: {e}")
        return None
```

```
#epitopes['Full Sequence'] = epitopes['Epitope - Molecule Parent IRI'].
↳apply(fetch_full_sequence)
epitopes = pd.read_csv(r'/Users/tariq/Documents/capstone/data/epitope_full_seq.
↳csv')
```

```
/var/folders/5j/4p7c5_1x2fg18bk0nf74_hg40000gn/T/ipykernel_31212/1845155022.py:1
5: DtypeWarning: Columns (13,14,45,46,47,48,49,54,55,56,57,60,65,66,67,68,69,70,
71,72,73,74,75,76,77,78,79,82,83,84,85,86,87,88,89,90,91,92,93,94,95,96,97,98,99
,100,101,102,105,106,107,108,109,110,111,112,113,115,120,123,128,132,134,135,141
,142,143,144,145,149,152) have mixed types. Specify dtype option on import or
set low_memory=False.
```

```
assays = pd.read_csv(r'/Users/tariq/Documents/capstone/data/tcell_table_export
_1740279970.csv')
```

```
[2]: # make all the column names snake case
epitopes.columns = epitopes.columns.str.lower()
assays.columns = assays.columns.str.lower()

# remove spaces from column names
epitopes.columns = epitopes.columns.str.replace(' ', '')
epitopes.columns = epitopes.columns.str.replace('-', '_')
epitopes.columns = epitopes.columns.str.replace(' ', '_')

assays.columns = assays.columns.str.replace(' ', '')
assays.columns = assays.columns.str.replace('-', '_')
assays.columns = assays.columns.str.replace(' ', '_')

epitopes = epitopes.filter(['epitope_objecttype', 'epitope_name',
↳'fullsequence'])
assays = assays.filter(['epitope_name', 'epitope_moleculeparent', 'host_name',
↳'host_mhcrestriction', 'assay_method', 'assay_response',
↳'assay_qualitative_measurement', 'mhcrestriction_name',
↳'mhcrestriction_class', 'assayantigen_name'])

# map mhc name and class from the assays dataframe to a new column in the
↳epitopes dataframe based on epitope_name
mhc = assays.filter(['epitope_name', 'mhcrestriction_name',
↳'mhcrestriction_class'])
mhc = mhc.drop_duplicates(subset=['epitope_name'])
epitopes = epitopes.merge(mhc, on='epitope_name', how='left')
```

```
[3]: epitopes.head()
```

```
[3]:  epitope_objecttype epitope_name \
0      Linear peptide    AAGIGILTV
1      Linear peptide    AAGIGILTVI
2      Linear peptide    ACDPHSGHFV
```

```

3      Linear peptide    ADLVGFLLLK
4      Linear peptide    ADVEFCLSL

                                fullsequence mhcrestriction_name \
0  MPREDAHFIIYGYPKKGHGHSYTTAEAAAGIGILTIVILGVLLLLIGCW...      HLA-A2
1  MPREDAHFIIYGYPKKGHGHSYTTAEAAAGIGILTIVILGVLLLLIGCW...      HLA-A*02:01
2                                     NaN                        HLA-A2
3  MSLEQRSLHCKPEEALAQEALGLVCVQAATSSSSPLVLGTLEEV...      HLA-A*11:01
4  MLLAVLYCLLWSFQTSAGHFPRACVSSKNLMEKECCPPWSGDRSPC...      HLA-B*44:03

mhcrestriction_class
0                      I
1                      I
2                      I
3                      I
4                      I

```

```
[4]: epitopes.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 28681 entries, 0 to 28680
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   epitope_objecttype     28681 non-null object
1   epitope_name           28681 non-null object
2   fullsequence           7164 non-null  object
3   mhcrestriction_name    17613 non-null object
4   mhcrestriction_class   17613 non-null object
dtypes: object(5)
memory usage: 1.1+ MB

```

0.2 Feature Engineering

```
[5]: epitopes['epitope_length'] = epitopes['epitope_name'].str.len()
```

```

[6]: # Function to count amino acids in a peptide
def count_amino_acids(peptide):
    try:
        # Create a ProteinAnalysis object for the peptide
        analyzer = ProteinAnalysis(peptide)
        # Get amino acid counts and normalize to frequencies
        aa_count = analyzer.count_amino_acids()
        total_aa = sum(aa_count.values())
        aa_freq = {aa: count for aa, count in aa_count.items()}
        # Add the peptide itself to the results
        aa_freq['peptide'] = peptide
    
```

```

        return aa_freq
    except Exception as e:
        # Handle invalid peptides (e.g., with non-standard amino acids)
        result = {aa: 0 for aa in 'ACDEFGHIKLMNPQRSTVWY'}
        result['peptide'] = peptide
        return result

# Create analyzer function that will be used in the next cell
def analyzer(peptide):
    return count_amino_acids(peptide)

# Use both epitope name and peptide sequence in the DataFrame
epitope_composition_df = epitopes.apply(lambda row:
    ↪count_amino_acids(row['epitope_name']), axis=1).apply(pd.Series)

```

```
[7]: epitope_composition_df.head()
```

```

[7]:   A  C  D  E  F  G  H  I  K  L  ...  N  P  Q  R  S  T  V  W  Y    peptide
0  2  0  0  0  0  2  0  2  0  1  ...  0  0  0  0  0  1  1  0  0    AAGIGILTV
1  2  0  0  0  0  2  0  3  0  1  ...  0  0  0  0  0  1  1  0  0    AAGIGILTVI
2  1  1  1  0  1  1  2  0  0  0  ...  0  1  0  0  1  0  1  0  0    ACDPHSGHFV
3  1  0  1  0  1  1  0  0  1  4  ...  0  0  0  0  0  0  1  0  0    ADLVGFLLLK
4  1  1  1  1  1  0  0  0  0  2  ...  0  0  0  0  1  0  1  0  0    ADVEFCLSL

```

[5 rows x 21 columns]

```

[8]: # Example DataFrame with a 'peptide' column
df = pd.DataFrame({
    'peptide': ['ACDEFGHIK', 'LMNPQRSTV', 'WYFP']
})

# Kyte-Doolittle hydrophobicity scale
kyte_doolittle = {
    'I': 4.5, 'V': 4.2, 'L': 3.8, 'F': 2.8, 'C': 2.5,
    'M': 1.9, 'A': 1.8, 'G': -0.4, 'T': -0.7, 'S': -0.8,
    'W': -0.9, 'Y': -1.3, 'P': -1.6, 'H': -3.2, 'E': -3.5,
    'Q': -3.5, 'D': -3.5, 'N': -3.5, 'K': -3.9, 'R': -4.5
}

def compute_avg_hydrophobicity(peptide):
    # Get hydrophobicity scores for each amino acid; default to 0 if missing
    scores = [kyte_doolittle.get(aa, 0) for aa in peptide]
    return sum(scores) / len(scores) if scores else 0

# Apply the function to the 'peptide' column to create a new column 'avg_hydro'
epitopes['epitope_avg_hydro'] = epitopes['epitope_name'].
    ↪apply(compute_avg_hydrophobicity)

```

```
[9]: # Import the molecular_weight function from Bio.SeqUtils

def calculate_molecular_weight(peptide):
    """Calculate the molecular weight of a peptide sequence using Biopython."""
    try:
        # ProteinAnalysis only works with standard amino acids
        protein = ProteinAnalysis(peptide)
        return protein.molecular_weight()
    except Exception as e:
        # Handle peptides with non-standard amino acids
        return None

# Apply the function to calculate molecular weight for each epitope
epitopes['molecular_weight'] = epitopes['epitope_name'].
    ↪apply(calculate_molecular_weight)
```

```
[10]: def calculate_aromaticity(peptide):
    """Calculate the aromaticity of a peptide sequence using Biopython."""
    try:
        # ProteinAnalysis only works with standard amino acids
        protein = ProteinAnalysis(peptide)
        return protein.aromaticity()
    except Exception as e:
        # Handle peptides with non-standard amino acids
        return None

# Apply the function to calculate molecular weight for each epitope
epitopes['aromaticity'] = epitopes['epitope_name'].apply(calculate_aromaticity)
```

```
[11]: def calculate_isoelectric_point(peptide):
    """Calculate the isoelectric point of a peptide sequence using Biopython."""
    try:
        # ProteinAnalysis only works with standard amino acids
        protein = ProteinAnalysis(peptide)
        return protein.isoelectric_point()
    except Exception as e:
        # Handle peptides with non-standard amino acids
        return None

# Apply the function to calculate molecular weight for each epitope
epitopes['isoelectric_point'] = epitopes['epitope_name'].
    ↪apply(calculate_isoelectric_point)
```

```
[12]: def calculate_instability(peptide):
    """Calculate the instability of a peptide sequence using Biopython."""
    try:
```

```

        # ProteinAnalysis only works with standard amino acids
        protein = ProteinAnalysis(peptide)
        return protein.instability_index()
    except Exception as e:
        # Handle peptides with non-standard amino acids
        return None

# Apply the function to calculate molecular weight for each epitope
epitopes['instability'] = epitopes['epitope_name'].apply(calculate_instability)

```

```

[13]: def calculate_charge_at_pH7(peptide):
        """Calculate the charge of a peptide sequence at pH 7 using Biopython."""
        try:
            # ProteinAnalysis only works with standard amino acids
            protein = ProteinAnalysis(peptide)
            return protein.charge_at_pH(7)
        except Exception as e:
            # Handle peptides with non-standard amino acids
            return None

# Apply the function to calculate molecular weight for each epitope
epitopes['charge_at_pH7'] = epitopes['epitope_name'].
    ↪ apply(calculate_charge_at_pH7)

```

```

[14]: epitopes.head()

```

```

[14]: epitope_objecttype epitope_name \
0      Linear peptide    AAGIGILTV
1      Linear peptide    AAGIGILTVI
2      Linear peptide    ACDPHSGHFV
3      Linear peptide    ADLVGFLLLK
4      Linear peptide    ADVEFCLSL

                                fullsequence mhcrestriction_name \
0  MPREDAHFIYGYPKKGHGHSYTTAEAAAGIGILTVILGVLLIGCW...          HLA-A2
1  MPREDAHFIYGYPKKGHGHSYTTAEAAAGIGILTVILGVLLIGCW...        HLA-A*02:01
2                                     NaN                      HLA-A2
3  MSLEQRSLHCKPEEALEAQQEALGLVCVQAATSSSSPLVLGTLEEV...        HLA-A*11:01
4  MLLAVLYCLLWSFQTSAGHFPRACVSSKNLMEKECCPPWSGDRSPC...        HLA-B*44:03

mhcrestriction_class  epitope_length  epitope_avg_hydro  molecular_weight \
0                    I                9          2.122222          813.9814
1                    I               10          2.360000          927.1390
2                    I               10         -0.140000         1069.1507
3                    I               10          1.620000         1088.3394
4                    I                9          1.233333          996.1348

```

	aromaticity	isoelectric_point	instability	charge_at_pH7
0	0.000000	5.570017	11.422222	-0.204125
1	0.000000	5.570017	11.280000	-0.204125
2	0.100000	5.972266	61.830000	-1.038557
3	0.100000	5.880358	-16.470000	-0.204004
4	0.111111	4.050028	20.855556	-2.210095

0.3 Generation of Negative Samples

```
[15]: def generate_negatives(row):
    epitope = row["epitope_name"]
    full_seq = row["fullsequence"]
    mhc = row["mhcrestriction_name"]

    # Handle missing or empty sequences
    if pd.isnull(full_seq) or full_seq == "":
        return []

    epitope = str(epitope)
    full_seq = str(full_seq)
    ep_len = len(epitope)

    negatives = []
    for i in range(len(full_seq) - ep_len + 1):
        window = full_seq[i:i+ep_len]
        if window != epitope:
            negatives.append({"peptide": window, "mhc": mhc})
    return negatives

# Apply the function to each row
'''
negatives = pd.DataFrame()
negatives['negatives'] = epitopes.apply(generate_negatives, axis=1)
negatives = negatives[["negatives"]].explode("negatives").reset_index(drop=True)
negatives.dropna(subset=["negatives"], inplace=True)

# Remove duplicate peptide-mhc combinations
print(f"Shape before removing duplicates: {negatives.shape}")
negatives = negatives.drop_duplicates(subset=['negatives'])
print(f"Shape after removing duplicates: {negatives.shape}")

# Check for any remaining NaN values
print(f"Number of NaN values in negatives: {negatives['negatives'].isna().
      ↪sum()}")

# Extract peptide and mhc into separate columns
```

```

negatives['peptide'] = negatives['negatives'].apply(lambda x: x['peptide'])
negatives['mhc'] = negatives['negatives'].apply(lambda x: x['mhc'])

# Calculate features on the peptide column
negatives['peptide_length'] = negatives['peptide'].apply(len)
negatives['peptide_avg_hydro'] = negatives['peptide'].
    ↪ apply(compute_avg_hydrophobicity)
negatives['molecular_weight'] = negatives['peptide'].
    ↪ apply(calculate_molecular_weight)
negatives['aromaticity'] = negatives['peptide'].apply(calculate_aromaticity)
negatives['isoelectric_point'] = negatives['peptide'].
    ↪ apply(calculate_isoelectric_point)
negatives['instability'] = negatives['peptide'].apply(calculate_instability)
negatives['charge_at_pH7'] = negatives['peptide'].apply(calculate_charge_at_pH7)

# Drop the original dictionary column if no longer needed
negatives.drop('negatives', axis=1, inplace=True)
'''

```

```

[15]: '\nnegatives = pd.DataFrame()\nnegatives[\negatives\'] =
epitopes.apply(generate_negatives, axis=1)\nnegatives = negatives[["negatives"]]
.explode("negatives").reset_index(drop=True)\nnegatives.dropna(subset=["negative
s"], inplace=True)\n\n# Remove duplicate peptide-mhc
combinations\nprint(f"Shape before removing duplicates:
{negatives.shape}")\nnegatives =
negatives.drop_duplicates(subset=[\negatives\'])\nprint(f"Shape after removing
duplicates: {negatives.shape}")\n\n# Check for any remaining NaN
values\nprint(f"Number of NaN values in negatives:
{negatives[\negatives\'].isna().sum()}")\n\n# Extract peptide and mhc into
separate columns\nnegatives[\ peptide\'] = negatives[\negatives\'].apply(lambda
x: x[\ peptide\'])\nnegatives[\ mhc\'] = negatives[\negatives\'].apply(lambda
x: x[\ mhc\'])\n\n# Calculate features on the peptide
column\nnegatives[\ peptide_length\'] =
negatives[\ peptide\'].apply(len)\nnegatives[\ peptide_avg_hydro\'] = negatives[
\ peptide\'].apply(compute_avg_hydrophobicity)\nnegatives[\ molecular_weight\']
= negatives[\ peptide\'].apply(calculate_molecular_weight)\nnegatives[\ aromatic
ity\'] = negatives[\ peptide\'].apply(calculate_aromaticity)\nnegatives[\ isoele
ctric_point\'] = negatives[\ peptide\'].apply(calculate_isoelectric_point)\nnega
tives[\ instability\'] = negatives[\ peptide\'].apply(calculate_instability)\nne
gatives[\ charge_at_pH7\'] =
negatives[\ peptide\'].apply(calculate_charge_at_pH7)\n\n# Drop the original
dictionary column if no longer needed\nnegatives.drop(\negatives\ ', axis=1,
inplace=True)\n'

```

```

[16]: negatives = pd.read_csv("data/negatives_MHC.csv")

```

/var/folders/5j/4p7c5_1x2fg18bk0nf74_hg40000gn/T/ipykernel_31212/1811011591.py:1


```
: DtypeWarning: Columns (1) have mixed types. Specify dtype option on import or
set low_memory=False.
```

```
negatives = pd.read_csv("data/negatives_MHC.csv")
```

```
[17]: nine_mers = epitopes[epitopes['epitope_length'] == 9]
nine_mers.to_csv("data/ninemer_epitopes.csv", index=False)
```

```
[18]: ninemer_negatives = negatives[negatives['peptide_length'] == 9]
ninemer_negatives_trimmed = ninemer_negatives[:50000]
ninemer_negatives_trimmed.to_csv("data/ninemer_negatives_trimmed.csv",
↪index=False)
```

0.4 EDA

0.4.1 Data Summary

```
[19]: epitopes.head()
```

```
[19]: epitope_objecttype epitope_name \
0      Linear peptide    AAGIGILTV
1      Linear peptide    AAGIGILTVI
2      Linear peptide    ACDPHSGHFV
3      Linear peptide    ADLVGFLLLK
4      Linear peptide    ADVEFCLSL
```

```
                                fullsequence mhcrestriction_name \
0  MPREDAHFIYGYPKKGHGHSYTTAEAAAGIGILTVILGVLLIGCW...      HLA-A2
1  MPREDAHFIYGYPKKGHGHSYTTAEAAAGIGILTVILGVLLIGCW...  HLA-A*02:01
2                                     NaN                HLA-A2
3  MSLEQRSLSHCKPEEALEAQEALGLVCVQAATSSSSPLVLGTLEEV...  HLA-A*11:01
4  MLLAVLYCLLWSFQTSAGHFPRACVSSKNLMEKECCPPWSGDRSPC...  HLA-B*44:03
```

```
mhcrestriction_class epitope_length epitope_avg_hydro molecular_weight \
0                    I              9          2.122222          813.9814
1                    I             10          2.360000          927.1390
2                    I             10         -0.140000         1069.1507
3                    I             10          1.620000         1088.3394
4                    I              9          1.233333          996.1348
```

```
aromaticity isoelectric_point instability charge_at_pH7
0    0.000000      5.570017    11.422222    -0.204125
1    0.000000      5.570017    11.280000    -0.204125
2    0.100000      5.972266    61.830000    -1.038557
3    0.100000      5.880358   -16.470000    -0.204004
4    0.111111      4.050028    20.855556    -2.210095
```

```
[21]: epitopes.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 28681 entries, 0 to 28680
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   epitope_objecttype                    28681 non-null  object
1   epitope_name                          28681 non-null  object
2   fullsequence                         7164 non-null   object
3   mhcrestriction_name                  17613 non-null  object
4   mhcrestriction_class                 17613 non-null  object
5   epitope_length                       28681 non-null  int64
6   epitope_avg_hydro                    28681 non-null  float64
7   molecular_weight                     28623 non-null  float64
8   aromaticity                         28681 non-null  float64
9   isoelectric_point                   28681 non-null  float64
10  instability                           28623 non-null  float64
11  charge_at_pH7                       28681 non-null  float64
dtypes: float64(6), int64(1), object(5)
memory usage: 2.6+ MB

```

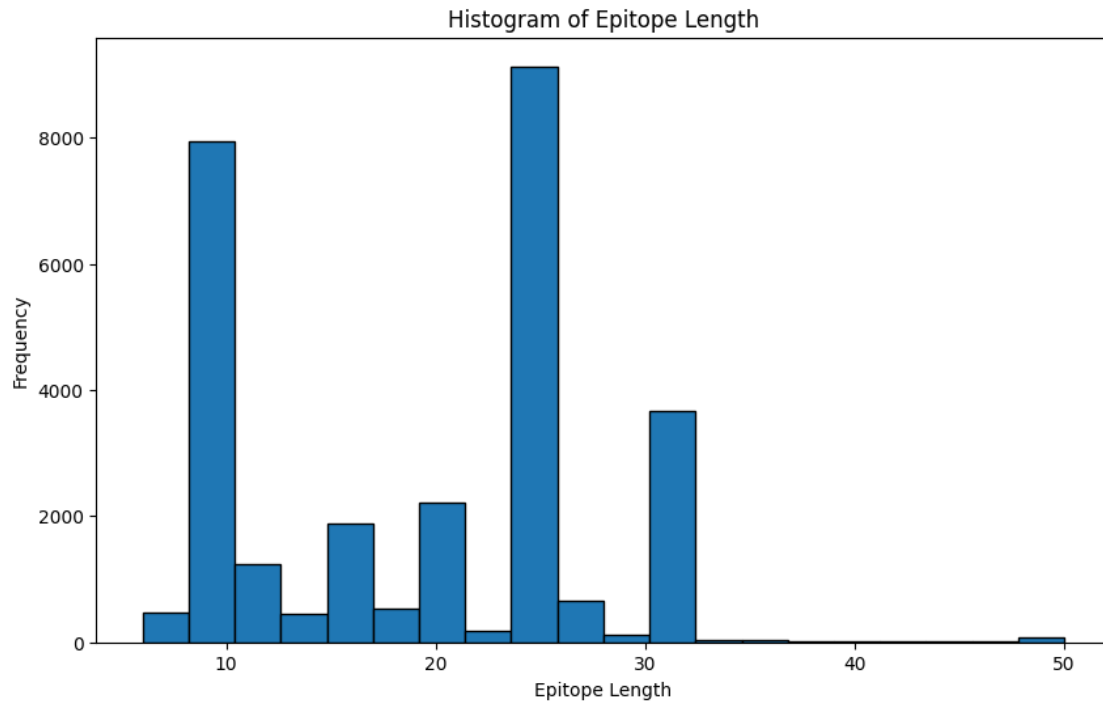
0.4.2 Properties of Epitopes

Length

```

[23]: # hist of epitope length
plt.figure(figsize=(10, 6))
plt.hist(epitopes['epitope_length'], bins=20, edgecolor='black')
plt.xlabel('Epitope Length')
plt.ylabel('Frequency')
plt.title('Histogram of Epitope Length')
plt.show()

```

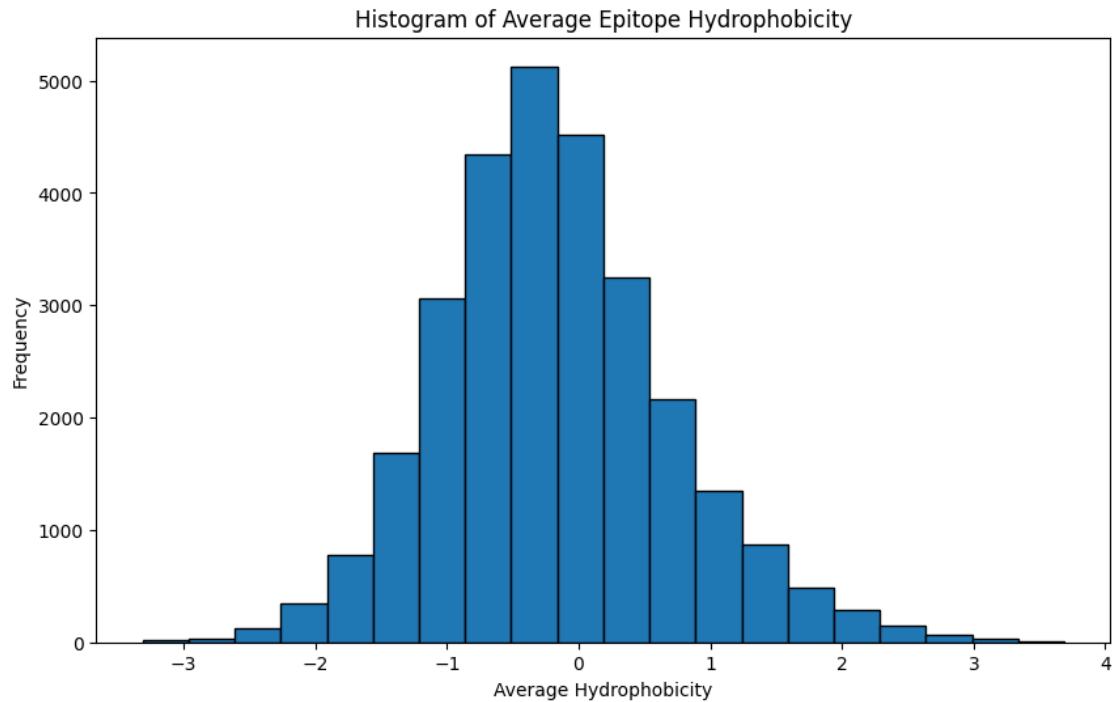


```
[24]: epitopes['epitope_length'].describe()
```

```
[24]: count    28681.000000
      mean      19.389422
      std       8.255925
      min       6.000000
      25%      10.000000
      50%      20.000000
      75%      25.000000
      max      50.000000
      Name: epitope_length, dtype: float64
```

Hydrophobicity

```
[25]: # histogram of average hydrophobicity
      plt.figure(figsize=(10, 6))
      plt.hist(epitopes['epitope_avg_hydro'], bins=20, edgecolor='black')
      plt.xlabel('Average Hydrophobicity')
      plt.ylabel('Frequency')
      plt.title('Histogram of Average Epitope Hydrophobicity')
      plt.show()
```



```
[26]: epitopes['epitope_avg_hydro'].describe()
```

```
[26]: count    28681.000000
      mean      -0.178410
      std       0.883064
      min      -3.312000
      25%      -0.762500
      50%      -0.240000
      75%       0.333333
      max       3.688889
      Name: epitope_avg_hydro, dtype: float64
```

Composition

```
[27]: # plot the composition of the epitopes, sort by the composition of the amino
      ↪ acids
      # Calculate mean composition and sort

      '''
      mean_composition = epitope_composition_df.mean().sort_values(ascending=False)

      # Plot the sorted composition
      plt.figure(figsize=(10, 6))
      plt.bar(mean_composition.index, mean_composition.values)
```

```
plt.xlabel('Amino Acid')
plt.ylabel('Composition')
plt.title('Composition of Epitopes')
plt.show()

'''
```

```
[27]: "\nmean_composition =
epitope_composition_df.mean().sort_values(ascending=False)\n\n# Plot the sorted
composition\nplt.figure(figsize=(10, 6))\nplt.bar(mean_composition.index,
mean_composition.values)\nplt.xlabel('Amino
Acid')\nplt.ylabel('Composition')\nplt.title('Composition of
Epitopes')\nplt.show()\n\n"
```

n-gram frequency analysis

```
[28]: def ngram_frequency(peptides, n=2):
    ngrams = []
    for peptide in peptides:
        if len(peptide) < n:
            continue
        for i in range(len(peptide) - n + 1):
            ngram = peptide[i:i+n]
            ngrams.append(ngram)
    return Counter(ngrams)

dipeptide_freq = ngram_frequency(epitopes['epitope_name'], n=2)

df_ngram = pd.DataFrame(dipeptide_freq.items(), columns=['ngram', 'count'])
df_ngram = df_ngram.sort_values('count', ascending=False)

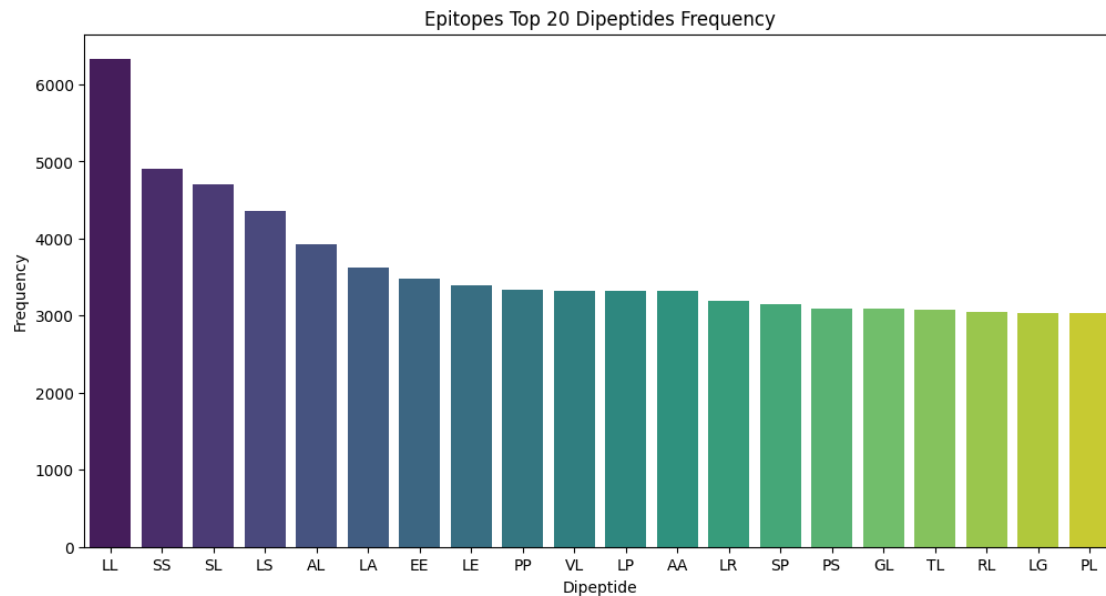
top_n = 20
top_ngram = df_ngram.head(top_n)

plt.figure(figsize=(12, 6))
sns.barplot(x='ngram', y='count', data=top_ngram, palette="viridis")
plt.title(f"Epitopes Top {top_n} Dipeptides Frequency")
plt.xlabel("Dipeptide")
plt.ylabel("Frequency")
plt.show()
```

```
/var/folders/5j/4p7c5_1x2fg18bk0nf74_hg40000gn/T/ipykernel_31212/733366050.py:20
: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

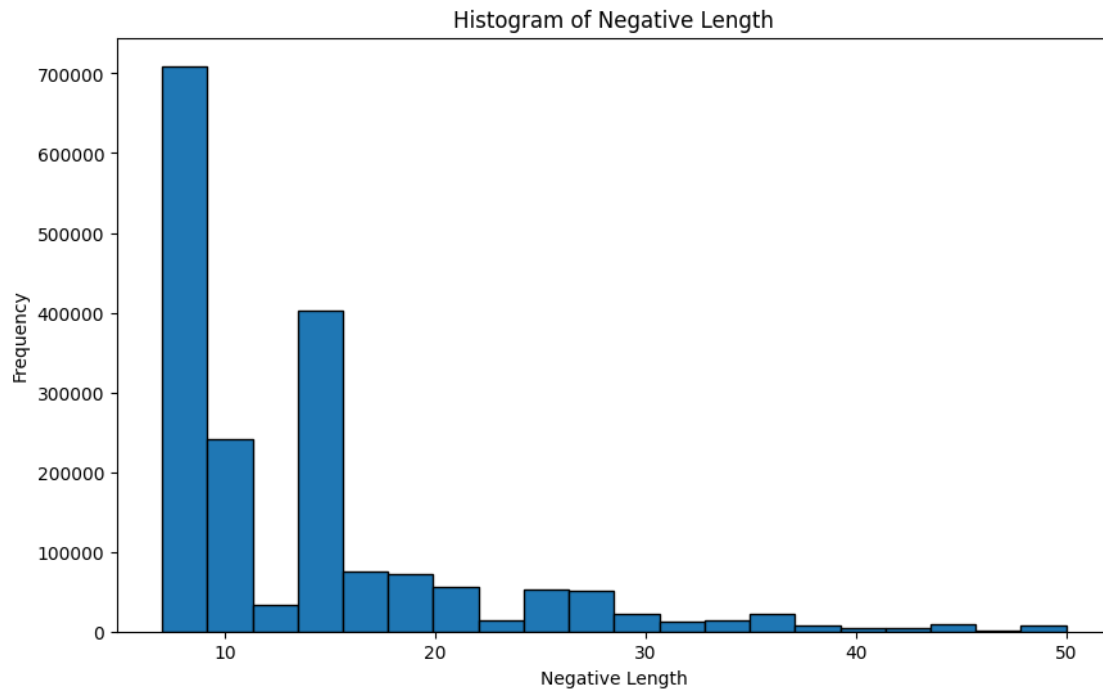
```
sns.barplot(x='ngram', y='count', data=top_ngram, palette="viridis")
```



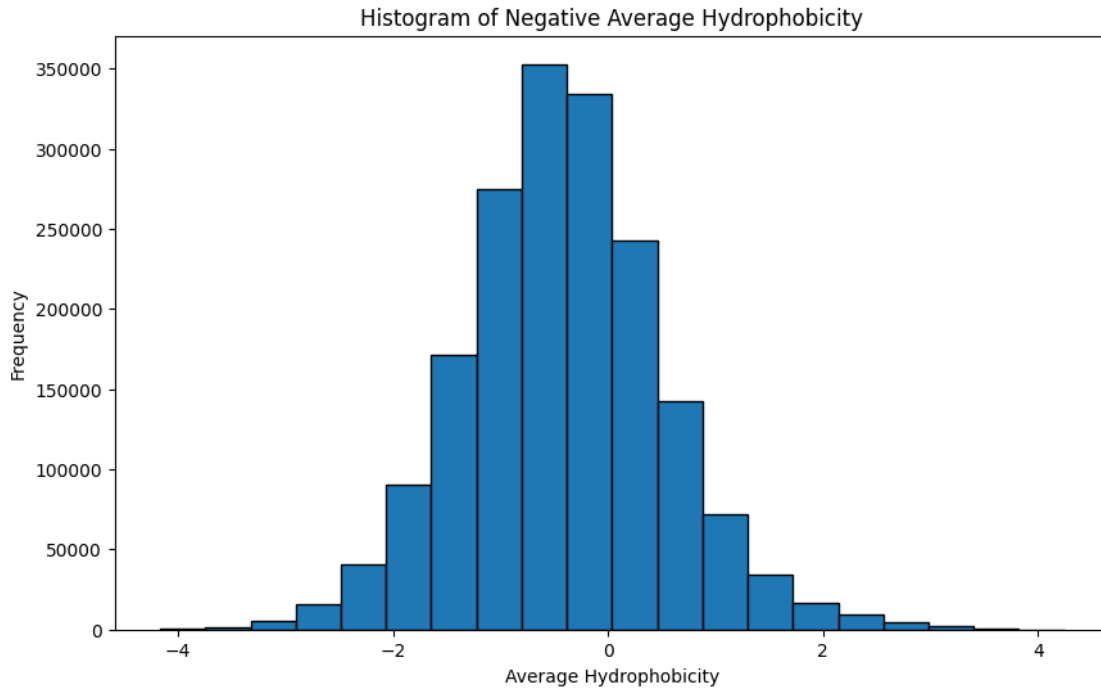
MHC Binding Affinity

0.4.3 Properties of negative samples

```
[29]: # hist of negative length
plt.figure(figsize=(10, 6))
plt.hist(negatives['peptide_length'], bins=20, edgecolor='black')
plt.xlabel('Negative Length')
plt.ylabel('Frequency')
plt.title('Histogram of Negative Length')
plt.show()
```



```
[30]: # histogram of average hydrophobicity
plt.figure(figsize=(10, 6))
plt.hist(negatives['peptide_avg_hydro'], bins=20, edgecolor='black')
plt.xlabel('Average Hydrophobicity')
plt.ylabel('Frequency')
plt.title('Histogram of Negative Average Hydrophobicity')
plt.show()
```



```
[31]: negatives['peptide_avg_hydro'].mean()
```

```
[31]: -0.4169864724628861
```

```
[32]: # plot the composition of the negatives, sort by the composition of the amino
      ↪ acids
      # Calculate mean composition and sort

      '''
      mean_composition = negatives_composition_df.mean().sort_values(ascending=False)

      # Plot the sorted composition
      plt.figure(figsize=(10, 6))
      plt.bar(mean_composition.index, mean_composition.values)
      plt.xlabel('Amino Acid')
      plt.ylabel('Composition')
      plt.title('Composition of Negative Samples')
      plt.show()
      '''
```

```
[32]: "\nmean_composition =
negatives_composition_df.mean().sort_values(ascending=False)\n\n# Plot the
sorted composition\nplt.figure(figsize=(10, 6))\nplt.bar(mean_composition.index,
mean_composition.values)\nplt.xlabel('Amino
```



```
Acid')\nplt.ylabel('Composition')\nplt.title('Composition of Negative
Samples')\nplt.show()\n"
```

0.5 Modeling

0.5.1 Data Preprocessing

```
[33]: epitopes = pd.read_csv("data/ninemer_epitopes.csv")
epitopes = epitopes.drop(columns=['epitope_objecttype', 'fullsequence',
    ↳ 'mhcrestriction_name', 'mhcrestriction_class', 'epitope_length'])
epitopes = epitopes.rename(columns={'epitope_name': 'peptide',
    ↳ 'epitope_avg_hydro': 'peptide_avg_hydro'})
epitopes_BA_pred = pd.read_csv("data/ninemer_epitopes_BA_pred.csv")
epitopes_composition = epitopes.apply(lambda row:
    ↳ count_amino_acids(row['peptide']), axis=1).apply(pd.Series)

negatives = pd.read_csv("data/ninemer_negatives_trimmed.csv")
negatives = negatives.drop(columns=['mhc', 'peptide_length'])
negatives = negatives.rename(columns={'peptide': 'peptide'})
negatives = negatives.drop_duplicates(subset=['peptide'])
negatives_BA_pred = pd.read_csv("data/ninemer_negatives_trimmed_BA_pred.csv")
negatives_BA_pred = negatives_BA_pred.drop_duplicates(subset=['peptide'])
negatives_composition = negatives.apply(lambda row:
    ↳ count_amino_acids(row['peptide']), axis=1).apply(pd.Series)

[34]: # Merge the 'Score_BA' column from epitopes_BA_pred into the epitopes dataframe
epitopes = pd.merge(epitopes, epitopes_BA_pred[['peptide', 'Score_BA',
    ↳ 'ic50']], on='peptide', how='left')
#epitopes = pd.merge(epitopes, epitopes_composition, on='peptide', how='left')

negatives = pd.merge(negatives, negatives_BA_pred[['peptide', 'Score_BA',
    ↳ 'ic50']], on='peptide', how='left')
#negatives = pd.merge(negatives, negatives_composition, on='peptide',
    ↳ how='left')

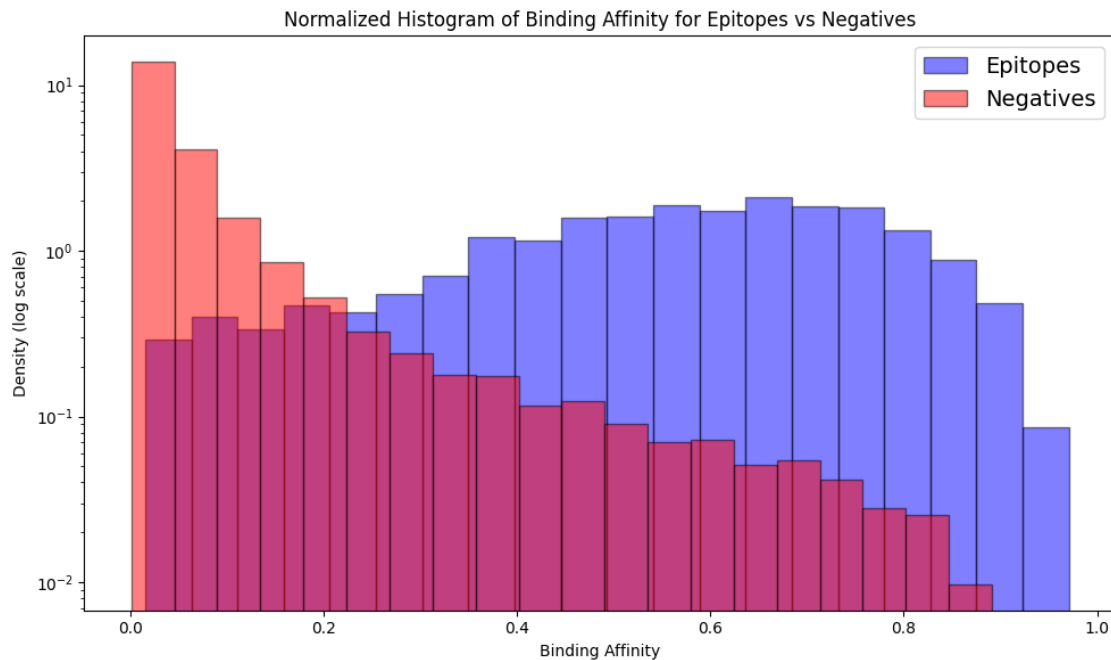
[35]: # plot Score_BA for epitopes and negatives overlaid on the same plot
plt.figure(figsize=(10, 6))

# Use density instead of raw counts to normalize the histograms
plt.hist(epitopes['Score_BA'], bins=20, alpha=0.5, color='blue',
    ↳ edgecolor='black',
        label='Epitopes', density=True)
plt.hist(negatives['Score_BA'], bins=20, alpha=0.5, color='red',
    ↳ edgecolor='black',
        label='Negatives', density=True)

# Alternative approach: use log scale for y-axis
```

```
plt.yscale('log')

plt.xlabel('Binding Affinity')
plt.ylabel('Density (log scale)')
plt.title('Normalized Histogram of Binding Affinity for Epitopes vs Negatives')
plt.legend(prop={'size': 14}) # Increased legend font size
plt.tight_layout()
plt.show()
```



```
[ ]:
```

```
[36]: # Add label column to epitopes dataframe (positive class = 1)
      epitopes['label'] = 1

      # Add label column to negatives dataframe (negative class = 0)
      negatives['label'] = 0

      # Combine the positive and negative examples
      combined_data = pd.concat([epitopes, negatives], ignore_index=True)

      # Shuffle the combined dataset
      combined_data = combined_data.sample(frac=1, random_state=42).
      ↪reset_index(drop=True)

      # Define features and target
```

```

X = combined_data.drop(columns=['peptide', 'label'])
y = combined_data['label']

# Identify numerical columns to scale (exclude one-hot encoded amino acid
↳ columns)
numerical_cols = ['peptide_avg_hydro', 'molecular_weight', 'aromaticity',
↳ 'isoelectric_point', 'instability', 'Score_BA', 'charge_at_pH7']
amino_acid_cols = [col for col in X.columns if col not in numerical_cols]

# Split the data into training and testing sets (80% train, 20% test)
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Scale numerical features using StandardScaler
scaler = StandardScaler()
X_train[numerical_cols] = scaler.fit_transform(X_train[numerical_cols])
X_test[numerical_cols] = scaler.transform(X_test[numerical_cols])

# Print the shapes to verify the split
print(f"Training set: {X_train.shape[0]} samples")
print(f"Testing set: {X_test.shape[0]} samples")
print(f"Positive samples in training: {sum(y_train == 1)}")
print(f"Negative samples in training: {sum(y_train == 0)}")
print(f"Positive samples in testing: {sum(y_test == 1)}")
print(f"Negative samples in testing: {sum(y_test == 0)}")
print(f"Scaled numerical features: {numerical_cols}")

```

```

Training set: 20502 samples
Testing set: 5126 samples
Positive samples in training: 4236
Negative samples in training: 16266
Positive samples in testing: 1059
Negative samples in testing: 4067
Scaled numerical features: ['peptide_avg_hydro', 'molecular_weight',
'aromaticity', 'isoelectric_point', 'instability', 'Score_BA', 'charge_at_pH7']

```

```

[37]: # drop the Score_BA column
      #X_train = X_train.drop(columns=['Score_BA'])
      #X_test = X_test.drop(columns=['Score_BA'])

```

```

[38]: # Initialize the Random Forest Classifier
      rf_model = RandomForestClassifier(
          n_estimators=100, # Number of trees

```

```

    max_depth=None,      # Maximum depth of trees
    min_samples_split=2,
    min_samples_leaf=1,
    random_state=42
)

# Train the model
rf_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = rf_model.predict(X_test)
y_pred_proba = rf_model.predict_proba(X_test)[:, 1] # Probability estimates_
↳ for positive class

# Evaluate the model
print("Random Forest Model Evaluation:")
print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:")
print(cm)

# Calculate ROC AUC
roc_auc = roc_auc_score(y_test, y_pred_proba)
print(f"\nROC AUC Score: {roc_auc:.4f}")

# Plot ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'Random Forest (AUC = {roc_auc:.4f})')
plt.plot([0, 1], [0, 1], 'k--', label='Random (AUC = 0.5)')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Random Forest')
plt.legend()
plt.grid(True, alpha=0.3)
plt.show()

# Feature importance
feature_importance = pd.DataFrame({
    'Feature': X_train.columns,
    'Importance': rf_model.feature_importances_
})

```

```

feature_importance = feature_importance.sort_values('Importance',
    ↪ascending=False)

# Plot top 15 features
plt.figure(figsize=(10, 6))
top_features = feature_importance.head(15)
plt.barh(np.arange(len(top_features)), top_features['Importance'],
    ↪align='center')
plt.yticks(np.arange(len(top_features)), top_features['Feature'])
plt.xlabel('Importance')
plt.title('Top 15 Feature Importance - Random Forest')
plt.tight_layout()
plt.show()

```

Random Forest Model Evaluation:

Accuracy: 0.9138

Classification Report:

	precision	recall	f1-score	support
0	0.94	0.96	0.95	4067
1	0.82	0.75	0.78	1059
accuracy			0.91	5126
macro avg	0.88	0.85	0.86	5126
weighted avg	0.91	0.91	0.91	5126

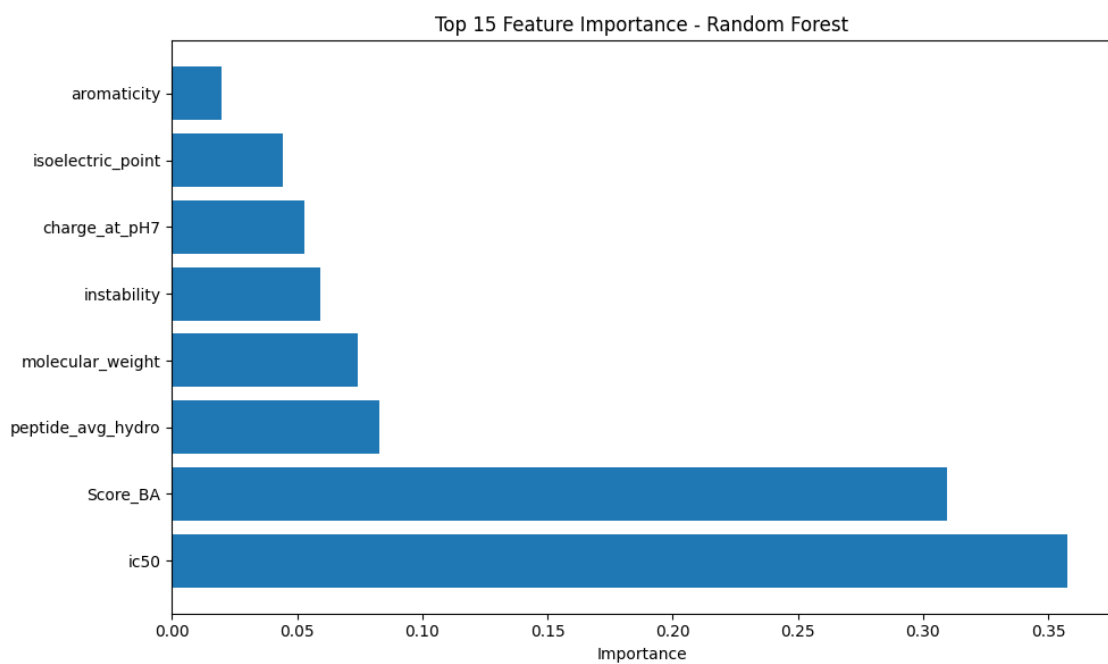
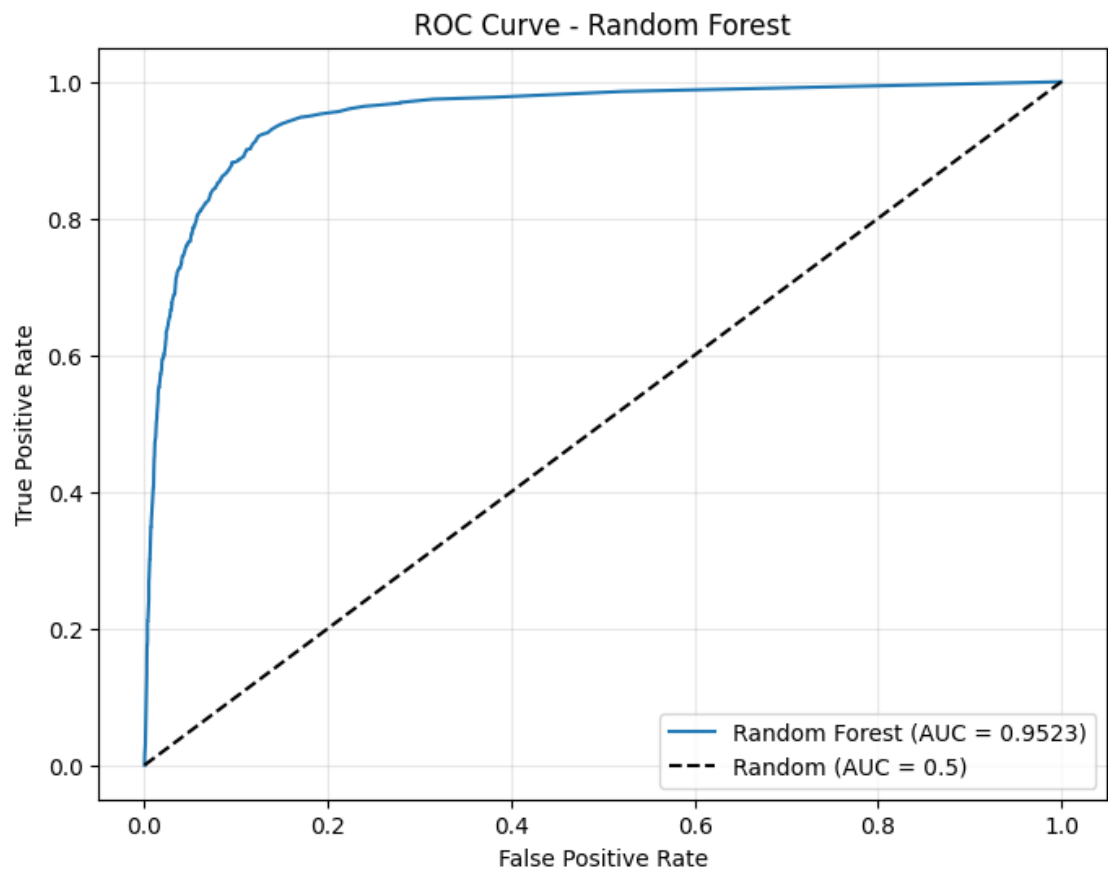
Confusion Matrix:

```

[[3893  174]
 [ 268  791]]

```

ROC AUC Score: 0.9523



0.5.2 Clustering

```
[39]: # Example clustering approach
from sklearn.cluster import KMeans, DBSCAN, AgglomerativeClustering
from sklearn.manifold import TSNE
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.impute import SimpleImputer

# Create feature matrix (using your existing features)
X = pd.concat([epitopes[['peptide_avg_hydro', 'molecular_weight', 'aromaticity',
                        'isoelectric_point', 'instability', 'charge_at_pH7',
                        'Score_BA']],
              # Add amino acid composition features
              pd.get_dummies(epitopes['peptide'].apply(lambda x: ''.join(x)),
                              prefix='pos')], axis=1)

# Handle missing values
print("Number of NaN values in dataset:", X.isna().sum().sum())
imputer = SimpleImputer(strategy='mean')
X_imputed = imputer.fit_transform(X)

# Option 1: K-means clustering
kmeans = KMeans(n_clusters=5, random_state=42) # Adjust number of clusters
clusters = kmeans.fit_predict(X_imputed)
epitopes['cluster'] = clusters

# Option 2: Hierarchical clustering
# hclust = AgglomerativeClustering(n_clusters=5)
# clusters = hclust.fit_predict(X_imputed)

# Visualize with t-SNE
tsne = TSNE(n_components=2, random_state=42)
X_tsne = tsne.fit_transform(X_imputed)

plt.figure(figsize=(10, 8))
sns.scatterplot(x=X_tsne[:, 0], y=X_tsne[:, 1], hue=clusters, palette='viridis')
plt.title('Epitope Clusters Visualization')
plt.show()

# Analyze cluster characteristics
for cluster_id in range(5):
    cluster_peptides = epitopes[epitopes['cluster'] == cluster_id]
    print(f"Cluster {cluster_id}: {len(cluster_peptides)} peptides")
    print(f"Average binding score: {cluster_peptides['Score_BA'].mean():.2f}")
```

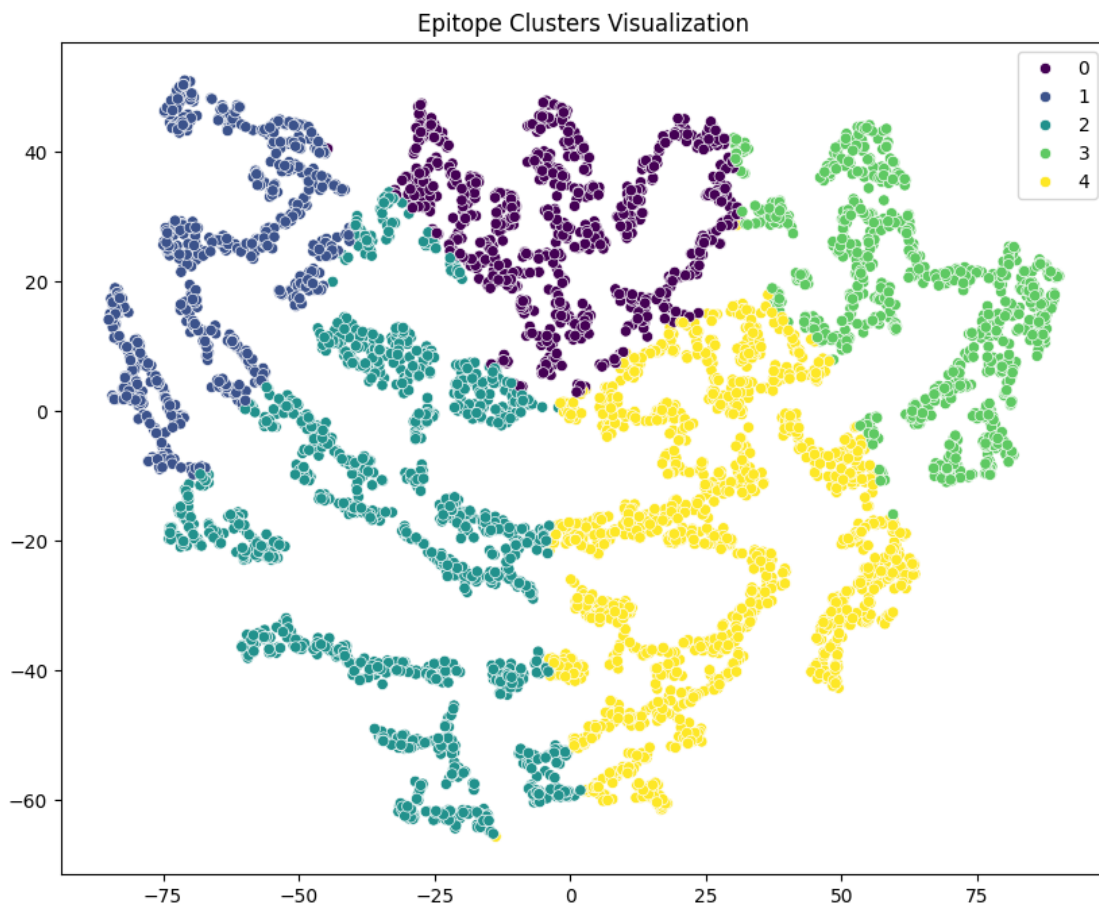
```

print(f"Average hydrophobicity: {cluster_peptides['peptide_avg_hydro'].
↪mean():.2f}")

# Find sequence motifs in cluster
motif_analysis = pd.DataFrame()
for i in range(9): # For 9-mer peptides
    aa_counts = cluster_peptides['peptide'].str[i].
↪value_counts(normalize=True)
    motif_analysis[f'Position_{i+1}'] = aa_counts
print("Top amino acids at each position:")
for col in motif_analysis.columns:
    top_aas = motif_analysis[col].nlargest(3)
    print(f"{col}: {'', ' '.join([f'{aa}({freq:.2f})' for aa, freq in top_aas.
↪items()])}")
print("\n")

```

Number of NaN values in dataset: 937



Cluster 0: 827 peptides

Average binding score: 0.55
Average hydrophobicity: 0.02
Top amino acids at each position:
Position_1: S(0.12), L(0.11), A(0.07)
Position_2: L(0.22), P(0.13), S(0.09)
Position_3: S(0.12), L(0.10), A(0.08)
Position_4: S(0.13), E(0.13), P(0.13)
Position_5: S(0.10), L(0.10), R(0.07)
Position_6: S(0.12), L(0.12), P(0.09)
Position_7: P(0.13), S(0.13), L(0.09)
Position_8: S(0.12), P(0.12), L(0.09)
Position_9: L(0.31), V(0.16), I(0.10)

Cluster 1: 732 peptides
Average binding score: 0.53
Average hydrophobicity: 0.75
Top amino acids at each position:
Position_1: A(0.20), G(0.14), S(0.14)
Position_2: L(0.31), A(0.13), P(0.13)
Position_3: A(0.16), G(0.11), S(0.11)
Position_4: G(0.17), P(0.16), A(0.15)
Position_5: G(0.19), A(0.16), V(0.11)
Position_6: G(0.16), S(0.14), L(0.12)
Position_7: A(0.13), P(0.12), S(0.11)
Position_8: A(0.18), S(0.16), G(0.12)
Position_9: L(0.29), V(0.25), A(0.14)

Cluster 2: 1419 peptides
Average binding score: 0.56
Average hydrophobicity: 0.57
Top amino acids at each position:
Position_1: L(0.10), A(0.10), S(0.09)
Position_2: L(0.32), V(0.10), T(0.08)
Position_3: L(0.14), A(0.09), S(0.09)
Position_4: S(0.09), G(0.09), L(0.09)
Position_5: L(0.11), G(0.11), A(0.09)
Position_6: L(0.14), V(0.09), S(0.09)
Position_7: L(0.14), V(0.11), A(0.09)
Position_8: L(0.11), A(0.10), S(0.10)
Position_9: L(0.28), V(0.20), K(0.12)

Cluster 3: 863 peptides
Average binding score: 0.57
Average hydrophobicity: -0.31
Top amino acids at each position:

Position_1: R(0.14), F(0.12), Y(0.12)
Position_2: L(0.19), Y(0.15), R(0.10)
Position_3: Y(0.11), F(0.11), L(0.10)
Position_4: E(0.11), R(0.10), L(0.08)
Position_5: R(0.14), F(0.11), L(0.09)
Position_6: L(0.12), F(0.10), R(0.08)
Position_7: L(0.12), F(0.09), R(0.08)
Position_8: L(0.10), E(0.09), R(0.09)
Position_9: L(0.26), F(0.18), Y(0.12)

Cluster 4: 1454 peptides

Average binding score: 0.58

Average hydrophobicity: 0.19

Top amino acids at each position:

Position_1: F(0.12), K(0.11), R(0.10)
Position_2: L(0.27), V(0.09), Y(0.08)
Position_3: L(0.13), F(0.08), D(0.07)
Position_4: E(0.11), D(0.08), L(0.08)
Position_5: L(0.11), F(0.08), V(0.08)
Position_6: L(0.15), V(0.09), I(0.09)
Position_7: L(0.14), F(0.09), V(0.06)
Position_8: L(0.13), S(0.08), F(0.08)
Position_9: L(0.28), V(0.15), F(0.11)

0.5.3 New Model