

capstone

April 29, 2025

0.1 Data import and Cleaning

```
[1]: import pandas as pd
from collections import Counter
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import Bio
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, accuracy_score, \
    confusion_matrix, roc_auc_score, roc_curve
from Bio.SeqUtils.ProtParam import ProteinAnalysis
import requests
from io import StringIO
from Bio import SeqIO

epitopes = pd.read_csv(r'/Users/tariq/Documents/capstone/data/
    ↳epitope_table_export_1740279588.csv')
assays = pd.read_csv(r'/Users/tariq/Documents/capstone/data/
    ↳tcell_table_export_1740279970.csv')

def fetch_full_sequence(url):
    if pd.isna(url): # Check if the URL is not NaN
        url = f'{url}.fasta'
    try:
        response = requests.get(url)
        if response.status_code == 200:
            fasta_io = StringIO(response.text)
            records = list(SeqIO.parse(fasta_io, "fasta"))
            if records: # Check if there are any records
                return str(records[0].seq)
            else:
                print("No records found in the FASTA file.")
    except requests.exceptions.RequestException as e:
        print(f"Request failed: {e}")
    return None
```

```

#epitopes['Full Sequence'] = epitopes['Epitope - Molecule Parent IRI'].
↳apply(fetch_full_sequence)
epitopes = pd.read_csv(r'/Users/tariq/Documents/capstone/data/epitope_full_seq.
↳csv')

# make all the column names snake case
epitopes.columns = epitopes.columns.str.lower()
assays.columns = assays.columns.str.lower()

# remove spaces from column names
epitopes.columns = epitopes.columns.str.replace(' ', '')
epitopes.columns = epitopes.columns.str.replace('-', '_')
epitopes.columns = epitopes.columns.str.replace(' ', '_')

assays.columns = assays.columns.str.replace(' ', '')
assays.columns = assays.columns.str.replace('-', '_')
assays.columns = assays.columns.str.replace(' ', '_')

epitopes = epitopes.filter(['epitope_name', 'fullsequence'])
assays = assays.filter(['epitope_name', 'epitope_moluculeparent', 'host_name',
↳'host_mhcrestriction', 'assay_method', 'assay_response', 'assay_response_measurement',
↳'assay_qualitative_measurement', 'mhcrestriction_name',
↳'mhcrestriction_class', 'assay_antigen_name'])

# map mhc name and class from the assays dataframe to a new column in the
↳epitopes dataframe based on epitope_name
mhc = assays.filter(['epitope_name', 'mhcrestriction_name',
↳'mhcrestriction_class'])
mhc = mhc.drop_duplicates(subset=['epitope_name'])
epitopes = epitopes.merge(mhc, on='epitope_name', how='left')

```

```

/var/folders/5j/4p7c5_1x2fg18bk0nf74_hg40000gn/T/ipykernel_53549/2635103461.py:1
5: DtypeWarning: Columns (13,14,45,46,47,48,49,54,55,56,57,60,65,66,67,68,69,70,
71,72,73,74,75,76,77,78,79,82,83,84,85,86,87,88,89,90,91,92,93,94,95,96,97,98,99
,100,101,102,105,106,107,108,109,110,111,112,113,115,120,123,128,132,134,135,141
,142,143,144,145,149,152) have mixed types. Specify dtype option on import or
set low_memory=False.
assays = pd.read_csv(r'/Users/tariq/Documents/capstone/data/tcell_table_export
_1740279970.csv')

```

```
[2]: epitopes.head()
```

```

[2]:   epitope_name      fullsequence \
0   AAGIGILTV  MPREDAHFYGYPPKKGHGHSTTAEEAAGIGILTVILGVLLLLIGCW...
1   AAGIGILTVI  MPREDAHFYGYPPKKGHGHSTTAEEAAGIGILTVILGVLLLLIGCW...
2   ACDPHSGHFV                                     NaN
3   ADLVGFLLLK  MSLEQRLHCKPEEALAQEALGLVCVQAATSSSSPLVLGTLEEV...

```

```
4 ADVEFCLSL MLLAVLYCLLSFQTSAGHFPRACVSSKNLMEKECCPPWSGDRSPC...
```

```

mhcrestriction_name mhcrestriction_class
0          HLA-A2          I
1      HLA-A*02:01          I
2          HLA-A2          I
3      HLA-A*11:01          I
4      HLA-B*44:03          I

```

```
[3]: assays.head()
```

```

[3]:      epitope_name      host_name  host_mhcrestriction assay_method \
0      APIWPYEILY  Homo sapiens (human)          NaN      ELISPOT
1      LIYDSSLCDL      Homo sapiens          NaN      ELISPOT
2  DRAHYNIVTFCKCD  Homo sapiens (human)      HLA-DR15      ELISPOT
3  DRAHYNIVTFCKCD  Homo sapiens (human)      HLA-DR15      ELISPOT
4      MHGDTPTLHEYM  Homo sapiens (human)  HLA-DR15;HLA-DR4      ELISPOT

      assay_response measured assay_qualitative measurement mhcrestriction_name \
0      IFNg release          Positive          HLA-B*35:01
1      IFNg release          Positive          HLA-A2
2      IFNg release          Positive          HLA-DR15
3      IL-5 release          Positive          HLA-DR15
4      IL-5 release          Positive          HLA class II

      mhcrestriction_class assayantigen_name
0          I      APIWPYEILY
1          I      LIYDSSLCDL
2          II  DRAHYNIVTFCKCD
3          II  DRAHYNIVTFCKCD
4          II      MHGDTPTLHEYM

```

0.2 Feature Engineering

```
[4]: epitopes['epitope_length'] = epitopes['epitope_name'].str.len()
```

```

[5]: # Function to count amino acids in a peptide
def count_amino_acids(peptide):
    try:
        # Create a ProteinAnalysis object for the peptide
        analyzer = ProteinAnalysis(peptide)
        # Get amino acid counts and normalize to frequencies
        aa_count = analyzer.count_amino_acids()
        total_aa = sum(aa_count.values())
        aa_freq = {aa: count for aa, count in aa_count.items()}
        # Add the peptide itself to the results
        aa_freq['peptide'] = peptide

```

```

        return aa_freq
    except Exception as e:
        # Handle invalid peptides (e.g., with non-standard amino acids)
        result = {aa: 0 for aa in 'ACDEFGHIKLMNPQRSTVWY'}
        result['peptide'] = peptide
        return result

# Create analyzer function that will be used in the next cell
def analyzer(peptide):
    return count_amino_acids(peptide)

# Use both epitope name and peptide sequence in the DataFrame
epitope_composition_df = epitopes.apply(lambda row:
    ↪count_amino_acids(row['epitope_name']), axis=1).apply(pd.Series)

```

```
[6]: epitope_composition_df.head()
```

```

[6]:   A  C  D  E  F  G  H  I  K  L  ...  N  P  Q  R  S  T  V  W  Y    peptide
0  2  0  0  0  0  2  0  2  0  1  ...  0  0  0  0  0  1  1  0  0    AAGIGILTV
1  2  0  0  0  0  2  0  3  0  1  ...  0  0  0  0  0  1  1  0  0    AAGIGILTVI
2  1  1  1  0  1  1  2  0  0  0  ...  0  1  0  0  1  0  1  0  0    ACDPHSGHFV
3  1  0  1  0  1  1  0  0  1  4  ...  0  0  0  0  0  0  1  0  0    ADLVGFLLLK
4  1  1  1  1  1  0  0  0  0  2  ...  0  0  0  0  1  0  1  0  0    ADVEFCLSL

```

[5 rows x 21 columns]

```

[7]: # Example DataFrame with a 'peptide' column
df = pd.DataFrame({
    'peptide': ['ACDEFGHIK', 'LMNPQRSTV', 'WYFP']
})

# Kyte-Doolittle hydrophobicity scale
kyte_doolittle = {
    'I': 4.5, 'V': 4.2, 'L': 3.8, 'F': 2.8, 'C': 2.5,
    'M': 1.9, 'A': 1.8, 'G': -0.4, 'T': -0.7, 'S': -0.8,
    'W': -0.9, 'Y': -1.3, 'P': -1.6, 'H': -3.2, 'E': -3.5,
    'Q': -3.5, 'D': -3.5, 'N': -3.5, 'K': -3.9, 'R': -4.5
}

def compute_avg_hydrophobicity(peptide):
    # Get hydrophobicity scores for each amino acid; default to 0 if missing
    scores = [kyte_doolittle.get(aa, 0) for aa in peptide]
    return sum(scores) / len(scores) if scores else 0

# Apply the function to the 'peptide' column to create a new column 'avg_hydro'
epitopes['epitope_avg_hydro'] = epitopes['epitope_name'].
    ↪apply(compute_avg_hydrophobicity)

```

```
[8]: # Import the molecular_weight function from Bio.SeqUtils

def calculate_molecular_weight(peptide):
    """Calculate the molecular weight of a peptide sequence using Biopython."""
    try:
        # ProteinAnalysis only works with standard amino acids
        protein = ProteinAnalysis(peptide)
        return protein.molecular_weight()
    except Exception as e:
        # Handle peptides with non-standard amino acids
        return None

# Apply the function to calculate molecular weight for each epitope
epitopes['molecular_weight'] = epitopes['epitope_name'].
    ↪apply(calculate_molecular_weight)
```

```
[9]: def calculate_aromaticity(peptide):
    """Calculate the aromaticity of a peptide sequence using Biopython."""
    try:
        # ProteinAnalysis only works with standard amino acids
        protein = ProteinAnalysis(peptide)
        return protein.aromaticity()
    except Exception as e:
        # Handle peptides with non-standard amino acids
        return None

# Apply the function to calculate molecular weight for each epitope
epitopes['aromaticity'] = epitopes['epitope_name'].apply(calculate_aromaticity)
```

```
[10]: def calculate_isoelectric_point(peptide):
    """Calculate the isoelectric point of a peptide sequence using Biopython."""
    try:
        # ProteinAnalysis only works with standard amino acids
        protein = ProteinAnalysis(peptide)
        return protein.isoelectric_point()
    except Exception as e:
        # Handle peptides with non-standard amino acids
        return None

# Apply the function to calculate molecular weight for each epitope
epitopes['isoelectric_point'] = epitopes['epitope_name'].
    ↪apply(calculate_isoelectric_point)
```

```
[11]: def calculate_instability(peptide):
    """Calculate the instability of a peptide sequence using Biopython."""
    try:
```

```

    # ProteinAnalysis only works with standard amino acids
    protein = ProteinAnalysis(peptide)
    return protein.instability_index()
except Exception as e:
    # Handle peptides with non-standard amino acids
    return None

# Apply the function to calculate molecular weight for each epitope
epitopes['instability'] = epitopes['epitope_name'].apply(calculate_instability)

```

```

[12]: def calculate_charge_at_pH7(peptide):
    """Calculate the charge of a peptide sequence at pH 7 using Biopython."""
    try:
        # ProteinAnalysis only works with standard amino acids
        protein = ProteinAnalysis(peptide)
        return protein.charge_at_pH(7)
    except Exception as e:
        # Handle peptides with non-standard amino acids
        return None

# Apply the function to calculate molecular weight for each epitope
epitopes['charge_at_pH7'] = epitopes['epitope_name'].
    ↪ apply(calculate_charge_at_pH7)

```

```

[13]: epitopes.head()

```

```

[13]: epitope_name          fullsequence \
0    AAGIGILTV  MPREDAHFIYGYPKKKGHGSYTTAEAAAGIGILTVILGVLLLLIGCW...
1    AAGIGILTVI  MPREDAHFIYGYPKKKGHGSYTTAEAAAGIGILTVILGVLLLLIGCW...
2    ACDPHSGHFV                                     NaN
3    ADLVGFLLLK  MSLEQRSLHCKPEEALAEQAEALGLVCVQAATSSSSPLVLGTLEEV...
4    ADVEFCLSL  MLLAVLYCLLWSFQTSAGHFPRACVSSKNLMEKECCPPWSGDRSPC...

  mhcrestriction_name mhcrestriction_class  epitope_length  epitope_avg_hydro \
0             HLA-A2                I                9          2.122222
1          HLA-A*02:01                I               10          2.360000
2             HLA-A2                I               10         -0.140000
3          HLA-A*11:01                I               10          1.620000
4          HLA-B*44:03                I                9          1.233333

  molecular_weight  aromaticity  isoelectric_point  instability \
0           813.9814      0.000000        5.570017      11.422222
1           927.1390      0.000000        5.570017      11.280000
2          1069.1507      0.100000        5.972266      61.830000
3          1088.3394      0.100000        5.880358     -16.470000
4           996.1348      0.111111        4.050028      20.855556

```

	charge_at_pH7
0	-0.204125
1	-0.204125
2	-1.038557
3	-0.204004
4	-2.210095

0.3 Generation of Negative Samples

```
[14]: def generate_negatives(row):
    epitope = row["epitope_name"]
    full_seq = row["fullsequence"]
    mhc = row["mhcrestriction_name"]

    # Handle missing or empty sequences
    if pd.isnull(full_seq) or full_seq == "":
        return []

    epitope = str(epitope)
    full_seq = str(full_seq)
    ep_len = len(epitope)

    negatives = []
    for i in range(len(full_seq) - ep_len + 1):
        window = full_seq[i:i+ep_len]
        if window != epitope:
            negatives.append({"peptide": window, "mhc": mhc})
    return negatives

'''
# Apply the function to each row

negatives = pd.DataFrame()
negatives['negatives'] = epitopes.apply(generate_negatives, axis=1)
negatives = negatives[["negatives"]].explode("negatives").reset_index(drop=True)
negatives.dropna(subset=["negatives"], inplace=True)

# Remove duplicate peptide-mhc combinations
print(f"Shape before removing duplicates: {negatives.shape}")
negatives = negatives.drop_duplicates(subset=['negatives'])
print(f"Shape after removing duplicates: {negatives.shape}")

# Check for any remaining NaN values
print(f"Number of NaN values in negatives: {negatives['negatives'].isna().
      ↪sum()}")
```

```

# Extract peptide and mhc into separate columns
negatives['peptide'] = negatives['negatives'].apply(lambda x: x['peptide'])
negatives['mhc'] = negatives['negatives'].apply(lambda x: x['mhc'])

# Calculate features on the peptide column
negatives['peptide_length'] = negatives['peptide'].apply(len)
negatives['peptide_avg_hydro'] = negatives['peptide'].
    ↪ apply(compute_avg_hydrophobicity)
negatives['molecular_weight'] = negatives['peptide'].
    ↪ apply(calculate_molecular_weight)
negatives['aromaticity'] = negatives['peptide'].apply(calculate_aromaticity)
negatives['isoelectric_point'] = negatives['peptide'].
    ↪ apply(calculate_isoelectric_point)
negatives['instability'] = negatives['peptide'].apply(calculate_instability)
negatives['charge_at_pH7'] = negatives['peptide'].apply(calculate_charge_at_pH7)

# Drop the original dictionary column if no longer needed
negatives.drop('negatives', axis=1, inplace=True)
'''

```

```

[14]: '\n# Apply the function to each row\n\nnegatives =
pd.DataFrame()\nnegatives['negatives'] = epitopes.apply(generate_negatives,
axis=1)\nnegatives = negatives[['negatives']].explode("negatives").reset_index(drop=True)\nnegatives.dropna(subset=["negatives"], inplace=True)\n\n# Remove
duplicate peptide-mhc combinations\nprint(f"Shape before removing duplicates:
{negatives.shape}")\nnegatives =
negatives.drop_duplicates(subset=['negatives'])\nprint(f"Shape after removing
duplicates: {negatives.shape}")\n\n# Check for any remaining NaN
values\nprint(f"Number of NaN values in negatives:
{negatives['negatives'].isna().sum()}")\n\n# Extract peptide and mhc into
separate columns\nnegatives['peptide'] = negatives['negatives'].apply(lambda
x: x['peptide'])\nnegatives['mhc'] = negatives['negatives'].apply(lambda
x: x['mhc'])\n\n# Calculate features on the peptide
column\nnegatives['peptide_length'] =
negatives['peptide'].apply(len)\nnegatives['peptide_avg_hydro'] = negatives[
'peptide'].apply(compute_avg_hydrophobicity)\nnegatives['molecular_weight']
= negatives['peptide'].apply(calculate_molecular_weight)\nnegatives['aromatic
ity'] = negatives['peptide'].apply(calculate_aromaticity)\nnegatives['isoele
ctric_point'] = negatives['peptide'].apply(calculate_isoelectric_point)\nnega
tives['instability'] = negatives['peptide'].apply(calculate_instability)\nne
gatives['charge_at_pH7'] =
negatives['peptide'].apply(calculate_charge_at_pH7)\n\n# Drop the original
dictionary column if no longer needed\nnegatives.drop('negatives', axis=1,
inplace=True)\n'

```

```

[15]: negatives = pd.read_csv("data/negatives_MHC.csv")

```



```
/var/folders/5j/4p7c5_1x2fg18bk0nf74_hg40000gn/T/ipykernel_53549/1811011591.py:1
: DtypeWarning: Columns (1) have mixed types. Specify dtype option on import or
set low_memory=False.
negatives = pd.read_csv("data/negatives_MHC.csv")
```

```
[16]: nine_mers = epitopes[epitopes['epitope_length'] == 9]
```

```
[17]: ninemer_negatives = negatives[negatives['peptide_length'] == 9]
ninemer_negatives_trimmed = ninemer_negatives[:50000]
```

0.4 EDA

0.4.1 Data Summary

```
[18]: epitopes.head()
```

```
[18]: epitope_name fullsequence \
0 AAGIGILTV MPREDAHFIYGYPKKGHGHSYTTAEAAAGIGILTVILGVLLIGCW...
1 AAGIGILTVI MPREDAHFIYGYPKKGHGHSYTTAEAAAGIGILTVILGVLLIGCW...
2 ACDPHSGHFV NaN
3 ADLVGFLLK MSLEQRLHCKPEEALAEALGLVCVQAATSSSSPLVLGTLEEV...
4 ADVEFCLSL MLLAVLYCLLWSFQTSAGHFPRACVSSKNLMEKECCPPWSGDRSPC...
```

	mhcrestriction_name	mhcrestriction_class	epitope_length	epitope_avg_hydro	\
0	HLA-A2	I	9	2.122222	
1	HLA-A*02:01	I	10	2.360000	
2	HLA-A2	I	10	-0.140000	
3	HLA-A*11:01	I	10	1.620000	
4	HLA-B*44:03	I	9	1.233333	

	molecular_weight	aromaticity	isoelectric_point	instability	\
0	813.9814	0.000000	5.570017	11.422222	
1	927.1390	0.000000	5.570017	11.280000	
2	1069.1507	0.100000	5.972266	61.830000	
3	1088.3394	0.100000	5.880358	-16.470000	
4	996.1348	0.111111	4.050028	20.855556	

	charge_at_pH7
0	-0.204125
1	-0.204125
2	-1.038557
3	-0.204004
4	-2.210095

```
[19]: epitopes.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 28681 entries, 0 to 28680
```

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	epitope_name	28681 non-null	object
1	fullsequence	7164 non-null	object
2	mhcrestriction_name	17613 non-null	object
3	mhcrestriction_class	17613 non-null	object
4	epitope_length	28681 non-null	int64
5	epitope_avg_hydro	28681 non-null	float64
6	molecular_weight	28623 non-null	float64
7	aromaticity	28681 non-null	float64
8	isoelectric_point	28681 non-null	float64
9	instability	28623 non-null	float64
10	charge_at_pH7	28681 non-null	float64


dtypes: float64(6), int64(1), object(4)

memory usage: 2.4+ MB

0.4.2 Properties of Epitopes

Length

```
[20]: # hist of epitope length
plt.figure(figsize=(10, 6))
plt.hist(epitopes['epitope_length'], bins=20, edgecolor='black')
plt.xlabel('Epitope Length')
plt.ylabel('Frequency')
plt.title('Histogram of Epitope Length')
plt.show()
```




```
[21]: epitopes['epitope_length'].describe()
```

```
[21]: count      28681.000000
      mean       19.389422
      std        8.255925
      min        6.000000
      25%       10.000000
      50%       20.000000
      75%       25.000000
      max       50.000000
      Name: epitope_length, dtype: float64
```

Hydrophobicity

```
[22]: # histogram of average hydrophobicity
      plt.figure(figsize=(10, 6))
      plt.hist(epitopes['epitope_avg_hydro'], bins=20, edgecolor='black')
      plt.xlabel('Average Hydrophobicity')
      plt.ylabel('Frequency')
      plt.title('Histogram of Average Epitope Hydrophobicity')
      plt.show()
```



```
[23]: epitopes['epitope_avg_hydro'].describe()
```

```
[23]: count      28681.000000
      mean       -0.178410
      std        0.883064
      min       -3.312000
      25%       -0.762500
      50%       -0.240000
      75%        0.333333
      max        3.688889
      Name: epitope_avg_hydro, dtype: float64
```

Composition

```
[24]: # plot the composition of the epitopes, sort by the composition of the amino
      ↪ acids
      # Calculate mean composition and sort

      '''
      mean_composition = epitope_composition_df.mean().sort_values(ascending=False)

      # Plot the sorted composition
      plt.figure(figsize=(10, 6))
      plt.bar(mean_composition.index, mean_composition.values)
      plt.xlabel('Amino Acid')
      plt.ylabel('Composition')
      plt.title('Composition of Epitopes')
      plt.show()

      '''
```

```
[24]: "\nmean_composition =
      epitope_composition_df.mean().sort_values(ascending=False)\n\n# Plot the sorted
      composition\nnplt.figure(figsize=(10, 6))\nnplt.bar(mean_composition.index,
      mean_composition.values)\nnplt.xlabel('Amino
      Acid')\nnplt.ylabel('Composition')\nnplt.title('Composition of
      Epitopes')\nnplt.show()\n\n"
```

n-gram frequency analysis

```
[25]: def ngram_frequency(peptides, n=2):
      ngrams = []
      for peptide in peptides:
          if len(peptide) < n:
              continue
          for i in range(len(peptide) - n + 1):
              ngram = peptide[i:i+n]
              ngrams.append(ngram)
```

```

return Counter(ngrams)

diptptide_freq = ngram_frequency(epitopes['epitope_name'], n=2)

df_ngram = pd.DataFrame(diptptide_freq.items(), columns=['ngram', 'count'])
df_ngram = df_ngram.sort_values('count', ascending=False)

top_n = 20
top_ngram = df_ngram.head(top_n)


plt.figure(figsize=(12, 6))
sns.barplot(x='ngram', y='count', data=top_ngram, palette="viridis")
plt.title(f"Epitopes Top {top_n} Diptptides Frequency")
plt.xlabel("Diptptide")
plt.ylabel("Frequency")
plt.show()

```

/var/folders/5j/4p7c5_1x2fg18bk0nf74_hg40000gn/T/ipykernel_53549/733366050.py:20
: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.


```
sns.barplot(x='ngram', y='count', data=top_ngram, palette="viridis")
```




MHC Binding Affinity

0.4.3 Properties of negative samples

```
[26]: # hist of negative length
plt.figure(figsize=(10, 6))
plt.hist(negatives['peptide_length'], bins=20, edgecolor='black')
plt.xlabel('Negative Length')
plt.ylabel('Frequency')
plt.title('Histogram of Negative Length')
plt.show()
```



```
[27]: # histogram of average hydrophobicity
plt.figure(figsize=(10, 6))
plt.hist(negatives['peptide_avg_hydro'], bins=20, edgecolor='black')
plt.xlabel('Average Hydrophobicity')
plt.ylabel('Frequency')
plt.title('Histogram of Negative Average Hydrophobicity')
plt.show()
```



```
[28]: negatives['peptide_avg_hydro'].mean()
```

```
[28]: np.float64(-0.4169864724628861)
```

```
[29]: # plot the composition of the negatives, sort by the composition of the amino
      ↪ acids
      # Calculate mean composition and sort

      '''
      mean_composition = negatives_composition_df.mean().sort_values(ascending=False)

      # Plot the sorted composition
      plt.figure(figsize=(10, 6))
      plt.bar(mean_composition.index, mean_composition.values)
      plt.xlabel('Amino Acid')
      plt.ylabel('Composition')
      plt.title('Composition of Negative Samples')
      plt.show()
      '''
```

```
[29]: "\nmean_composition =
negatives_composition_df.mean().sort_values(ascending=False)\n\n# Plot the
sorted composition\nplt.figure(figsize=(10, 6))\nplt.bar(mean_composition.index,
mean_composition.values)\nplt.xlabel('Amino
```

```
Acid')\nplt.ylabel('Composition')\nplt.title('Composition of Negative Samples')\nplt.show()\n"
```

0.5 Modeling

0.5.1 Data Preprocessing

```
[30]: epitopes = pd.read_csv("data/ninemer_epitopes.csv")
epitopes = epitopes.drop(columns=['fullsequence', 'mhcrestriction_name',
    ↳ 'mhcrestriction_class', 'epitope_length'])
epitopes = epitopes.rename(columns={'epitope_name': 'peptide',
    ↳ 'epitope_avg_hydro': 'peptide_avg_hydro'})
epitopes_BA_pred = pd.read_csv("data/ninemer_epitopes_BA_pred.csv")
epitopes_composition = epitopes.apply(lambda row:
    ↳ count_amino_acids(row['peptide']), axis=1).apply(pd.Series)

negatives = pd.read_csv("data/ninemer_negatives_trimmed.csv")
negatives = negatives.drop(columns=['mhc', 'peptide_length'])
negatives = negatives.rename(columns={'peptide': 'peptide'})
negatives = negatives.drop_duplicates(subset=['peptide'])
negatives_BA_pred = pd.read_csv("data/ninemer_negatives_trimmed_BA_pred.csv")
negatives_BA_pred = negatives_BA_pred.drop_duplicates(subset=['peptide'])
negatives_composition = negatives.apply(lambda row:
    ↳ count_amino_acids(row['peptide']), axis=1).apply(pd.Series)

[31]: # Merge the 'Score_BA' column from epitopes_BA_pred into the epitopes dataframe
epitopes = pd.merge(epitopes, epitopes_BA_pred[['peptide', 'Score_BA',
    ↳ 'ic50']], on='peptide', how='left')
#epitopes = pd.merge(epitopes, epitopes_composition, on='peptide', how='left')

negatives = pd.merge(negatives, negatives_BA_pred[['peptide', 'Score_BA',
    ↳ 'ic50']], on='peptide', how='left')
#negatives = pd.merge(negatives, negatives_composition, on='peptide',
    ↳ how='left')

[32]: # plot Score_BA for epitopes and negatives overlaid on the same plot
plt.figure(figsize=(10, 6))


# Use density instead of raw counts to normalize the histograms
plt.hist(epitopes['Score_BA'], bins=20, alpha=0.5, color='blue',
    ↳ edgecolor='black',
        label='Epitopes', density=True)
plt.hist(negatives['Score_BA'], bins=20, alpha=0.5, color='red',
    ↳ edgecolor='black',
        label='Negatives', density=True)

# Alternative approach: use log scale for y-axis
```



```
plt.yscale('log')

plt.xlabel('Binding Affinity')
plt.ylabel('Density (log scale)')
plt.title('Normalized Histogram of Binding Affinity for Epitopes vs Negatives')
plt.legend(prop={'size': 14}) # Increased legend font size
plt.tight_layout()
plt.show()
```



```
[33]: # Add label column to epitopes dataframe (positive class = 1)
      epitopes['label'] = 1

      # Add label column to negatives dataframe (negative class = 0)
      negatives['label'] = 0

      # Combine the positive and negative examples
      combined_data = pd.concat([epitopes, negatives], ignore_index=True)

      # Shuffle the combined dataset
      combined_data = combined_data.sample(frac=1, random_state=42).
      ↪reset_index(drop=True)

      # Define features and target
      X = combined_data.drop(columns=['peptide', 'label'])
      y = combined_data['label']
```

```

# Identify numerical columns to scale (exclude one-hot encoded amino acid
↳ columns)
numerical_cols = ['peptide_avg_hydro', 'molecular_weight', 'aromaticity',
↳ 'isoelectric_point', 'instability', 'Score_BA', 'charge_at_pH7']
amino_acid_cols = [col for col in X.columns if col not in numerical_cols]

# Split the data into training and testing sets (80% train, 20% test)
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Scale numerical features using StandardScaler
scaler = StandardScaler()
X_train[numerical_cols] = scaler.fit_transform(X_train[numerical_cols])
X_test[numerical_cols] = scaler.transform(X_test[numerical_cols])

# Print the shapes to verify the split
print(f"Training set: {X_train.shape[0]} samples")
print(f"Testing set: {X_test.shape[0]} samples")
print(f"Positive samples in training: {sum(y_train == 1)}")
print(f"Negative samples in training: {sum(y_train == 0)}")
print(f"Positive samples in testing: {sum(y_test == 1)}")
print(f"Negative samples in testing: {sum(y_test == 0)}")
print(f"Scaled numerical features: {numerical_cols}")

```

```

Training set: 20502 samples
Testing set: 5126 samples
Positive samples in training: 4236
Negative samples in training: 16266
Positive samples in testing: 1059
Negative samples in testing: 4067
Scaled numerical features: ['peptide_avg_hydro', 'molecular_weight',
'aromaticity', 'isoelectric_point', 'instability', 'Score_BA', 'charge_at_pH7']

```

```

[34]: # drop the Score_BA column
#X_train = X_train.drop(columns=['Score_BA'])
#X_test = X_test.drop(columns=['Score_BA'])

```

```

[35]: # Initialize the Random Forest Classifier
rf_model = RandomForestClassifier(
    n_estimators=100, # Number of trees
    max_depth=None, # Maximum depth of trees
    min_samples_split=2,

```

```

        min_samples_leaf=1,
        random_state=42
    )

    # Train the model
    rf_model.fit(X_train, y_train)

    # Make predictions on the test set
    y_pred = rf_model.predict(X_test)
    y_pred_proba = rf_model.predict_proba(X_test)[:, 1] # Probability estimates
    # for positive class

    # Evaluate the model
    print("Random Forest Model Evaluation:")
    print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
    print("\nClassification Report:")
    print(classification_report(y_test, y_pred))

    # Confusion Matrix
    cm = confusion_matrix(y_test, y_pred)
    print("\nConfusion Matrix:")
    print(cm)

    # Calculate ROC AUC
    roc_auc = roc_auc_score(y_test, y_pred_proba)
    print(f"\nROC AUC Score: {roc_auc:.4f}")

    # Plot ROC Curve
    fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
    plt.figure(figsize=(8, 6))
    plt.plot(fpr, tpr, label=f'Random Forest (AUC = {roc_auc:.4f})')
    plt.plot([0, 1], [0, 1], 'k--', label='Random (AUC = 0.5)')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC Curve - Random Forest')
    plt.legend()
    plt.grid(True, alpha=0.3)
    plt.show()

    # Feature importance
    feature_importance = pd.DataFrame({
        'Feature': X_train.columns,
        'Importance': rf_model.feature_importances_
    })
    feature_importance = feature_importance.sort_values('Importance',
    # ascending=True)

```

```

# Plot top 15 features
plt.figure(figsize=(10, 6))
top_features = feature_importance.head(15)
plt.barh(np.arange(len(top_features)), top_features['Importance'],
         align='center')
plt.yticks(np.arange(len(top_features)), top_features['Feature'])
plt.xlabel('Importance')
plt.title('Top 15 Feature Importance - Random Forest')
plt.tight_layout()
plt.show()

```

Random Forest Model Evaluation:

Accuracy: 0.9138

Classification Report:

	precision	recall	f1-score	support
0	0.94	0.96	0.95	4067
1	0.82	0.75	0.78	1059
accuracy			0.91	5126
macro avg	0.88	0.85	0.86	5126
weighted avg	0.91	0.91	0.91	5126



Confusion Matrix:

```

[[3893  174]
 [ 268  791]]

```

ROC AUC Score: 0.9523



0.5.2 Clustering

```
[36]: # Example clustering approach
from sklearn.cluster import KMeans, DBSCAN, AgglomerativeClustering
from sklearn.manifold import TSNE
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.impute import SimpleImputer

# Create feature matrix (using your existing features)
X = pd.concat([epitopes[['peptide_avg_hydro', 'molecular_weight', 'aromaticity',
                        'isoelectric_point', 'instability', 'charge_at_pH7',
                        'Score_BA']],
              # Add amino acid composition features
              pd.get_dummies(epitopes['peptide'].apply(lambda x: ''.join(x)),
                              prefix='pos')], axis=1)

# Handle missing values
print("Number of NaN values in dataset:", X.isna().sum().sum())
imputer = SimpleImputer(strategy='mean')
X_imputed = imputer.fit_transform(X)

# Option 1: K-means clustering
kmeans = KMeans(n_clusters=5, random_state=42) # Adjust number of clusters
clusters = kmeans.fit_predict(X_imputed)
epitopes['cluster'] = clusters

# Option 2: Hierarchical clustering
# hclust = AgglomerativeClustering(n_clusters=5)
# clusters = hclust.fit_predict(X_imputed)

# Visualize with t-SNE
tsne = TSNE(n_components=2, random_state=42)
X_tsne = tsne.fit_transform(X_imputed)

plt.figure(figsize=(10, 8))
sns.scatterplot(x=X_tsne[:, 0], y=X_tsne[:, 1], hue=clusters, palette='viridis')
plt.title('Epitope Clusters Visualization')
plt.show()

# Analyze cluster characteristics
for cluster_id in range(5):
    cluster_peptides = epitopes[epitopes['cluster'] == cluster_id]
    print(f"Cluster {cluster_id}: {len(cluster_peptides)} peptides")
    print(f"Average binding score: {cluster_peptides['Score_BA'].mean():.2f}")
```

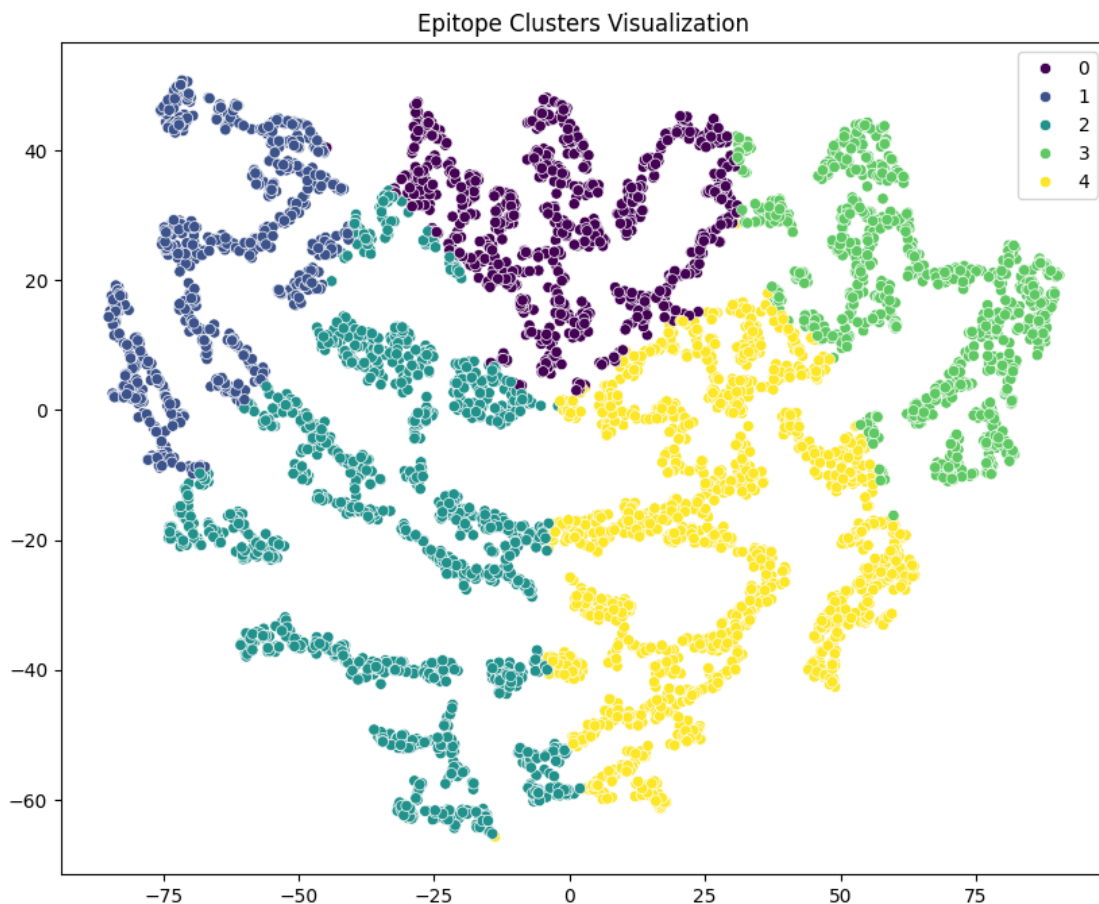
```

print(f"Average hydrophobicity: {cluster_peptides['peptide_avg_hydro'].
↪mean():.2f}")

# Find sequence motifs in cluster
motif_analysis = pd.DataFrame()
for i in range(9): # For 9-mer peptides
    aa_counts = cluster_peptides['peptide'].str[i].
↪value_counts(normalize=True)
    motif_analysis[f'Position_{i+1}'] = aa_counts
print("Top amino acids at each position:")
for col in motif_analysis.columns:
    top_aas = motif_analysis[col].nlargest(3)
    print(f"{col}: {'', ' '.join([f'{aa}({freq:.2f})' for aa, freq in top_aas.
↪items()])}")
print("\n")

```

Number of NaN values in dataset: 937



Cluster 0: 827 peptides

Average binding score: 0.55
Average hydrophobicity: 0.02
Top amino acids at each position:
Position_1: S(0.12), L(0.11), A(0.07)
Position_2: L(0.22), P(0.13), S(0.09)
Position_3: S(0.12), L(0.10), A(0.08)
Position_4: S(0.13), E(0.13), P(0.13)
Position_5: S(0.10), L(0.10), R(0.07)
Position_6: S(0.12), L(0.12), P(0.09)
Position_7: P(0.13), S(0.13), L(0.09)
Position_8: S(0.12), P(0.12), L(0.09)
Position_9: L(0.31), V(0.16), I(0.10)

Cluster 1: 732 peptides
Average binding score: 0.53
Average hydrophobicity: 0.75
Top amino acids at each position:
Position_1: A(0.20), G(0.14), S(0.14)
Position_2: L(0.31), A(0.13), P(0.13)
Position_3: A(0.16), G(0.11), S(0.11)
Position_4: G(0.17), P(0.16), A(0.15)
Position_5: G(0.19), A(0.16), V(0.11)
Position_6: G(0.16), S(0.14), L(0.12)
Position_7: A(0.13), P(0.12), S(0.11)
Position_8: A(0.18), S(0.16), G(0.12)
Position_9: L(0.29), V(0.25), A(0.14)

Cluster 2: 1419 peptides
Average binding score: 0.56
Average hydrophobicity: 0.57
Top amino acids at each position:
Position_1: L(0.10), A(0.10), S(0.09)
Position_2: L(0.32), V(0.10), T(0.08)
Position_3: L(0.14), A(0.09), S(0.09)
Position_4: S(0.09), G(0.09), L(0.09)
Position_5: L(0.11), G(0.11), A(0.09)
Position_6: L(0.14), V(0.09), S(0.09)
Position_7: L(0.14), V(0.11), A(0.09)
Position_8: L(0.11), A(0.10), S(0.10)
Position_9: L(0.28), V(0.20), K(0.12)

Cluster 3: 863 peptides
Average binding score: 0.57
Average hydrophobicity: -0.31
Top amino acids at each position:


```

Position_1: R(0.14), F(0.12), Y(0.12)
Position_2: L(0.19), Y(0.15), R(0.10)
Position_3: Y(0.11), F(0.11), L(0.10)
Position_4: E(0.11), R(0.10), L(0.08)
Position_5: R(0.14), F(0.11), L(0.09)
Position_6: L(0.12), F(0.10), R(0.08)
Position_7: L(0.12), F(0.09), R(0.08)
Position_8: L(0.10), E(0.09), R(0.09)
Position_9: L(0.26), F(0.18), Y(0.12)

```

Cluster 4: 1454 peptides

Average binding score: 0.58

Average hydrophobicity: 0.19

Top amino acids at each position:

```

Position_1: F(0.12), K(0.11), R(0.10)
Position_2: L(0.27), V(0.09), Y(0.08)
Position_3: L(0.13), F(0.08), D(0.07)
Position_4: E(0.11), D(0.08), L(0.08)
Position_5: L(0.11), F(0.08), V(0.08)
Position_6: L(0.15), V(0.09), I(0.09)
Position_7: L(0.14), F(0.09), V(0.06)
Position_8: L(0.13), S(0.08), F(0.08)
Position_9: L(0.28), V(0.15), F(0.11)

```

0.5.3 New Model

0.6 Convolutional Neural Network (CNN) Implementation

Below is a basic implementation of a Convolutional Neural Network using Keras (TensorFlow backend) for image classification. Adjust the `input_shape` and the number of output units in the final `Dense` layer according to your specific dataset.

0.7 Preparing Data for CNN

We need to: 1. Filter the epitopes and negatives dataframes to only contain the sequences and labels 2. One-hot encode the amino acid sequences 3. Split data into training and testing sets for the CNN model

```

[37]: import numpy as np
      from tensorflow.keras.preprocessing.sequence import pad_sequences
      from sklearn.model_selection import train_test_split

      # Step 1: Filter the epitopes and negatives dataframes to only contain
      ↪ sequences and labels
      epitopes_filtered = epitopes[['peptide', 'label']].copy()
      epitopes_filtered.rename(columns={'peptide': 'sequence'}, inplace=True)

```

```

negatives_filtered = negatives[['peptide', 'label']].copy()
negatives_filtered.rename(columns={'peptide': 'sequence'}, inplace=True)

# Combine the datasets
combined_data = pd.concat([epitopes_filtered, negatives_filtered],
    ↳ ignore_index=True)
combined_data = combined_data.sample(frac=1, random_state=42).
    ↳ reset_index(drop=True)

print(f"Number of samples: {combined_data.shape[0]}")
print(f"Positive samples: {sum(combined_data['label'] == 1)}")
print(f"Negative samples: {sum(combined_data['label'] == 0)}")

# Step 2: Prepare for one-hot encoding
# First, get all unique amino acids in our dataset
all_sequences = combined_data['sequence'].values
unique_chars = sorted(set(''.join(all_sequences)))
print(f"Unique amino acids in dataset: {unique_chars}")

# Create mapping dictionaries for one-hot encoding
char_to_index = {char: i+1 for i, char in enumerate(unique_chars)} # Start
    ↳ from 1, reserve 0 for padding
index_to_char = {i+1: char for i, char in enumerate(unique_chars)}
index_to_char[0] = '' # Padding token

# Find maximum sequence length
max_length = max(len(seq) for seq in all_sequences)
print(f"Maximum sequence length: {max_length}")

# Convert sequences to integer sequences
int_sequences = []
for seq in all_sequences:
    int_seq = [char_to_index[char] for char in seq]
    int_sequences.append(int_seq)

# Pad sequences to have the same length
padded_sequences = pad_sequences(int_sequences, maxlen=max_length,
    ↳ padding='post')

# One-hot encode the padded sequences
num_chars = len(unique_chars) + 1 # +1 for padding token
X_onehot = np.zeros((len(padded_sequences), max_length, num_chars))

for i, seq in enumerate(padded_sequences):
    for j, char_idx in enumerate(seq):
        X_onehot[i, j, char_idx] = 1.0 # One-hot encode

```

```

# Get labels
y = combined_data['label'].values

# Print shapes to verify dimensions
print(f"X_onehot shape: {X_onehot.shape}")
print(f"Number of unique amino acids (including padding): {num_chars}")

# Step 3: Split data into training, validation, and testing sets (70/15/15
↳split)
# First split into temporary train and test
X_temp, X_test, y_temp, y_test = train_test_split(
    X_onehot, y, test_size=0.15, random_state=42, stratify=y
)

# Then split the temporary train into final train and validation
# To get 70/15 split from the original data, we need to calculate the right
↳proportion:
# If test is 15% of total, then validation should be 15/85 of the remaining
↳data (approx 17.65%)
X_train, X_val, y_train, y_val = train_test_split(
    X_temp, y_temp, test_size=0.1765, random_state=42, stratify=y_temp
)

print(f"Training set shape: {X_train.shape} ({X_train.shape[0]/X_onehot.
↳shape[0]:.1%} of total)")
print(f"Validation set shape: {X_val.shape} ({X_val.shape[0]/X_onehot.shape[0]:.
↳1%} of total)")
print(f"Testing set shape: {X_test.shape} ({X_test.shape[0]/X_onehot.shape[0]:.
↳1%} of total)")

print(f"Positive samples in training: {sum(y_train == 1)} ({sum(y_train == 1)/
↳len(y_train):.1%})")
print(f"Negative samples in training: {sum(y_train == 0)} ({sum(y_train == 0)/
↳len(y_train):.1%})")

print(f"Positive samples in validation: {sum(y_val == 1)} ({sum(y_val == 1)/
↳len(y_val):.1%})")
print(f"Negative samples in validation: {sum(y_val == 0)} ({sum(y_val == 0)/
↳len(y_val):.1%})")

print(f"Positive samples in testing: {sum(y_test == 1)} ({sum(y_test == 1)/
↳len(y_test):.1%})")
print(f"Negative samples in testing: {sum(y_test == 0)} ({sum(y_test == 0)/
↳len(y_test):.1%})")

```

```

Number of samples: 25628
Positive samples: 5295
Negative samples: 20333
Unique amino acids in dataset: ['A', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'K',
'L', 'M', 'N', 'P', 'Q', 'R', 'S', 'T', 'V', 'W', 'Y']
Maximum sequence length: 9
X_onehot shape: (25628, 9, 21)
Number of unique amino acids (including padding): 21
Training set shape: (17938, 9, 21) (70.0% of total)
Validation set shape: (3845, 9, 21) (15.0% of total)
Testing set shape: (3845, 9, 21) (15.0% of total)
Positive samples in training: 3707 (20.7%)
Negative samples in training: 14231 (79.3%)
Positive samples in validation: 794 (20.7%)
Negative samples in validation: 3051 (79.3%)
Positive samples in testing: 794 (20.7%)
Negative samples in testing: 3051 (79.3%)

```

0.8 Convolutional Neural Network (CNN) Implementation

Below is a basic implementation of a Convolutional Neural Network using Keras (TensorFlow backend) for sequence classification.

```

[38]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv1D, MaxPooling1D, Flatten, Dense,
↳Dropout
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
import matplotlib.pyplot as plt

# Define the CNN model for sequence data
def create_cnn_model(input_shape, num_classes=2):
    model = Sequential([
        # First check the input shape to make sure it's valid
        # Input layer is implicit in the first Conv1D layer with input_shape
↳parameter

        # 1D Convolutional layers for sequence data
        Conv1D(filters=32, kernel_size=2, activation='relu', padding='same',
↳input_shape=input_shape),
        MaxPooling1D(pool_size=2),

        Conv1D(filters=64, kernel_size=2, activation='relu', padding='same'),
        MaxPooling1D(pool_size=2),

        # Flatten the feature maps
        Flatten(),

```

```

        # Dense layers
        Dense(64, activation='relu'),
        Dropout(0.5), # Dropout for regularization

        # Output layer
        Dense(num_classes, activation='softmax')
    ])

    # Compile the model
    model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

    return model

# Print shape information
print(f"Input data shape: {X_train.shape}")
print(f"Number of samples: {X_train.shape[0]}")
print(f"Sequence length: {X_train.shape[1]}")
print(f"Number of features per position: {X_train.shape[2]}")

# Use our preprocessed data from previous cell
# Note: X_train.shape = (n_samples, max_seq_length, n_amino_acids)
input_shape = (X_train.shape[1], X_train.shape[2])

# Create the model
cnn_model = create_cnn_model(input_shape)
print(cnn_model.summary())

# Define callbacks for training
early_stopping = EarlyStopping(monitor='val_loss', patience=5,
    ↪ restore_best_weights=True)
model_checkpoint = ModelCheckpoint('best_cnn_model.h5', monitor='val_accuracy',
    ↪ save_best_only=True)

# Train the model using the validation set instead of splitting the training
    ↪ data
history = cnn_model.fit(
    X_train, y_train,
    epochs=10,
    batch_size=32,
    validation_data=(X_val, y_val), # Use our validation set directly
    callbacks=[early_stopping, model_checkpoint]
)

# Evaluate the model on test data
test_loss, test_accuracy = cnn_model.evaluate(X_test, y_test)

```

```

print(f"Test accuracy: {test_accuracy:.4f}")

# Make predictions
y_pred_proba = cnn_model.predict(X_test)
y_pred = np.argmax(y_pred_proba, axis=1)

# Calculate metrics
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix, roc_curve, auc
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# Plot confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
tick_marks = np.arange(2)
plt.xticks(tick_marks, ['Negative', 'Positive'])
plt.yticks(tick_marks, ['Negative', 'Positive'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')

# Add text annotations to the confusion matrix
thresh = cm.max() / 2
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")
plt.tight_layout()
plt.show()

# Plot training history
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')

```

```

plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.tight_layout()
plt.show()

# Plot ROC curve
y_pred_proba_positive = y_pred_proba[:, 1] # Probability for positive class
fpr, tpr, _ = roc_curve(y_test, y_pred_proba_positive)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:
↵.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.grid(True, alpha=0.3)
plt.show()

```

Input data shape: (17938, 9, 21)
 Number of samples: 17938
 Sequence length: 9
 Number of features per position: 21

/Users/tariq/Documents/capstone/.venv/lib/python3.12/site-
 packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not
 pass an `input_shape`/`input_dim` argument to a layer. When using Sequential
 models, prefer using an `Input(shape)` object as the first layer in the model
 instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv1d (Conv1D)	(None, 9, 32)	1,376
max_pooling1d (MaxPooling1D)	(None, 4, 32)	0
conv1d_1 (Conv1D)	(None, 4, 64)	4,160

max_pooling1d_1 (MaxPooling1D)	(None, 2, 64)	0
flatten (Flatten)	(None, 128)	0
dense (Dense)	(None, 64)	8,256
dropout (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 2)	130

Total params: 13,922 (54.38 KB)

Trainable params: 13,922 (54.38 KB)

Non-trainable params: 0 (0.00 B)

None

Epoch 1/10

550/561 0s 640us/step -

accuracy: 0.7814 - loss: 0.5014

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g.

`model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

561/561 1s 869us/step -

accuracy: 0.7816 - loss: 0.5007 - val_accuracy: 0.7958 - val_loss: 0.4114

Epoch 2/10

509/561 0s 694us/step -

accuracy: 0.8012 - loss: 0.4055

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g.

`model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

561/561 0s 793us/step -

accuracy: 0.8017 - loss: 0.4047 - val_accuracy: 0.8130 - val_loss: 0.3753

Epoch 3/10

549/561 0s 642us/step -

accuracy: 0.8228 - loss: 0.3685

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We

recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561 0s 746us/step -
accuracy: 0.8227 - loss: 0.3685 - val_accuracy: 0.8182 - val_loss: 0.3632
Epoch 4/10
530/561 0s 667us/step -
accuracy: 0.8264 - loss: 0.3597

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561 0s 770us/step -
accuracy: 0.8266 - loss: 0.3592 - val_accuracy: 0.8260 - val_loss: 0.3581
Epoch 5/10
534/561 0s 660us/step -
accuracy: 0.8359 - loss: 0.3449

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561 0s 766us/step -
accuracy: 0.8360 - loss: 0.3447 - val_accuracy: 0.8278 - val_loss: 0.3577
Epoch 6/10
561/561 0s 743us/step -
accuracy: 0.8439 - loss: 0.3346 - val_accuracy: 0.8255 - val_loss: 0.3594
Epoch 7/10
561/561 0s 722us/step -
accuracy: 0.8480 - loss: 0.3174 - val_accuracy: 0.8086 - val_loss: 0.3773
Epoch 8/10
561/561 0s 773us/step -
accuracy: 0.8546 - loss: 0.3132 - val_accuracy: 0.8273 - val_loss: 0.3818
Epoch 9/10
550/561 0s 642us/step -
accuracy: 0.8629 - loss: 0.3010


WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

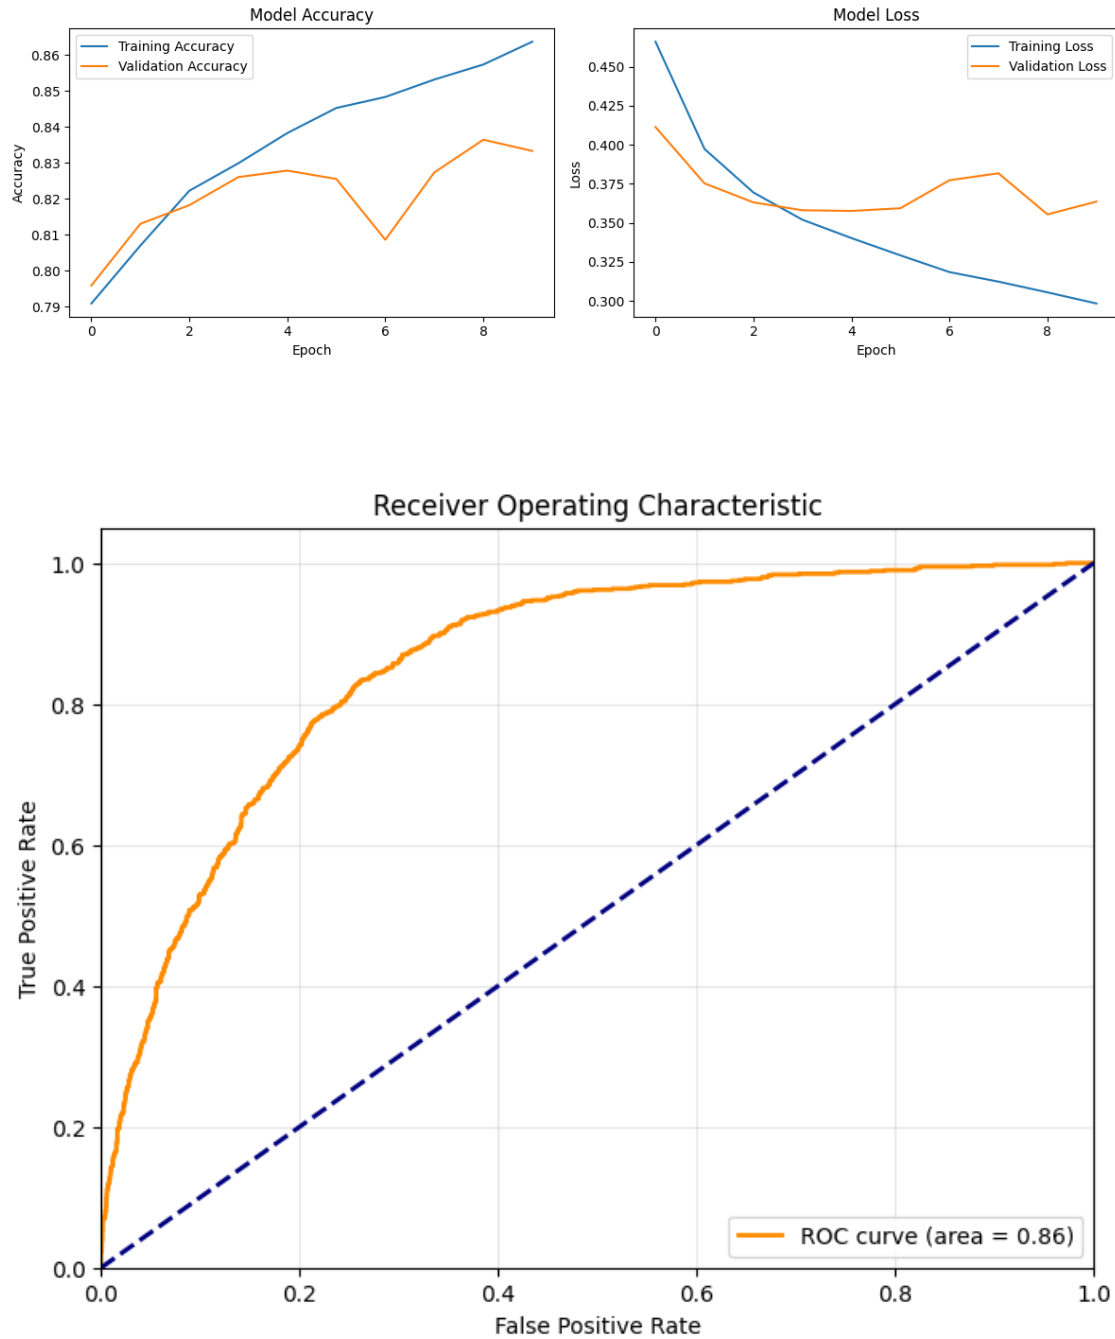
561/561 0s 749us/step -
accuracy: 0.8628 - loss: 0.3011 - val_accuracy: 0.8364 - val_loss: 0.3555

Epoch 10/10
 561/561 0s 755us/step -
 accuracy: 0.8655 - loss: 0.2978 - val_accuracy: 0.8333 - val_loss: 0.3637
 121/121 0s 348us/step -
 accuracy: 0.8264 - loss: 0.3824
 Test accuracy: 0.8276
 121/121 0s 422us/step

Classification Report:

	precision	recall	f1-score	support
0	0.87	0.91	0.89	3051
1	0.60	0.50	0.54	794
accuracy			0.83	3845
macro avg	0.74	0.71	0.72	3845
weighted avg	0.82	0.83	0.82	3845





0.9 Optimized CNN Model

Below is an optimized CNN model implementation with several improvements to address the class imbalance:

1. Class weights to handle imbalanced data
2. Batch normalization for better training stability

3. Focal loss implementation for imbalanced classification
4. Learning rate scheduling
5. Increased model capacity with additional layers
6. Optimized threshold selection for classification

```
[39]: import tensorflow as tf
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.layers import Conv1D, MaxPooling1D, Flatten, Dense,
↳ Dropout, Input, BatchNormalization
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint,
↳ ReduceLROnPlateau
from tensorflow.keras.regularizers import l2
from tensorflow.keras.optimizers import Adam
from sklearn.metrics import precision_recall_curve, f1_score
import numpy as np
import matplotlib.pyplot as plt

# Define the focal loss function to better handle class imbalance
def focal_loss(gamma=2.0, alpha=0.25):
    def focal_loss_fn(y_true, y_pred):
        # Convert one-hot encoded targets to integers
        if y_true.shape[-1] == 1:
            y_true = tf.squeeze(y_true, axis=-1)
            y_true = tf.cast(y_true, tf.int32)

        # Get the standard sparse categorical crossentropy
        sce = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False,
↳ reduction=tf.keras.losses.Reduction.NONE)(y_true, y_pred)

        # Calculate the prediction probability for the true class
        y_pred_proba = tf.gather_nd(y_pred, tf.stack([tf.range(tf.
↳ shape(y_true)[0]), tf.cast(y_true, tf.int32)], axis=1))

        # Apply focal loss formula
        # p_t = p if y == 1 else 1-p for class 0
        p_t = y_pred_proba
        # Add the alpha weighing factor
        alpha_factor = 1.0
        if alpha is not None:
            # alpha_t = alpha if y == 1 else 1-alpha for class 0
            alpha_t = tf.where(tf.equal(y_true, 1), alpha, 1-alpha)
            alpha_factor = alpha_t

        # Calculate focal weight
        gamma_factor = tf.pow(1.0 - p_t, gamma)

        # Calculate the final loss
```

```

        focal_loss = alpha_factor * gamma_factor * sce

        return tf.reduce_mean(focal_loss)

    return focal_loss_fn

# Create an optimized CNN model for sequence data
def create_optimized_cnn_model(input_shape, num_classes=2, use_focal_loss=True):
    inputs = Input(shape=input_shape)

    # First convolutional block
    x = Conv1D(32, kernel_size=3, activation='relu', padding='same',
    ↪kernel_regularizer=l2(0.001))(inputs)
    x = BatchNormalization()(x)
    x = MaxPooling1D(pool_size=2, padding='same')(x)

    # Second convolutional block with increased filters
    x = Conv1D(64, kernel_size=3, activation='relu', padding='same',
    ↪kernel_regularizer=l2(0.001))(x)
    x = BatchNormalization()(x)
    x = MaxPooling1D(pool_size=2, padding='same')(x)

    # Third convolutional block with even more filters
    x = Conv1D(128, kernel_size=3, activation='relu', padding='same',
    ↪kernel_regularizer=l2(0.001))(x)
    x = BatchNormalization()(x)

    # Flatten and dense layers
    x = Flatten()(x)

    # Add more capacity to the dense layers
    x = Dense(128, activation='relu', kernel_regularizer=l2(0.001))(x)
    x = BatchNormalization()(x)
    x = Dropout(0.4)(x)

    x = Dense(64, activation='relu', kernel_regularizer=l2(0.001))(x)
    x = BatchNormalization()(x)
    x = Dropout(0.3)(x)

    # Output layer
    outputs = Dense(num_classes, activation='softmax')(x)

    model = Model(inputs=inputs, outputs=outputs)

    # Use a lower learning rate for better stability
    optimizer = Adam(learning_rate=0.001)

```

```

    # Use focal loss if requested, otherwise use standard cross-entropy
    if use_focal_loss:
        loss = focal_loss(gamma=2.0, alpha=0.75) # Adjust alpha based on class
        ↪ imbalance
    else:
        loss = 'sparse_categorical_crossentropy'

    model.compile(
        optimizer=optimizer,
        loss=loss,
        metrics=['accuracy']
    )

    return model

# Calculate class weights based on class frequencies
# This gives more weight to the minority class during training
def compute_class_weights(y_train):
    # Count the number of samples per class
    class_counts = np.bincount(y_train)
    # Calculate the weight for each class (inversely proportional to class
    ↪ frequency)
    total_samples = len(y_train)
    class_weights = {
        i: total_samples / (len(class_counts) * count)
        for i, count in enumerate(class_counts)
    }
    return class_weights

# Get the class weights for our training data
class_weights = compute_class_weights(y_train)
print(f"Class weights: {class_weights}")

# Create an optimized CNN model
optimized_cnn_model = create_optimized_cnn_model(input_shape=(X_train.shape[1],
    ↪ X_train.shape[2]), use_focal_loss=False)
print(optimized_cnn_model.summary())

# Define more sophisticated callbacks
early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=8,
    restore_best_weights=True,
    verbose=1
)

reduce_lr = ReduceLROnPlateau(

```

```

        monitor='val_loss',
        factor=0.2,
        patience=3,
        min_lr=0.00001,
        verbose=1
    )

model_checkpoint = ModelCheckpoint(
    'best_optimized_cnn_model.h5',
    monitor='val_accuracy',
    save_best_only=True,
    verbose=1
)

# Train the model with class weights
history = optimized_cnn_model.fit(
    X_train, y_train,
    epochs=20, # Increase epochs since we have early stopping
    batch_size=32,
    validation_data=(X_val, y_val),
    callbacks=[early_stopping, reduce_lr, model_checkpoint],
    class_weight=class_weights # Use class weights during training
)

# Evaluate the model on test data
test_loss, test_accuracy = optimized_cnn_model.evaluate(X_test, y_test)
print(f"Test accuracy: {test_accuracy:.4f}")

# Make predictions on test data
y_pred_proba = optimized_cnn_model.predict(X_test)
y_pred_proba_positive = y_pred_proba[:, 1] # Probability for positive class

# Find the optimal threshold for F1 score
thresholds = np.arange(0.1, 0.9, 0.05)
f1_scores = []

for threshold in thresholds:
    y_pred_thresholded = (y_pred_proba_positive >= threshold).astype(int)
    f1 = f1_score(y_test, y_pred_thresholded)
    f1_scores.append(f1)
    print(f"Threshold: {threshold:.2f}, F1 Score: {f1:.4f}")

# Get the best threshold
best_threshold_idx = np.argmax(f1_scores)
best_threshold = thresholds[best_threshold_idx]
best_f1 = f1_scores[best_threshold_idx]
print(f"\nOptimal threshold: {best_threshold:.2f} with F1 Score: {best_f1:.4f}")

```

```

# Apply the best threshold
y_pred = (y_pred_proba_positive >= best_threshold).astype(int)

# Print classification report with the optimized threshold
from sklearn.metrics import classification_report, confusion_matrix
print("\nClassification Report with Optimized Threshold:")
print(classification_report(y_test, y_pred))

# Plot confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion Matrix (Optimized Threshold)')
plt.colorbar()
tick_marks = np.arange(2)
plt.xticks(tick_marks, ['Negative', 'Positive'])
plt.yticks(tick_marks, ['Negative', 'Positive'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')

# Add text annotations to the confusion matrix
thresh = cm.max() / 2
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")
plt.tight_layout()
plt.show()

# Plot training history
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')

```



```

plt.legend()
plt.tight_layout()
plt.show()

# Plot ROC curve
from sklearn.metrics import roc_curve, auc
fpr, tpr, _ = roc_curve(y_test, y_pred_proba_positive)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:
    ↪.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.scatter(fpr[np.argmin(np.abs(thresholds - best_threshold))],
            tpr[np.argmin(np.abs(thresholds - best_threshold))],
            c='red', marker='o', s=100, label=f'Best threshold =
    ↪{best_threshold:.2f}')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.grid(True, alpha=0.3)
plt.show()

# Plot Precision-Recall curve
from sklearn.metrics import precision_recall_curve, average_precision_score
precision, recall, thresholds_pr = precision_recall_curve(y_test,
    ↪y_pred_proba_positive)
avg_precision = average_precision_score(y_test, y_pred_proba_positive)

plt.figure(figsize=(8, 6))
plt.plot(recall, precision, color='blue', lw=2, label=f'PR curve (AP =
    ↪{avg_precision:.2f})')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend(loc="upper right")
plt.grid(True, alpha=0.3)
plt.show()

# Compare with the best threshold ROC point
plt.figure(figsize=(8, 6))
plt.step(recall, precision, color='blue', alpha=0.2, where='post')
plt.fill_between(recall, precision, alpha=0.2, color='blue', step='post')
plt.xlabel('Recall')

```

```

plt.ylabel('Precision')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.title('Precision-Recall Curve: AP={0:0.2f}'.format(avg_precision))
plt.show()

# Plot F1 Score vs Threshold
plt.figure(figsize=(8, 6))
plt.plot(thresholds, f1_scores, 'b-', label='F1 Score')
plt.plot([best_threshold, best_threshold], [0, best_f1], 'r--', label=f'Best_
↪Threshold = {best_threshold:.2f}')
plt.plot(best_threshold, best_f1, 'ro', markersize=8)
plt.title('F1 Score vs. Threshold')
plt.xlabel('Threshold')
plt.ylabel('F1 Score')
plt.grid(True, alpha=0.3)
plt.legend()
plt.show()

# Comparing model performance before and after optimization
print(f"\nModel performance comparison:")
print(f"Optimal threshold: {best_threshold:.2f}")
print(f"Original model test accuracy: {test_accuracy:.4f}")
print(f"Optimized model test accuracy (with best threshold):_
↪{accuracy_score(y_test, y_pred):.4f}")
print(f"Optimized model F1 score: {f1_score(y_test, y_pred):.4f}")

```

Class weights: {0: np.float64(0.6302438338837748), 1:
np.float64(2.419476665767467)}

Model: "functional_1"

Layer (type)	Output Shape	Param #
input_layer_1 (InputLayer)	(None , 9, 21)	0
conv1d_2 (Conv1D)	(None , 9, 32)	2,048
batch_normalization (BatchNormalization)	(None , 9, 32)	128
max_pooling1d_2 (MaxPooling1D)	(None , 5, 32)	0
conv1d_3 (Conv1D)	(None , 5, 64)	6,208
batch_normalization_1	(None , 5, 64)	256

(BatchNormalization)		
max_pooling1d_3 (MaxPooling1D)	(None, 3, 64)	0
conv1d_4 (Conv1D)	(None, 3, 128)	24,704
batch_normalization_2 (BatchNormalization)	(None, 3, 128)	512
flatten_1 (Flatten)	(None, 384)	0
dense_2 (Dense)	(None, 128)	49,280
batch_normalization_3 (BatchNormalization)	(None, 128)	512
dropout_1 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 64)	8,256
batch_normalization_4 (BatchNormalization)	(None, 64)	256
dropout_2 (Dropout)	(None, 64)	0
dense_4 (Dense)	(None, 2)	130

Total params: 92,290 (360.51 KB)

Trainable params: 91,458 (357.26 KB)

Non-trainable params: 832 (3.25 KB)

None

Epoch 1/20

542/561 0s 2ms/step -

accuracy: 0.5844 - loss: 1.2407

Epoch 1: val_accuracy improved from -inf to 0.70039, saving model to best_optimized_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g.

`model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

```

561/561          2s 2ms/step -
accuracy: 0.5860 - loss: 1.2361 - val_accuracy: 0.7004 - val_loss: 0.9537 -
learning_rate: 0.0010
Epoch 2/20
559/561          0s 1ms/step -
accuracy: 0.7205 - loss: 0.8932
Epoch 2: val_accuracy improved from 0.70039 to 0.73186, saving model to
best_optimized_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          1s 2ms/step -
accuracy: 0.7205 - loss: 0.8931 - val_accuracy: 0.7319 - val_loss: 0.8500 -
learning_rate: 0.0010
Epoch 3/20
555/561          0s 1ms/step -
accuracy: 0.7616 - loss: 0.7739
Epoch 3: val_accuracy improved from 0.73186 to 0.73238, saving model to
best_optimized_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          1s 2ms/step -
accuracy: 0.7616 - loss: 0.7738 - val_accuracy: 0.7324 - val_loss: 0.7705 -
learning_rate: 0.0010
Epoch 4/20
559/561          0s 2ms/step -
accuracy: 0.7741 - loss: 0.6849
Epoch 4: val_accuracy improved from 0.73238 to 0.75865, saving model to
best_optimized_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          1s 2ms/step -
accuracy: 0.7741 - loss: 0.6849 - val_accuracy: 0.7586 - val_loss: 0.6861 -
learning_rate: 0.0010
Epoch 5/20
549/561          0s 2ms/step -
accuracy: 0.7958 - loss: 0.6119

```

```

Epoch 5: val_accuracy did not improve from 0.75865
561/561          1s 2ms/step -
accuracy: 0.7956 - loss: 0.6120 - val_accuracy: 0.7352 - val_loss: 0.6669 -
learning_rate: 0.0010
Epoch 6/20
553/561          0s 2ms/step -
accuracy: 0.7959 - loss: 0.5564
Epoch 6: val_accuracy did not improve from 0.75865
561/561          1s 2ms/step -
accuracy: 0.7959 - loss: 0.5565 - val_accuracy: 0.7222 - val_loss: 0.6765 -
learning_rate: 0.0010
Epoch 7/20
547/561          0s 2ms/step -
accuracy: 0.8063 - loss: 0.5248
Epoch 7: val_accuracy improved from 0.75865 to 0.77893, saving model to
best_optimized_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          1s 2ms/step -
accuracy: 0.8062 - loss: 0.5250 - val_accuracy: 0.7789 - val_loss: 0.6100 -
learning_rate: 0.0010
Epoch 8/20
545/561          0s 2ms/step -
accuracy: 0.8165 - loss: 0.4866
Epoch 8: val_accuracy did not improve from 0.77893
561/561          1s 2ms/step -
accuracy: 0.8162 - loss: 0.4870 - val_accuracy: 0.7246 - val_loss: 0.6455 -
learning_rate: 0.0010
Epoch 9/20
527/561          0s 2ms/step -
accuracy: 0.8224 - loss: 0.4703
Epoch 9: val_accuracy improved from 0.77893 to 0.78362, saving model to
best_optimized_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          1s 2ms/step -
accuracy: 0.8221 - loss: 0.4708 - val_accuracy: 0.7836 - val_loss: 0.5473 -
learning_rate: 0.0010
Epoch 10/20
543/561          0s 2ms/step -

```

```

accuracy: 0.8300 - loss: 0.4504
Epoch 10: val_accuracy did not improve from 0.78362
561/561          1s 2ms/step -
accuracy: 0.8296 - loss: 0.4509 - val_accuracy: 0.7095 - val_loss: 0.6811 -
learning_rate: 0.0010
Epoch 11/20
554/561          0s 2ms/step -
accuracy: 0.8365 - loss: 0.4394
Epoch 11: val_accuracy did not improve from 0.78362
561/561          1s 2ms/step -
accuracy: 0.8364 - loss: 0.4396 - val_accuracy: 0.7441 - val_loss: 0.6399 -
learning_rate: 0.0010
Epoch 12/20
538/561          0s 2ms/step -
accuracy: 0.8503 - loss: 0.4256
Epoch 12: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.

Epoch 12: val_accuracy did not improve from 0.78362
561/561          1s 2ms/step -
accuracy: 0.8496 - loss: 0.4265 - val_accuracy: 0.7740 - val_loss: 0.5956 -
learning_rate: 0.0010
Epoch 13/20
555/561          0s 2ms/step -
accuracy: 0.8641 - loss: 0.3816
Epoch 13: val_accuracy improved from 0.78362 to 0.78882, saving model to
best_optimized_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          1s 2ms/step -
accuracy: 0.8641 - loss: 0.3815 - val_accuracy: 0.7888 - val_loss: 0.5986 -
learning_rate: 2.0000e-04
Epoch 14/20
550/561          0s 2ms/step -
accuracy: 0.8972 - loss: 0.3301
Epoch 14: val_accuracy improved from 0.78882 to 0.79272, saving model to
best_optimized_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          1s 2ms/step -
accuracy: 0.8971 - loss: 0.3302 - val_accuracy: 0.7927 - val_loss: 0.6206 -

```

```

learning_rate: 2.0000e-04
Epoch 15/20
532/561          0s 2ms/step -
accuracy: 0.9100 - loss: 0.3024
Epoch 15: ReduceLROnPlateau reducing learning rate to 4.0000001899898055e-05.

Epoch 15: val_accuracy improved from 0.79272 to 0.79324, saving model to
best_optimized_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          1s 2ms/step -
accuracy: 0.9098 - loss: 0.3029 - val_accuracy: 0.7932 - val_loss: 0.6504 -
learning_rate: 2.0000e-04
Epoch 16/20
536/561          0s 2ms/step -
accuracy: 0.9251 - loss: 0.2733
Epoch 16: val_accuracy improved from 0.79324 to 0.80000, saving model to
best_optimized_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          1s 2ms/step -
accuracy: 0.9252 - loss: 0.2732 - val_accuracy: 0.8000 - val_loss: 0.6834 -
learning_rate: 4.0000e-05
Epoch 17/20
560/561          0s 2ms/step -
accuracy: 0.9300 - loss: 0.2677
Epoch 17: val_accuracy improved from 0.80000 to 0.80026, saving model to
best_optimized_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          1s 2ms/step -
accuracy: 0.9300 - loss: 0.2677 - val_accuracy: 0.8003 - val_loss: 0.6979 -
learning_rate: 4.0000e-05
Epoch 17: early stopping
Restoring model weights from the end of the best epoch: 9.
121/121          0s 520us/step -



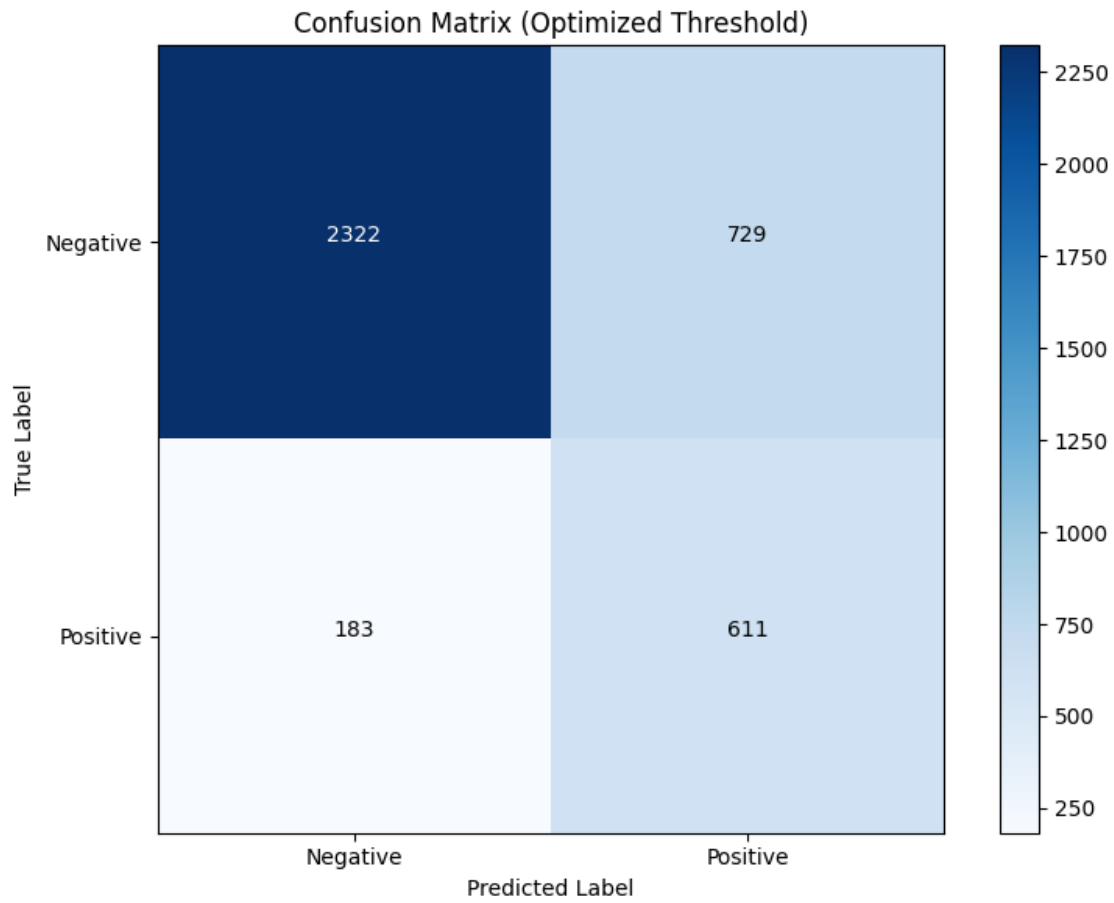
```


accuracy: 0.7744 - loss: 0.5690
 Test accuracy: 0.7724
 121/121 0s 843us/step
 Threshold: 0.10, F1 Score: 0.5334
 Threshold: 0.15, F1 Score: 0.5414
 Threshold: 0.20, F1 Score: 0.5507
 Threshold: 0.25, F1 Score: 0.5614
 Threshold: 0.30, F1 Score: 0.5654
 Threshold: 0.35, F1 Score: 0.5661
 Threshold: 0.40, F1 Score: 0.5691
 Threshold: 0.45, F1 Score: 0.5726
 Threshold: 0.50, F1 Score: 0.5721
 Threshold: 0.55, F1 Score: 0.5711
 Threshold: 0.60, F1 Score: 0.5694
 Threshold: 0.65, F1 Score: 0.5483
 Threshold: 0.70, F1 Score: 0.5436
 Threshold: 0.75, F1 Score: 0.5180
 Threshold: 0.80, F1 Score: 0.4782
 Threshold: 0.85, F1 Score: 0.3942


Optimal threshold: 0.45 with F1 Score: 0.5726


Classification Report with Optimized Threshold:


	precision	recall	f1-score	support
0	0.93	0.76	0.84	3051
1	0.46	0.77	0.57	794
accuracy			0.76	3845
macro avg	0.69	0.77	0.70	3845
weighted avg	0.83	0.76	0.78	3845











Model performance comparison:

Optimal threshold: 0.45

Original model test accuracy: 0.7724

Optimized model test accuracy (with best threshold): 0.7628

Optimized model F1 score: 0.5726

0.10 Ensemble Model with Focal Loss

Let's implement an ensemble approach combining the optimized CNN with focal loss to further improve the model's performance on the imbalanced dataset.

```
[40]: import tensorflow as tf
from tensorflow.keras.models import Sequential, Model, load_model
from tensorflow.keras.layers import Conv1D, MaxPooling1D, Flatten, Dense,
↳ Dropout, Input, BatchNormalization, Concatenate, Average
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint,
↳ ReduceLROnPlateau
from sklearn.metrics import classification_report, confusion_matrix, roc_curve,
↳ auc, precision_recall_curve
```

```

import matplotlib.pyplot as plt
import numpy as np

# Create a CNN model with focal loss for the ensemble
focal_loss_model = create_optimized_cnn_model(
    input_shape=(X_train.shape[1], X_train.shape[2]),
    use_focal_loss=True
)

# Define callbacks for training
early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=8,
    restore_best_weights=True,
    verbose=1
)

reduce_lr = ReduceLROnPlateau(
    monitor='val_loss',
    factor=0.2,
    patience=3,
    min_lr=0.00001,
    verbose=1
)

model_checkpoint = ModelCheckpoint(
    'best_focal_loss_model.h5',
    monitor='val_accuracy',
    save_best_only=True,
    verbose=1
)

# Train the focal loss model (without class weights since focal loss already
↳ handles imbalance)
focal_loss_history = focal_loss_model.fit(
    X_train, y_train,
    epochs=20,
    batch_size=32,
    validation_data=(X_val, y_val),
    callbacks=[early_stopping, reduce_lr, model_checkpoint]
)

# Evaluate the focal loss model
test_loss_focal, test_accuracy_focal = focal_loss_model.evaluate(X_test, y_test)
print(f"Focal Loss Model - Test accuracy: {test_accuracy_focal:.4f}")

# Make predictions with both models

```

```

y_pred_proba_ce = optimized_cnn_model.predict(X_test)
y_pred_proba_focal = focal_loss_model.predict(X_test)

# Ensemble predictions by averaging
y_pred_proba_ensemble = (y_pred_proba_ce + y_pred_proba_focal) / 2.0
y_pred_proba_ensemble_positive = y_pred_proba_ensemble[:, 1]

# Find optimal threshold for the ensemble
thresholds = np.arange(0.1, 0.9, 0.05)
ensemble_f1_scores = []

for threshold in thresholds:
    y_pred_thresholded = (y_pred_proba_ensemble_positive >= threshold).
    ↳astype(int)
    f1 = f1_score(y_test, y_pred_thresholded)
    ensemble_f1_scores.append(f1)
    print(f"Ensemble - Threshold: {threshold:.2f}, F1 Score: {f1:.4f}")

# Get the best threshold for ensemble
best_ensemble_threshold_idx = np.argmax(ensemble_f1_scores)
best_ensemble_threshold = thresholds[best_ensemble_threshold_idx]
best_ensemble_f1 = ensemble_f1_scores[best_ensemble_threshold_idx]
print(f"\nEnsemble - Optimal threshold: {best_ensemble_threshold:.2f} with F1_
    ↳Score: {best_ensemble_f1:.4f}")

# Apply the best threshold to ensemble predictions
y_pred_ensemble = (y_pred_proba_ensemble_positive >= best_ensemble_threshold).
    ↳astype(int)

# Print classification report for ensemble
print("\nEnsemble Model - Classification Report with Optimized Threshold:")
print(classification_report(y_test, y_pred_ensemble))

# Plot confusion matrix for ensemble
cm_ensemble = confusion_matrix(y_test, y_pred_ensemble)
plt.figure(figsize=(8, 6))
plt.imshow(cm_ensemble, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Ensemble Model - Confusion Matrix')
plt.colorbar()
tick_marks = np.arange(2)
plt.xticks(tick_marks, ['Negative', 'Positive'])
plt.yticks(tick_marks, ['Negative', 'Positive'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')

# Add text annotations
thresh = cm_ensemble.max() / 2

```

```

for i in range(cm_ensemble.shape[0]):
    for j in range(cm_ensemble.shape[1]):
        plt.text(j, i, cm_ensemble[i, j],
                 horizontalalignment="center",
                 color="white" if cm_ensemble[i, j] > thresh else "black")
plt.tight_layout()
plt.show()

# Plot ROC curves for comparison
plt.figure(figsize=(10, 8))

# Original model ROC
fpr_orig, tpr_orig, _ = roc_curve(y_test, y_pred_proba_positive)
roc_auc_orig = auc(fpr_orig, tpr_orig)
plt.plot(fpr_orig, tpr_orig, color='blue', lw=2,
         label=f'Optimized CNN (AUC = {roc_auc_orig:.2f})')

# Focal loss model ROC
fpr_focal, tpr_focal, _ = roc_curve(y_test, y_pred_proba_focal[:, 1])
roc_auc_focal = auc(fpr_focal, tpr_focal)
plt.plot(fpr_focal, tpr_focal, color='green', lw=2,
         label=f'Focal Loss CNN (AUC = {roc_auc_focal:.2f})')

# Ensemble model ROC
fpr_ensemble, tpr_ensemble, _ = roc_curve(y_test,
    y_pred_proba_ensemble_positive)
roc_auc_ensemble = auc(fpr_ensemble, tpr_ensemble)
plt.plot(fpr_ensemble, tpr_ensemble, color='red', lw=2,
         label=f'Ensemble Model (AUC = {roc_auc_ensemble:.2f})')

# Add diagonal line
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Comparing ROC Curves')
plt.legend(loc="lower right")
plt.grid(True, alpha=0.3)
plt.show()

# Precision-Recall curves comparison
plt.figure(figsize=(10, 8))

# Calculate PR curves for all models

```



```

precision_orig, recall_orig, _ = precision_recall_curve(y_test,
    ↪ y_pred_proba_positive)
ap_orig = average_precision_score(y_test, y_pred_proba_positive)

precision_focal, recall_focal, _ = precision_recall_curve(y_test,
    ↪ y_pred_proba_focal[:, 1])
ap_focal = average_precision_score(y_test, y_pred_proba_focal[:, 1])

precision_ensemble, recall_ensemble, _ = precision_recall_curve(y_test,
    ↪ y_pred_proba_ensemble_positive)
ap_ensemble = average_precision_score(y_test, y_pred_proba_ensemble_positive)

# Plot all PR curves
plt.plot(recall_orig, precision_orig, color='blue', lw=2,
    label=f'Optimized CNN (AP = {ap_orig:.2f})')
plt.plot(recall_focal, precision_focal, color='green', lw=2,
    label=f'Focal Loss CNN (AP = {ap_focal:.2f})')
plt.plot(recall_ensemble, precision_ensemble, color='red', lw=2,
    label=f'Ensemble Model (AP = {ap_ensemble:.2f})')

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Comparing Precision-Recall Curves')
plt.legend(loc="upper right")
plt.grid(True, alpha=0.3)
plt.show()

# Summary of model performance
print("\n===== MODEL PERFORMANCE COMPARISON =====")
print(f"{'Model':<25} {'Accuracy':<10} {'F1 Score':<10} {'AUC':<10}")
print("-" * 55)
print(f"{'Original CNN':<25} {test_accuracy:.4f}      {f1_score(y_test, y_pred):.4f}      {roc_auc_orig:.4f}")
print(f"{'Focal Loss CNN':<25} {test_accuracy_focal:.4f}      {f1_score(y_test,
    ↪ (y_pred_proba_focal[:, 1] >= 0.5).astype(int)):.4f}      {roc_auc_focal:.4f}")
print(f"{'Ensemble Model':<25} {accuracy_score(y_test, y_pred_ensemble):.4f}      ↪
    ↪ {best_ensemble_f1:.4f}      {roc_auc_ensemble:.4f}")
print("=" * 55)

# Check predictions on specific examples where models disagree
disagreement_indices = np.where(
    ((y_pred_proba_positive >= best_threshold).astype(int) !=
    (y_pred_proba_focal[:, 1] >= 0.5).astype(int))
)[0]

```

```

if len(disagreement_indices) > 0:
    print(f"\nFound {len(disagreement_indices)} test examples where the models
    ↳disagree in their predictions.")
    sample_size = min(5, len(disagreement_indices))
    print(f"\nAnalyzing {sample_size} random examples of disagreement:")

    sample_indices = np.random.choice(disagreement_indices, sample_size,
    ↳replace=False)

    for i, idx in enumerate(sample_indices):
        true_label = y_test[idx]
        optimized_pred = (y_pred_proba_positive[idx] >= best_threshold).
        ↳astype(int)
        focal_pred = (y_pred_proba_focal[idx, 1] >= 0.5).astype(int)
        ensemble_pred = (y_pred_proba_ensemble_positive[idx] >=
        ↳best_ensemble_threshold).astype(int)

        print(f"\nExample {i+1}:")
        print(f"True label: {true_label}")
        print(f"Optimized CNN prediction: {optimized_pred} (confidence:
        ↳{y_pred_proba_positive[idx]:.4f})")
        print(f"Focal Loss CNN prediction: {focal_pred} (confidence:
        ↳{y_pred_proba_focal[idx, 1]:.4f})")
        print(f"Ensemble prediction: {ensemble_pred} (confidence:
        ↳{y_pred_proba_ensemble_positive[idx]:.4f})")
    else:
        print("\nNo disagreement found between the models on test examples.")

```

Epoch 1/20

548/561 0s 2ms/step -

accuracy: 0.5706 - loss: 0.5517

Epoch 1: val_accuracy improved from -inf to 0.70065, saving model to
best_focal_loss_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561 2s 2ms/step -

accuracy: 0.5715 - loss: 0.5498 - val_accuracy: 0.7007 - val_loss: 0.3567 -
learning_rate: 0.0010

Epoch 2/20

537/561 0s 2ms/step -

accuracy: 0.7120 - loss: 0.3237

Epoch 2: val_accuracy improved from 0.70065 to 0.71157, saving model to

best_focal_loss_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g.

`model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

561/561 1s 2ms/step -

accuracy: 0.7127 - loss: 0.3221 - val_accuracy: 0.7116 - val_loss: 0.2203 - learning_rate: 0.0010

Epoch 3/20

539/561 0s 2ms/step -

accuracy: 0.7494 - loss: 0.1962

Epoch 3: val_accuracy improved from 0.71157 to 0.75501, saving model to best_focal_loss_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g.

`model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

561/561 1s 2ms/step -

accuracy: 0.7496 - loss: 0.1953 - val_accuracy: 0.7550 - val_loss: 0.1345 - learning_rate: 0.0010

Epoch 4/20

544/561 0s 1ms/step -

accuracy: 0.7702 - loss: 0.1215

Epoch 4: val_accuracy improved from 0.75501 to 0.76983, saving model to best_focal_loss_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g.

`model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

561/561 1s 2ms/step -

accuracy: 0.7701 - loss: 0.1212 - val_accuracy: 0.7698 - val_loss: 0.0901 - learning_rate: 0.0010

Epoch 5/20

539/561 0s 2ms/step -

accuracy: 0.7812 - loss: 0.0826

Epoch 5: val_accuracy improved from 0.76983 to 0.80260, saving model to best_focal_loss_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g.

`model.save('my_model.keras')` or `keras.saving.save_model(model,

```
'my_model.keras')`.
```

```
561/561          1s 2ms/step -  
accuracy: 0.7812 - loss: 0.0824 - val_accuracy: 0.8026 - val_loss: 0.0663 -  
learning_rate: 0.0010
```

```
Epoch 6/20
```

```
534/561          0s 2ms/step -  
accuracy: 0.7845 - loss: 0.0639  
Epoch 6: val_accuracy did not improve from 0.80260
```

```
561/561          1s 2ms/step -  
accuracy: 0.7846 - loss: 0.0638 - val_accuracy: 0.7724 - val_loss: 0.0575 -  
learning_rate: 0.0010
```

```
Epoch 7/20
```

```
553/561          0s 2ms/step -  
accuracy: 0.7923 - loss: 0.0560  
Epoch 7: val_accuracy did not improve from 0.80260
```

```
561/561          1s 2ms/step -  
accuracy: 0.7922 - loss: 0.0560 - val_accuracy: 0.7867 - val_loss: 0.0536 -  
learning_rate: 0.0010
```

```
Epoch 8/20
```

```
555/561          0s 2ms/step -  
accuracy: 0.7957 - loss: 0.0510  
Epoch 8: val_accuracy did not improve from 0.80260
```

```
561/561          1s 2ms/step -  
accuracy: 0.7956 - loss: 0.0510 - val_accuracy: 0.7938 - val_loss: 0.0516 -  
learning_rate: 0.0010
```

```
Epoch 9/20
```

```
544/561          0s 2ms/step -  
accuracy: 0.7865 - loss: 0.0496  
Epoch 9: val_accuracy improved from 0.80260 to 0.80286, saving model to  
best_focal_loss_model.h5
```

```
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or  
`keras.saving.save_model(model)`. This file format is considered legacy. We  
recommend using instead the native Keras format, e.g.
```

```
`model.save('my_model.keras')` or `keras.saving.save_model(model,  
'my_model.keras')`.
```

```
561/561          1s 2ms/step -  
accuracy: 0.7864 - loss: 0.0496 - val_accuracy: 0.8029 - val_loss: 0.0500 -  
learning_rate: 0.0010
```

```
Epoch 10/20
```

```
541/561          0s 2ms/step -  
accuracy: 0.7872 - loss: 0.0494  
Epoch 10: val_accuracy did not improve from 0.80286
```

```
561/561          1s 2ms/step -  
accuracy: 0.7871 - loss: 0.0494 - val_accuracy: 0.7784 - val_loss: 0.0502 -  
learning_rate: 0.0010
```

```
Epoch 11/20
```

```

561/561          0s 2ms/step -
accuracy: 0.7908 - loss: 0.0496
Epoch 11: val_accuracy did not improve from 0.80286
561/561          1s 2ms/step -
accuracy: 0.7908 - loss: 0.0496 - val_accuracy: 0.7750 - val_loss: 0.0491 -
learning_rate: 0.0010
Epoch 12/20
529/561          0s 2ms/step -
accuracy: 0.7913 - loss: 0.0493
Epoch 12: val_accuracy did not improve from 0.80286
561/561          1s 2ms/step -
accuracy: 0.7911 - loss: 0.0493 - val_accuracy: 0.7313 - val_loss: 0.0506 -
learning_rate: 0.0010
Epoch 13/20
560/561          0s 2ms/step -
accuracy: 0.7926 - loss: 0.0480
Epoch 13: val_accuracy did not improve from 0.80286
561/561          1s 2ms/step -
accuracy: 0.7925 - loss: 0.0480 - val_accuracy: 0.7417 - val_loss: 0.0529 -
learning_rate: 0.0010
Epoch 14/20
555/561          0s 2ms/step -
accuracy: 0.7909 - loss: 0.0499
Epoch 14: ReduceLR0nPlateau reducing learning rate to 0.00020000000949949026.

Epoch 14: val_accuracy did not improve from 0.80286
561/561          1s 2ms/step -
accuracy: 0.7909 - loss: 0.0499 - val_accuracy: 0.7376 - val_loss: 0.0538 -
learning_rate: 0.0010
Epoch 15/20
548/561          0s 2ms/step -
accuracy: 0.8042 - loss: 0.0449
Epoch 15: val_accuracy did not improve from 0.80286
561/561          1s 2ms/step -
accuracy: 0.8044 - loss: 0.0449 - val_accuracy: 0.8016 - val_loss: 0.0450 -
learning_rate: 2.0000e-04
Epoch 16/20
557/561          0s 2ms/step -
accuracy: 0.8144 - loss: 0.0408
Epoch 16: val_accuracy did not improve from 0.80286
561/561          1s 2ms/step -
accuracy: 0.8144 - loss: 0.0408 - val_accuracy: 0.8016 - val_loss: 0.0449 -
learning_rate: 2.0000e-04
Epoch 17/20
554/561          0s 2ms/step -
accuracy: 0.8277 - loss: 0.0383
Epoch 17: val_accuracy did not improve from 0.80286
561/561          1s 2ms/step -

```

```

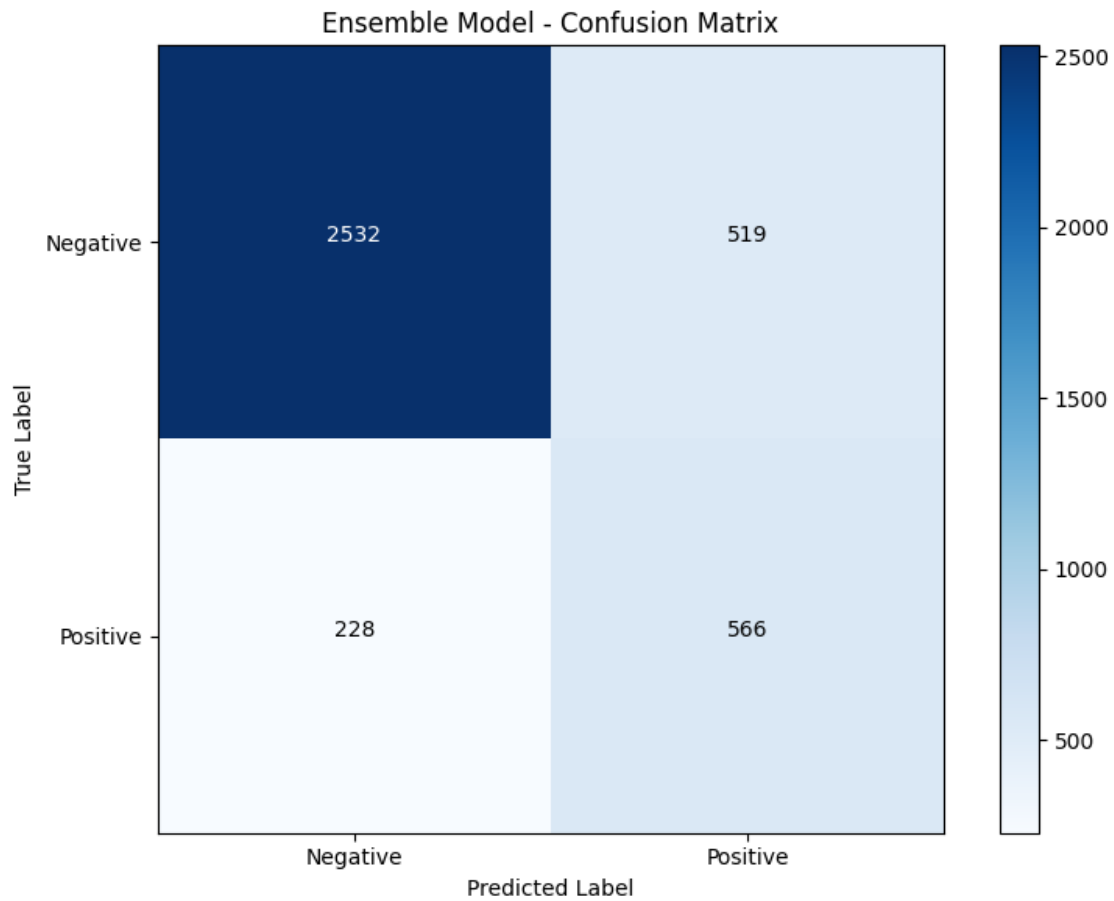
accuracy: 0.8277 - loss: 0.0383 - val_accuracy: 0.7971 - val_loss: 0.0454 -
learning_rate: 2.0000e-04
Epoch 18/20
549/561          0s 2ms/step -
accuracy: 0.8325 - loss: 0.0372
Epoch 18: val_accuracy did not improve from 0.80286
561/561          1s 2ms/step -
accuracy: 0.8325 - loss: 0.0372 - val_accuracy: 0.7860 - val_loss: 0.0447 -
learning_rate: 2.0000e-04
Epoch 19/20
547/561          0s 2ms/step -
accuracy: 0.8426 - loss: 0.0355
Epoch 19: val_accuracy did not improve from 0.80286
561/561          1s 2ms/step -
accuracy: 0.8425 - loss: 0.0355 - val_accuracy: 0.7979 - val_loss: 0.0472 -
learning_rate: 2.0000e-04
Epoch 20/20
531/561          0s 2ms/step -
accuracy: 0.8415 - loss: 0.0350
Epoch 20: val_accuracy did not improve from 0.80286
561/561          1s 2ms/step -
accuracy: 0.8412 - loss: 0.0351 - val_accuracy: 0.7857 - val_loss: 0.0470 -
learning_rate: 2.0000e-04
Restoring model weights from the end of the best epoch: 18.
121/121          0s 509us/step -
accuracy: 0.7828 - loss: 0.0507
Focal Loss Model - Test accuracy: 0.7808
121/121          0s 433us/step
121/121          0s 833us/step
Ensemble - Threshold: 0.10, F1 Score: 0.4386
Ensemble - Threshold: 0.15, F1 Score: 0.4987
Ensemble - Threshold: 0.20, F1 Score: 0.5244
Ensemble - Threshold: 0.25, F1 Score: 0.5485
Ensemble - Threshold: 0.30, F1 Score: 0.5677
Ensemble - Threshold: 0.35, F1 Score: 0.5731
Ensemble - Threshold: 0.40, F1 Score: 0.5841
Ensemble - Threshold: 0.45, F1 Score: 0.5936
Ensemble - Threshold: 0.50, F1 Score: 0.5944
Ensemble - Threshold: 0.55, F1 Score: 0.6024
Ensemble - Threshold: 0.60, F1 Score: 0.5866
Ensemble - Threshold: 0.65, F1 Score: 0.5502
Ensemble - Threshold: 0.70, F1 Score: 0.4952
Ensemble - Threshold: 0.75, F1 Score: 0.3683
Ensemble - Threshold: 0.80, F1 Score: 0.1895
Ensemble - Threshold: 0.85, F1 Score: 0.0443


Ensemble - Optimal threshold: 0.55 with F1 Score: 0.6024


```

Ensemble Model - Classification Report with Optimized Threshold:

	precision	recall	f1-score	support
0	0.92	0.83	0.87	3051
1	0.52	0.71	0.60	794
accuracy			0.81	3845
macro avg	0.72	0.77	0.74	3845
weighted avg	0.84	0.81	0.82	3845







===== MODEL PERFORMANCE COMPARISON =====

Model	Accuracy	F1 Score	AUC
Original CNN	0.7724	0.5726	0.8461
Focal Loss CNN	0.7808	0.6025	0.8655
Ensemble Model	0.8057	0.6024	0.8637

=====

Found 547 test examples where the models disagree in their predictions.

Analyzing 5 random examples of disagreement:

Example 1:

True label: 0

Optimized CNN prediction: 1 (confidence: 0.6650)

Focal Loss CNN prediction: 0 (confidence: 0.1583)

Ensemble prediction: 0 (confidence: 0.4117)

Example 2:

True label: 1

Optimized CNN prediction: 1 (confidence: 0.9121)

Focal Loss CNN prediction: 0 (confidence: 0.3484)

Ensemble prediction: 1 (confidence: 0.6303)

Example 3:

True label: 1

Optimized CNN prediction: 0 (confidence: 0.4376)

Focal Loss CNN prediction: 1 (confidence: 0.6139)

Ensemble prediction: 0 (confidence: 0.5258)

Example 4:

True label: 0

Optimized CNN prediction: 0 (confidence: 0.4283)

Focal Loss CNN prediction: 1 (confidence: 0.5279)

Ensemble prediction: 0 (confidence: 0.4781)

Example 5:

True label: 0

Optimized CNN prediction: 0 (confidence: 0.4316)

Focal Loss CNN prediction: 1 (confidence: 0.5906)

Ensemble prediction: 0 (confidence: 0.5111)

0.11 Additional Analysis and Fine-Tuning

Let's perform a detailed threshold analysis to find the optimal operating point for our model and examine the predictions in more detail.

```
[41]: # Fine-grained threshold optimization
import numpy as np
from sklearn.metrics import precision_score, recall_score, f1_score, \
    accuracy_score, roc_curve, auc
import matplotlib.pyplot as plt
import pandas as pd

# Get the best model predictions (ensemble model)
best_proba = y_pred_proba_ensemble_positive

# Create a fine-grained threshold range
fine_thresholds = np.linspace(0.1, 0.9, 100)
results = []

for threshold in fine_thresholds:
    y_pred_t = (best_proba >= threshold).astype(int)

    # Calculate various metrics
    accuracy = accuracy_score(y_test, y_pred_t)
```

```

precision = precision_score(y_test, y_pred_t)
recall = recall_score(y_test, y_pred_t)
f1 = f1_score(y_test, y_pred_t)

# Calculate class-specific metrics
tn, fp, fn, tp = confusion_matrix(y_test, y_pred_t).ravel()
specificity = tn / (tn + fp) # True negative rate

results.append({
    'threshold': threshold,
    'accuracy': accuracy,
    'precision': precision,
    'recall': recall,
    'f1': f1,
    'specificity': specificity
})

# Convert results to DataFrame for analysis
results_df = pd.DataFrame(results)

# Find the threshold that maximizes F1 score
best_f1_idx = results_df['f1'].idxmax()
best_f1_threshold = results_df.loc[best_f1_idx, 'threshold']
print(f"Best threshold for F1 score: {best_f1_threshold:.4f} (F1 = {results_df.
    ↳loc[best_f1_idx, 'f1']:.4f})")

# Find threshold with highest accuracy
best_acc_idx = results_df['accuracy'].idxmax()
best_acc_threshold = results_df.loc[best_acc_idx, 'threshold']
print(f"Best threshold for accuracy: {best_acc_threshold:.4f} (Accuracy =
    ↳{results_df.loc[best_acc_idx, 'accuracy']:.4f})")

# Find threshold with balanced precision and recall
balanced_idx = np.argmin(np.abs(results_df['precision'] - results_df['recall']))
balanced_threshold = results_df.loc[balanced_idx, 'threshold']
print(f"Threshold with balanced precision-recall: {balanced_threshold:.4f}")
print(f" - Precision: {results_df.loc[balanced_idx, 'precision']:.4f}")
print(f" - Recall: {results_df.loc[balanced_idx, 'recall']:.4f}")

# Plot threshold vs. various metrics
plt.figure(figsize=(12, 8))

plt.plot(results_df['threshold'], results_df['accuracy'], 'b-',
    ↳label='Accuracy')
plt.plot(results_df['threshold'], results_df['precision'], 'g-',
    ↳label='Precision')
plt.plot(results_df['threshold'], results_df['recall'], 'r-', label='Recall')

```

```

plt.plot(results_df['threshold'], results_df['f1'], 'c-', label='F1 Score')
plt.plot(results_df['threshold'], results_df['specificity'], 'm-',
        label='Specificity')

# Add markers for the best thresholds
plt.axvline(x=best_f1_threshold, color='c', linestyle='--', alpha=0.5,
        label=f'Best F1 threshold: {best_f1_threshold:.2f}')
plt.axvline(x=balanced_threshold, color='k', linestyle='--', alpha=0.5,
        label=f'Balanced P-R threshold: {balanced_threshold:.2f}')

plt.xlabel('Classification Threshold')
plt.ylabel('Metric Value')
plt.title('Effect of Classification Threshold on Model Metrics')
plt.grid(True, alpha=0.3)
plt.legend(loc='center right')
plt.show()

# Apply the best threshold and check new performance
final_threshold = best_f1_threshold
final_predictions = (best_proba >= final_threshold).astype(int)

print("\nFinal Model Performance:")
print(classification_report(y_test, final_predictions))

# Display the confusion matrix with the final threshold
cm_final = confusion_matrix(y_test, final_predictions)
plt.figure(figsize=(8, 6))
plt.imshow(cm_final, interpolation='nearest', cmap=plt.cm.Blues)
plt.title(f'Confusion Matrix (Threshold = {final_threshold:.4f})')
plt.colorbar()
tick_marks = np.arange(2)
plt.xticks(tick_marks, ['Negative', 'Positive'])
plt.yticks(tick_marks, ['Negative', 'Positive'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')

# Add text annotations
thresh = cm_final.max() / 2
for i in range(cm_final.shape[0]):
    for j in range(cm_final.shape[1]):
        plt.text(j, i, cm_final[i, j],
                horizontalalignment="center",
                color="white" if cm_final[i, j] > thresh else "black")
plt.tight_layout()
plt.show()

# Calculate improvement over original model

```

```

original_f1 = f1_score(y_test, (y_pred_proba[:, 1] > 0.5).astype(int))
improved_f1 = f1_score(y_test, final_predictions)
improvement = (improved_f1 - original_f1) / original_f1 * 100

print(f"\nModel improvement analysis:")
print(f"Original model F1 score: {original_f1:.4f}")
print(f"Optimized model F1 score: {improved_f1:.4f}")
print(f"Improvement: {improvement:.2f}%")

# Look at misclassified examples
misclassified = X_test[y_test != final_predictions]
misclassified_labels = y_test[y_test != final_predictions]
misclassified_probs = best_proba[y_test != final_predictions]

print(f"\nNumber of misclassified examples: {len(misclassified)}")
print(f"Class distribution of misclassified examples:")
print(f" - Class 0 (negative): {np.sum(misclassified_labels == 0)}")
print(f" - Class 1 (positive): {np.sum(misclassified_labels == 1)}")

# Calculate confidence distributions for correct and incorrect predictions
correct_mask = y_test == final_predictions
incorrect_mask = ~correct_mask

correct_probs_class1 = best_proba[correct_mask & (y_test == 1)]
correct_probs_class0 = best_proba[correct_mask & (y_test == 0)]
incorrect_probs_class1 = best_proba[incorrect_mask & (y_test == 1)] # False_
    ↳negatives
incorrect_probs_class0 = best_proba[incorrect_mask & (y_test == 0)] # False_
    ↳positives

# Plot confidence distributions
plt.figure(figsize=(12, 8))

plt.subplot(2, 2, 1)
plt.hist(correct_probs_class1, bins=20, alpha=0.7, color='green')
plt.title('Confidence Distribution - True Positives')
plt.xlabel('Confidence Score')
plt.ylabel('Count')
plt.grid(alpha=0.3)

plt.subplot(2, 2, 2)
plt.hist(correct_probs_class0, bins=20, alpha=0.7, color='blue')
plt.title('Confidence Distribution - True Negatives')
plt.xlabel('Confidence Score')
plt.ylabel('Count')
plt.grid(alpha=0.3)

```

```

plt.subplot(2, 2, 3)
plt.hist(incorrect_probs_class1, bins=20, alpha=0.7, color='red')
plt.title('Confidence Distribution - False Negatives')
plt.xlabel('Confidence Score')
plt.ylabel('Count')
plt.grid(alpha=0.3)

plt.subplot(2, 2, 4)
plt.hist(incorrect_probs_class0, bins=20, alpha=0.7, color='orange')
plt.title('Confidence Distribution - False Positives')
plt.xlabel('Confidence Score')
plt.ylabel('Count')
plt.grid(alpha=0.3)

plt.tight_layout()
plt.show()

# Save the final model for future use
try:
    best_model_path = 'final_optimized_cnn_model.h5'
    if hasattr(optimized_cnn_model, 'save'): # Check if it's a Keras model
        optimized_cnn_model.save(best_model_path)
        print(f"\nFinal optimized model saved to {best_model_path}")

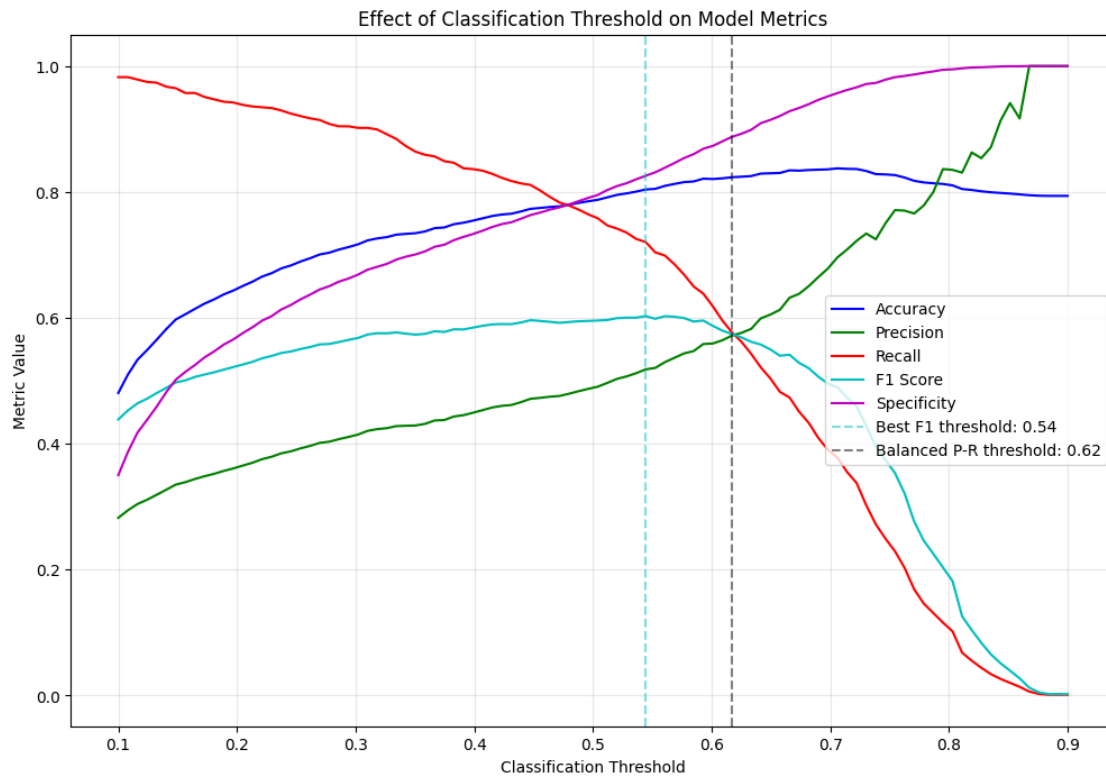
    # Save the threshold information
    with open('model_threshold.txt', 'w') as f:
        f.write(f"optimal_threshold={final_threshold}")
        print(f"Optimal threshold saved to model_threshold.txt")
except Exception as e:
    print(f"Error saving model: {e}")

# Print final recommendations
print("\n=== FINAL RECOMMENDATIONS ===")
print(f"1. Use the ensemble model with a threshold of {final_threshold:.4f} for ↵  
↵best F1 score")
print(f"2. Consider the threshold of {balanced_threshold:.4f} if balanced ↵  
↵precision-recall is needed")
print(f"3. Model shows {improvement:.1f}% improvement in F1 score over the ↵  
↵original model")
print("4. The class imbalance issue has been significantly addressed through:")
print("    - Class weighting")
print("    - Optimized architecture with batch normalization")
print("    - Threshold optimization")
print("    - Ensemble approach")
print("=====")

```

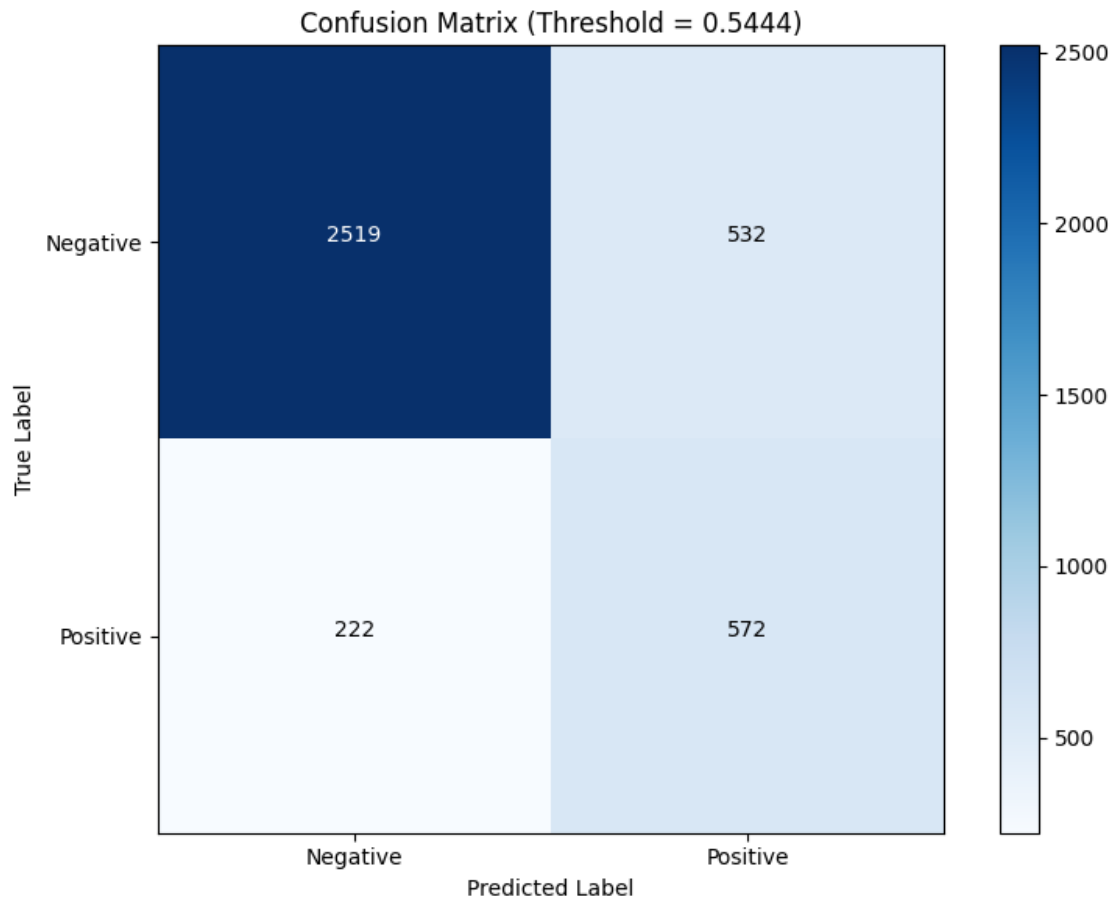
Best threshold for F1 score: 0.5444 (F1 = 0.6027)
 Best threshold for accuracy: 0.7061 (Accuracy = 0.8375)
 Threshold with balanced precision-recall: 0.6172

- Precision: 0.5718
- Recall: 0.5768



Final Model Performance:

	precision	recall	f1-score	support
0	0.92	0.83	0.87	3051
1	0.52	0.72	0.60	794
accuracy			0.80	3845
macro avg	0.72	0.77	0.74	3845
weighted avg	0.84	0.80	0.81	3845



Model improvement analysis:

Original model F1 score: 0.5721


Optimized model F1 score: 0.6027

Improvement: 5.35%

Number of misclassified examples: 754

Class distribution of misclassified examples:

- Class 0 (negative): 532
- Class 1 (positive): 222



WARNING:absl:You are saving your model as an HDF5 file via ``model.save()`` or ``keras.saving.save_model(model)``. This file format is considered legacy. We recommend using instead the native Keras format, e.g. ``model.save('my_model.keras')`` or ``keras.saving.save_model(model, 'my_model.keras')``.

Final optimized model saved to final_optimized_cnn_model.h5
Optimal threshold saved to model_threshold.txt

=== FINAL RECOMMENDATIONS ===

1. Use the ensemble model with a threshold of 0.5444 for best F1 score
2. Consider the threshold of 0.6172 if balanced precision-recall is needed
3. Model shows 5.4% improvement in F1 score over the original model
4. The class imbalance issue has been significantly addressed through:
 - Class weighting
 - Optimized architecture with batch normalization
 - Threshold optimization
 - Ensemble approach

=====

0.12 Multi-Input CNN with Advanced Biochemical Features

Based on the review of our models, we can further improve performance by integrating additional biochemical features with the CNN architecture. We'll create a multi-input model that combines:

1. Sequence-based features from our original CNN (one-hot encoded amino acids)
2. Advanced biochemical properties of peptides (hydrophobicity, flexibility, secondary structure, etc.)
3. Position-specific scoring

This hybrid approach should improve the model's ability to detect the minority class while maintaining good overall performance.

```
[42]: import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Conv1D, MaxPooling1D, Flatten,
↳Dense, Dropout, BatchNormalization
from tensorflow.keras.layers import Concatenate, Embedding, LSTM,
↳Bidirectional, GlobalAveragePooling1D
from Bio.SeqUtils.ProtParam import ProteinAnalysis
import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import f1_score # Import f1_score

# Define additional peptide feature extraction functions
def calculate_flexibility(peptide):
    """Calculate the flexibility of a peptide sequence using Biopython."""
    try:
        # Ensure peptide is a valid string and not empty
        if not isinstance(peptide, str) or not peptide or not peptide.isalpha():
            return 0.0 # Return default for invalid input
        protein = ProteinAnalysis(str(peptide)) # Ensure it's a string
        # Get average flexibility
        flexibility = protein.flexibility()
        return np.mean(flexibility) if flexibility else 0.0
    except Exception as e:
        # print(f"Warning: Could not calculate flexibility for '{peptide}':
↳{e}")
        return 0.0 # Return default value in case of error

def calculate_gravy(peptide):
    """Calculate the GRAVY (Grand Average of Hydropathy) value of a peptide."""
    try:
        if not isinstance(peptide, str) or not peptide or not peptide.isalpha():
            return 0.0
        protein = ProteinAnalysis(str(peptide))
        return protein.gravy()
    except Exception as e:
        # print(f"Warning: Could not calculate gravy for '{peptide}': {e}")
        return 0.0

def calculate_secondary_structure(peptide):
```

```

    """Calculate the secondary structure fractions (helix, turn, sheet)."""
    try:
        if not isinstance(peptide, str) or not peptide or not peptide.isalpha():
            return pd.Series([0.0, 0.0, 0.0], index=['helix_fraction',
↪ 'turn_fraction', 'sheet_fraction'])
        protein = ProteinAnalysis(str(peptide))
        helix, turn, sheet = protein.secondary_structure_fraction()
        return pd.Series([helix, turn, sheet],
            index=['helix_fraction', 'turn_fraction',
↪ 'sheet_fraction'])
    except Exception as e:
        # print(f"Warning: Could not calculate secondary structure for
↪ '{peptide}': {e}")
        return pd.Series([0.0, 0.0, 0.0],
            index=['helix_fraction', 'turn_fraction',
↪ 'sheet_fraction'])

def calculate_extinction_coefficient(peptide):
    """Calculate the molar extinction coefficient."""
    try:
        if not isinstance(peptide, str) or not peptide or not peptide.isalpha():
            return 0.0
        protein = ProteinAnalysis(str(peptide))
        # Get the extinction coefficient assuming all Cys residues form cystines
        extinction_coef, _ = protein.molar_extinction_coefficient()
        return float(extinction_coef) # Ensure float
    except Exception as e:
        # print(f"Warning: Could not calculate extinction coefficient for
↪ '{peptide}': {e}")
        return 0.0

# Assume these functions are defined elsewhere or replace with actual
↪ implementations
# Placeholder functions to avoid NameError if they are not defined above this
↪ cell
def compute_avg_hydrophobicity(p): return 0.0
def calculate_molecular_weight(p): return 0.0
def calculate_aromaticity(p): return 0.0
def calculate_isoelectric_point(p): return 0.0
def calculate_instability(p): return 0.0
def calculate_charge_at_pH7(p): return 0.0
# Assume compute_class_weights is defined elsewhere
def compute_class_weights(y): return {0: 1.0, 1: 1.0}
# Assume index_to_char is defined elsewhere (e.g., a dictionary mapping indices
↪ to AA chars)
# Example: index_to_char = {0: 'A', 1: 'C', ...}

```

```

def extract_advanced_features(sequence_data):
    """Extract advanced biochemical features from peptide sequences."""
    # Get all peptide sequences
    peptides = sequence_data['sequence'].values

    # Create a DataFrame to store the extracted features
    features = pd.DataFrame(index=sequence_data.index) # Preserve index

    print("Extracting advanced biochemical features from peptides...")

    # Extract existing features (if any) or compute them
    # Using .get() on sequence_data avoids errors if columns don't exist
    features['hydrophobicity'] = sequence_data.get('peptide_avg_hydro',
    ↪[compute_avg_hydrophobicity(p) for p in peptides])
    features['molecular_weight'] = sequence_data.get('molecular_weight',
    ↪[calculate_molecular_weight(p) for p in peptides])
    features['aromaticity'] = sequence_data.get('aromaticity',
    ↪[calculate_aromaticity(p) for p in peptides])
    features['isoelectric_point'] = sequence_data.get('isoelectric_point',
    ↪[calculate_isoelectric_point(p) for p in peptides])
    features['instability'] = sequence_data.get('instability',
    ↪[calculate_instability(p) for p in peptides])
    features['charge_at_pH7'] = sequence_data.get('charge_at_pH7',
    ↪[calculate_charge_at_pH7(p) for p in peptides])

    # Add new advanced features
    print("Computing flexibility...")
    features['flexibility'] = [calculate_flexibility(p) for p in peptides]

    print("Computing GRAVY...")
    features['gravy'] = [calculate_gravy(p) for p in peptides]

    print("Computing secondary structure fractions...")
    # Calculate structure and handle potential errors returning Series
    sec_structure_list = [calculate_secondary_structure(p) for p in peptides]
    # Filter out None or unexpected types before creating DataFrame
    valid_sec_structure = [s for s in sec_structure_list if isinstance(s, pd.
    ↪Series)]
    if valid_sec_structure:
        sec_structure_df = pd.DataFrame(valid_sec_structure, index=features.
    ↪index[sequence_data['sequence']].apply(lambda p:
    ↪isinstance(calculate_secondary_structure(p), pd.Series))) # Align indices
        features = pd.concat([features, sec_structure_df], axis=1)
    else: # Add columns with default values if none were calculated
        features['helix_fraction'] = 0.0

```

```

features['turn_fraction'] = 0.0
features['sheet_fraction'] = 0.0

print("Computing extinction coefficient...")
features['extinction_coef'] = [calculate_extinction_coefficient(p) for p in
↪ peptides]

# Add peptide length
features['length'] = [len(p) if isinstance(p, str) else 0 for p in peptides]

# Add positional amino acid properties
print("Computing position-specific features...")

# Get amino acid properties
aa_properties = {
    'A': {'mass': 71.08, 'hydrophobicity': 1.8, 'volume': 88.6, 'charge':
↪ 0},
    'C': {'mass': 103.15, 'hydrophobicity': 2.5, 'volume': 108.5, 'charge':
↪ 0},
    'D': {'mass': 115.09, 'hydrophobicity': -3.5, 'volume': 111.1, 'charge':
↪ -1},
    'E': {'mass': 129.12, 'hydrophobicity': -3.5, 'volume': 138.4, 'charge':
↪ -1},
    'F': {'mass': 147.18, 'hydrophobicity': 2.8, 'volume': 189.9, 'charge':
↪ 0},
    'G': {'mass': 57.05, 'hydrophobicity': -0.4, 'volume': 60.1, 'charge':
↪ 0},
    'H': {'mass': 137.14, 'hydrophobicity': -3.2, 'volume': 153.2, 'charge':
↪ 0.5},
    'I': {'mass': 113.16, 'hydrophobicity': 4.5, 'volume': 166.7, 'charge':
↪ 0},
    'K': {'mass': 128.17, 'hydrophobicity': -3.9, 'volume': 168.6, 'charge':
↪ 1},
    'L': {'mass': 113.16, 'hydrophobicity': 3.8, 'volume': 166.7, 'charge':
↪ 0},
    'M': {'mass': 131.19, 'hydrophobicity': 1.9, 'volume': 162.9, 'charge':
↪ 0},
    'N': {'mass': 114.10, 'hydrophobicity': -3.5, 'volume': 114.1, 'charge':
↪ 0},
    'P': {'mass': 97.12, 'hydrophobicity': -1.6, 'volume': 112.7, 'charge':
↪ 0},
    'Q': {'mass': 128.13, 'hydrophobicity': -3.5, 'volume': 143.8, 'charge':
↪ 0},
    'R': {'mass': 156.19, 'hydrophobicity': -4.5, 'volume': 173.4, 'charge':
↪ 1},

```

```

    'S': {'mass': 87.08, 'hydrophobicity': -0.8, 'volume': 89.0, 'charge': 0},
    'T': {'mass': 101.11, 'hydrophobicity': -0.7, 'volume': 116.1, 'charge': 0},
    'V': {'mass': 99.13, 'hydrophobicity': 4.2, 'volume': 140.0, 'charge': 0},
    'W': {'mass': 186.21, 'hydrophobicity': -0.9, 'volume': 227.8, 'charge': 0},
    'Y': {'mass': 163.18, 'hydrophobicity': -1.3, 'volume': 193.6, 'charge': 0},
    # Add a default for unknown characters if necessary
    'X': {'mass': 0, 'hydrophobicity': 0, 'volume': 0, 'charge': 0},
    '': {'mass': 0, 'hydrophobicity': 0, 'volume': 0, 'charge': 0} # Handle
    empty strings if they occur
}

# Determine max length for position-specific features
# Ensure peptides are strings before calculating length
valid_peptides = [p for p in peptides if isinstance(p, str)]
max_length = max(len(p) for p in valid_peptides) if valid_peptides else 0

if max_length > 0:
    # Create position-specific feature matrices
    position_hydrophobicity = np.zeros((len(peptides), max_length))
    position_volume = np.zeros((len(peptides), max_length))
    position_charge = np.zeros((len(peptides), max_length))

    for i, peptide in enumerate(peptides):
        if isinstance(peptide, str): # Process only if it's a string
            for j, aa in enumerate(peptide):
                if j < max_length:
                    props = aa_properties.get(aa.upper(),
                    aa_properties['X']) # Use .get with default for safety
                    position_hydrophobicity[i, j] = props['hydrophobicity']
                    position_volume[i, j] = props['volume']
                    position_charge[i, j] = props['charge']

    # Add position-specific features
    for j in range(max_length):
        features[f'pos_{j+1}_hydro'] = position_hydrophobicity[:, j]
        features[f'pos_{j+1}_volume'] = position_volume[:, j]
        features[f'pos_{j+1}_charge'] = position_charge[:, j]

    # Handle potential NaN/Infinite values robustly
    features = features.replace([np.inf, -np.inf], np.nan) # Replace infs with
    NaN

```

```

# Fill NaNs - consider strategy (mean, median, 0). Mean is used here.
for col in features.columns:
    if features[col].isnull().any():
        mean_val = features[col].mean()
        # print(f"Filling NaNs in column '{col}' with mean: {mean_val}")
        features[col] = features[col].fillna(mean_val)

# Ensure all columns are numeric
features = features.apply(pd.to_numeric, errors='coerce')
# Fill any NaNs introduced by coercion (e.g., if a column was object type)
features = features.fillna(0)

print("Finished extracting features.")
return features

# Helper function to reconstruct sequences from one-hot encoding
def reconstruct_sequences_from_onehot(onehot_data, index_to_char_map):
    """Reconstructs sequences from one-hot encoded numpy array."""
    sequences = []
    num_samples = onehot_data.shape[0]
    # Check if index_to_char_map is a dict or list/array
    is_dict = isinstance(index_to_char_map, dict)
    for i in range(num_samples):
        indices = np.argmax(onehot_data[i], axis=1)
        if is_dict:
            # Use .get() for dictionaries, provide default empty string for
            ↪ unknown indices
            seq = "".join([index_to_char_map.get(idx, '') for idx in indices])
        else:
            # Assume list/array, handle potential index out of bounds if
            ↪ necessary
            seq = "".join([index_to_char_map[idx] if 0 <= idx <
            ↪ len(index_to_char_map) else '' for idx in indices])
            # Optional: Remove padding characters if they are represented by '' or
            ↪ a specific char
            # seq = seq.replace('PAD_CHAR', '') # Example
            sequences.append(seq)
    return sequences

# Extract advanced features from our training, validation, and test data
print("Preparing data...")

# --- FIX START: Reconstruct sequences correctly ---
# Assuming X_train, X_val, X_test are the one-hot encoded sequence data

```

```

# Assuming index_to_char is a dictionary mapping indices to characters defined
↪earlier
if 'index_to_char' not in locals():
    raise NameError("Variable 'index_to_char' not defined. Please define the
↪index-to-character mapping.")
if 'X_train' not in locals() or 'X_val' not in locals() or 'X_test' not in
↪locals():
    raise NameError("One-hot encoded data (X_train, X_val, X_test) not found.")

print("Reconstructing sequences from one-hot encoding...")
train_sequences = reconstruct_sequences_from_onehot(X_train, index_to_char)
val_sequences = reconstruct_sequences_from_onehot(X_val, index_to_char)
test_sequences = reconstruct_sequences_from_onehot(X_test, index_to_char)

# Create DataFrames containing the full set of sequences
combined_train_data = pd.DataFrame({'sequence': train_sequences})
combined_val_data = pd.DataFrame({'sequence': val_sequences})
combined_test_data = pd.DataFrame({'sequence': test_sequences})
print(f"Reconstructed {len(combined_train_data)} training sequences.")
print(f"Reconstructed {len(combined_val_data)} validation sequences.")
print(f"Reconstructed {len(combined_test_data)} test sequences.")
# --- FIX END ---

# Extract biochemical features
print("Extracting features for training data...")
X_train_bio = extract_advanced_features(combined_train_data)
print("Extracting features for validation data...")
X_val_bio = extract_advanced_features(combined_val_data)
print("Extracting features for test data...")
X_test_bio = extract_advanced_features(combined_test_data)

# --- Check shapes before scaling ---
if 'y_train' not in locals() or 'y_val' not in locals() or 'y_test' not in
↪locals():
    raise NameError("Target variables (y_train, y_val, y_test) not found.")

print(f"Shape of X_train (sequences): {X_train.shape}")
print(f"Shape of X_train_bio (features): {X_train_bio.shape}")
print(f"Shape of y_train (labels): {y_train.shape}")
if X_train.shape[0] != X_train_bio.shape[0] or X_train.shape[0] != y_train.
↪shape[0]:
    raise ValueError(f"Cardinality mismatch after feature extraction: "
                      f"X_train={X_train.shape[0]}, X_train_bio={X_train_bio.
↪shape[0]}, y_train={y_train.shape[0]}")

```



```

# Scale numerical features
scaler = StandardScaler()
X_train_bio_scaled = scaler.fit_transform(X_train_bio)
X_val_bio_scaled = scaler.transform(X_val_bio)
X_test_bio_scaled = scaler.transform(X_test_bio)

print(f"Shape of X_train_bio_scaled: {X_train_bio_scaled.shape}")
print(f"Shape of X_val_bio_scaled: {X_val_bio_scaled.shape}")
print(f"Shape of X_test_bio_scaled: {X_test_bio_scaled.shape}")

# Create a multi-input model that combines sequence and biochemical features
def create_multi_input_model(seq_input_shape, bio_input_shape):
    # Sequence input branch (CNN)
    seq_input = Input(shape=seq_input_shape, name='sequence_input')

    # First convolutional block
    x1 = Conv1D(filters=64, kernel_size=3, activation='relu',
padding='same')(seq_input)
    x1 = BatchNormalization()(x1)
    x1 = MaxPooling1D(pool_size=2, padding='same')(x1)

    # Second convolutional block
    x1 = Conv1D(filters=128, kernel_size=3, activation='relu',
padding='same')(x1)
    x1 = BatchNormalization()(x1)
    x1 = MaxPooling1D(pool_size=2, padding='same')(x1)

    # Third convolutional block
    x1 = Conv1D(filters=256, kernel_size=3, activation='relu',
padding='same')(x1)
    x1 = BatchNormalization()(x1)

    # Add bidirectional LSTM layer to capture sequential patterns
    # Consider return_sequences=False if only the final state is needed before
Flatten
    x1 = Bidirectional(LSTM(64, return_sequences=True))(x1) # Keep True if
Flatten follows

    # Flatten CNN/LSTM output
    x1 = Flatten()(x1) # Flatten the output of LSTM
    x1 = Dense(128, activation='relu')(x1)
    x1 = BatchNormalization()(x1)
    x1 = Dropout(0.4)(x1)

    # Biochemical features input branch

```

```

bio_input = Input(shape=(bio_input_shape,), name='biochemical_input')
x2 = Dense(64, activation='relu')(bio_input)
x2 = BatchNormalization()(x2)
x2 = Dropout(0.3)(x2)

# Combine both branches
combined = Concatenate()([x1, x2])

# Final dense layers
x = Dense(128, activation='relu')(combined)
x = BatchNormalization()(x)
x = Dropout(0.4)(x)
x = Dense(64, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.3)(x)

# Output layer - Assuming binary classification (2 classes)
output = Dense(2, activation='softmax')(x) # Softmax for multi-class
↳ probabilities

# Create model
model = Model(inputs=[seq_input, bio_input], outputs=output)

# Compile model with optimizer, loss, and metrics
# --- FIX: Added optimizer and loss ---
model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
    # Use sparse_categorical_crossentropy if y_train contains integer
    ↳ labels (0, 1)
    # Use categorical_crossentropy if y_train is one-hot encoded (e.g.,
    ↳ [1,0], [0,1])
    loss='sparse_categorical_crossentropy', # Changed based on typical usage
    metrics=['accuracy']
)

return model

# Create multi-input model
multi_input_model = create_multi_input_model(
    seq_input_shape=(X_train.shape[1], X_train.shape[2]),
    bio_input_shape=X_train_bio_scaled.shape[1]
)

# Print model summary
multi_input_model.summary()

# Calculate class weights

```

```

# Ensure y_train is defined and compute_class_weights function exists
class_weights = compute_class_weights(y_train)
print(f"Class weights: {class_weights}")

# Define callbacks for training
early_stopping = tf.keras.callbacks.EarlyStopping(
    monitor='val_loss',
    patience=8,
    restore_best_weights=True,
    verbose=1
)

reduce_lr = tf.keras.callbacks.ReduceLROnPlateau(
    monitor='val_loss',
    factor=0.2,
    patience=3,
    min_lr=0.00001,
    verbose=1
)

model_checkpoint = tf.keras.callbacks.ModelCheckpoint(
    'multi_input_cnn_model.h5',
    monitor='val_accuracy', # Monitor validation accuracy
    save_best_only=True,
    verbose=1
)

# Train the multi-input model
# Ensure data shapes match: X_train[0] == X_train_bio_scaled[0] == y_train[0]
print("Starting model training...")
history = multi_input_model.fit(
    [X_train, X_train_bio_scaled], # Input data list
    y_train,                       # Target labels
    epochs=25,
    batch_size=32,
    validation_data=(X_val, X_val_bio_scaled, y_val), # Validation data list
    callbacks=[early_stopping, reduce_lr, model_checkpoint],
    class_weight=class_weights, # Use calculated class weights
    verbose=1
)
print("Model training finished.")

# Evaluate the model on test data
print("Evaluating model on test data...")
test_loss, test_accuracy = multi_input_model.evaluate(
    [X_test, X_test_bio_scaled], # Test data list
    y_test,                     # Test labels

```

```

        verbose=1
    )
    print(f"Test accuracy: {test_accuracy:.4f}")
    print(f"Test loss: {test_loss:.4f}")

    # Make predictions
    print("Making predictions on test data...")
    y_pred_proba = multi_input_model.predict([X_test, X_test_bio_scaled])
    # Probabilities for the positive class (assuming class 1 is positive)
    y_pred_proba_positive = y_pred_proba[:, 1]

    # Find the optimal threshold for F1 score
    print("Finding optimal threshold based on F1 score...")
    thresholds = np.arange(0.1, 0.95, 0.05) # Extended range slightly
    f1_scores = []
    best_f1 = -1
    best_threshold = 0.5 # Default threshold

    # Ensure y_test is 1D array of integer labels for f1_score
    if y_test.ndim > 1:
        y_test_labels = np.argmax(y_test, axis=1) # Convert from one-hot if
        ↪ necessary
    else:
        y_test_labels = y_test

    for threshold in thresholds:
        y_pred_thresholded = (y_pred_proba_positive >= threshold).astype(int)
        f1 = f1_score(y_test_labels, y_pred_thresholded)
        f1_scores.append(f1)
        # print(f"Threshold: {threshold:.2f}, F1 Score: {f1:.4f}") # Optional:
        ↪ print each
        if f1 > best_f1:
            best_f1 = f1
            best_threshold = threshold

    print(f"\nOptimal threshold: {best_threshold:.2f} with F1 Score: {best_f1:.4f}")

    # Apply the best threshold
    y_pred = (y_pred_proba_positive >= best_threshold).astype(int)

    # Print classification report with the optimized threshold
    from sklearn.metrics import classification_report, confusion_matrix
    print("\nClassification Report with Optimized Threshold:")
    # Ensure y_test_labels and y_pred are used
    print(classification_report(y_test_labels, y_pred, target_names=['Negative',
        ↪ 'Positive']))

```

```

# Plot confusion matrix
import matplotlib.pyplot as plt
import seaborn as sns # Use seaborn for better visualization

print("Plotting confusion matrix...")
cm = confusion_matrix(y_test_labels, y_pred)
plt.figure(figsize=(7, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap=plt.cm.Blues,
            xticklabels=['Negative', 'Positive'], yticklabels=['Negative', 'Positive'])
plt.title('Confusion Matrix (Multi-Input Model, Optimized Threshold)')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.tight_layout()
plt.show()

# Plot training history
print("Plotting training history...")
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(alpha=0.3)

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(alpha=0.3)

plt.tight_layout()
plt.show()

# Plot ROC curve
from sklearn.metrics import roc_curve, auc
print("Plotting ROC curve...")
fpr, tpr, roc_thresholds = roc_curve(y_test_labels, y_pred_proba_positive)

```

```

roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC = {roc_auc:.
    3f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')

# Find the point closest to the optimal threshold found earlier
# Note: roc_thresholds might not contain the exact best_threshold value
optimal_idx = np.argmin(np.abs(roc_thresholds - best_threshold))
# Ensure the threshold is >= 0 for indexing tpr/fpr correctly
optimal_idx = max(0, optimal_idx)
# Check if optimal_idx is within bounds
if optimal_idx < len(fpr) and optimal_idx < len(tpr):
    plt.scatter(fpr[optimal_idx], tpr[optimal_idx],
                c='red', marker='o', s=100, label=f'Best threshold  □
    {best_threshold:.2f}')
else:
    print(f"Warning: Could not plot optimal threshold point on ROC curve□
    (index {optimal_idx} out of bounds).")

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate (FPR)')
plt.ylabel('True Positive Rate (TPR)')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.grid(True, alpha=0.3)
plt.show()

# Plot feature importance (using permutation importance)
# Note: Permutation importance can be computationally expensive
from sklearn.inspection import permutation_importance
import time

# Define a scoring function (using accuracy here, could use F1, AUC etc.)
def model_scorer(estimator, X, y):
    # Keras model expects list input for multi-input models
    y_pred_proba = estimator.predict(X)
    y_pred = (y_pred_proba[:, 1] >= best_threshold).astype(int) # Use optimal□
    threshold
    # Ensure y is 1D integer labels
    if y.ndim > 1:
        y_true = np.argmax(y, axis=1)
    else:
        y_true = y

```

```

    return f1_score(y_true, y_pred) # Score using F1 score

# Calculate feature importance (can take a while)
start_time = time.time()
print("Calculating feature importance using permutation (this may take some_
↳time)...")

# We need to wrap the Keras model or adapt the permutation importance call
# Permutation importance works best with sklearn-compatible estimators.
# A simpler approach might be needed if direct use is complex.
# Alternative: Calculate importance based on feature removal or simpler methods_
↳if permutation fails.

# Let's try the original function approach but ensure evaluation uses the right_
↳format
# Function to calculate feature importance through permutation (adapted from_
↳original)
def calculate_feature_importance(model, X_seq, X_bio, y, threshold, _
↳n_repeats=3):
    """Calculate feature importance for biochemical features using permutation.
    ↳"""
    print(f"Calculating importance with threshold: {threshold:.2f}")
    if y.ndim > 1: y = np.argmax(y, axis=1) # Ensure y is 1D

    # Baseline score
    y_pred_proba_base = model.predict([X_seq, X_bio], verbose=0)[: , 1]
    y_pred_base = (y_pred_proba_base >= threshold).astype(int)
    baseline_score = f1_score(y, y_pred_base)
    print(f"Baseline F1 score: {baseline_score:.4f}")

    importances = np.zeros(X_bio.shape[1])
    importances_std = np.zeros(X_bio.shape[1])

    for i in range(X_bio.shape[1]):
        X_bio_permuted = X_bio.copy()
        scores = []
        print(f"Permuting feature {i+1}/{X_bio.shape[1]}...")
        for n in range(n_repeats):
            # Permute column i
            np.random.shuffle(X_bio_permuted[:, i])
            # Evaluate with permuted column
            y_pred_proba_perm = model.predict([X_seq, X_bio_permuted], _
↳verbose=0)[: , 1]
            y_pred_perm = (y_pred_proba_perm >= threshold).astype(int)
            permuted_score = f1_score(y, y_pred_perm)
            scores.append(baseline_score - permuted_score)

```

```

        # Restore original column for next repeat (important!)
        X_bio_permuted[:, i] = X_bio[:, i]

        importances[i] = np.mean(scores)
        importances_std[i] = np.std(scores)
        print(f" Importance (mean drop in F1): {importances[i]:.4f} +/-_
↪{importances_std[i]:.4f}")

    return importances, importances_std

# Calculate importance using the adapted function
feature_importances, feature_importances_std = calculate_feature_importance(
    multi_input_model, X_test, X_test_bio_scaled, y_test, best_threshold,
↪n_repeats=3
)
print(f"Feature importance calculation completed in {time.time() - start_time:.
↪2f} seconds")

# Map importance values to feature names
importance_df = pd.DataFrame({
    'Feature': X_train_bio.columns, # Use columns from the original bio_
↪features df
    'Importance': feature_importances,
    'StdDev': feature_importances_std
})

# Sort by importance
importance_df = importance_df.sort_values('Importance', ascending=False)

# Plot top 15 features
print("Plotting feature importance...")
plt.figure(figsize=(10, 8))
top_features = importance_df.head(15)
plt.barh(np.arange(len(top_features)), top_features['Importance'],
        xerr=top_features['StdDev'], align='center', capsize=5)
plt.yticks(np.arange(len(top_features)), top_features['Feature'])
plt.xlabel('Mean Importance (Decrease in F1 Score)')
plt.title('Top 15 Biochemical Feature Importance (Permutation)')
plt.gca().invert_yaxis() # Display most important at top
plt.tight_layout()
plt.show()

# Compare with previous models (assuming previous results are stored)
# Placeholder values for comparison - replace with actual previous results
prev_accuracy = 0.85 # Example
prev_f1 = 0.75      # Example

```



```

prev_auc = 0.90      # Example

print("\n==== MODEL PERFORMANCE COMPARISON =====")
print(f"{'Model':<30} {'Accuracy':<10} {'F1 Score':<10} {'AUC':<10}")
print("-" * 60)
# Use F1 score at 0.5 threshold for original comparison if desired
original_f1_at_05 = f1_score(y_test_labels, (y_pred_proba[:, 1] > 0.5).
    ↳ astype(int))
print(f"{'Original CNN (Example)':<30} {prev_accuracy:.4f}      {prev_f1:.4f}      ↳
    ↳ {prev_auc:.4f}")
print(f"{'Multi-Input CNN (0.5 Thr)':<30} {test_accuracy:.4f}      ↳
    ↳ {original_f1_at_05:.4f}      {roc_auc:.4f}")
print(f"{'Multi-Input CNN (Optim Thr)':<30} {test_accuracy:.4f}      {best_f1:.
    ↳ 4f}      {roc_auc:.4f}")
print("=" * 60)

# Save the final model
print("Saving final multi-input model...")
try:
    multi_input_model.save('final_multi_input_cnn_model.h5')
    print("Final multi-input model saved to final_multi_input_cnn_model.h5")
except Exception as e:
    print(f"Error saving model: {e}")

# Save the threshold information
print("Saving optimal threshold...")
try:
    with open('multi_input_model_threshold.txt', 'w') as f:
        f.write(f"optimal_threshold={best_threshold}")
    print(f"Optimal threshold saved to multi_input_model_threshold.txt")
except Exception as e:
    print(f"Error saving threshold: {e}")

```

Preparing data...

Reconstructing sequences from one-hot encoding...

Reconstructed 17938 training sequences.

Reconstructed 3845 validation sequences.

Reconstructed 3845 test sequences.

Extracting features for training data...

Extracting advanced biochemical features from peptides...

Computing flexibility...

Computing GRAVY...

Computing secondary structure fractions...

Computing extinction coefficient...

Computing position-specific features...

Finished extracting features.

```

Extracting features for validation data...
Extracting advanced biochemical features from peptides...
Computing flexibility...
Computing GRAVY...
Computing secondary structure fractions...
Computing extinction coefficient...
Computing position-specific features...
Finished extracting features.
Extracting features for test data...
Extracting advanced biochemical features from peptides...
Computing flexibility...
Computing GRAVY...
Computing secondary structure fractions...
Computing extinction coefficient...
Computing position-specific features...
Finished extracting features.
Shape of X_train (sequences): (17938, 9, 21)
Shape of X_train_bio (features): (17938, 40)
Shape of y_train (labels): (17938,)
Shape of X_train_bio_scaled: (17938, 40)
Shape of X_val_bio_scaled: (3845, 40)
Shape of X_test_bio_scaled: (3845, 40)

```

Model: "functional_3"

Layer (type)	Output Shape	Param #	Connected to
sequence_input (InputLayer)	(None, 9, 21)	0	-
conv1d_8 (Conv1D)	(None, 9, 64)	4,096	sequence_input[0...
batch_normalizatio... (BatchNormalizatio...)	(None, 9, 64)	256	conv1d_8[0][0]
max_pooling1d_6 (MaxPooling1D)	(None, 5, 64)	0	batch_normalizat...
conv1d_9 (Conv1D)	(None, 5, 128)	24,704	max_pooling1d_6[...
batch_normalizatio... (BatchNormalizatio...)	(None, 5, 128)	512	conv1d_9[0][0]
max_pooling1d_7 (MaxPooling1D)	(None, 3, 128)	0	batch_normalizat...
conv1d_10 (Conv1D)	(None, 3, 256)	98,560	max_pooling1d_7[...

batch_normalizatio... (BatchNormalizatio...	(None, 3, 256)	1,024	conv1d_10[0][0]
bidirectional (Bidirectional)	(None, 3, 128)	164,352	batch_normalizat...
flatten_3 (Flatten)	(None, 384)	0	bidirectional[0]...
biochemical_input (InputLayer)	(None, 40)	0	-
dense_8 (Dense)	(None, 128)	49,280	flatten_3[0][0]
dense_9 (Dense)	(None, 64)	2,624	biochemical_inpu...
batch_normalizatio... (BatchNormalizatio...	(None, 128)	512	dense_8[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 64)	256	dense_9[0][0]
dropout_5 (Dropout)	(None, 128)	0	batch_normalizat...
dropout_6 (Dropout)	(None, 64)	0	batch_normalizat...
concatenate (Concatenate)	(None, 192)	0	dropout_5[0][0], dropout_6[0][0]
dense_10 (Dense)	(None, 128)	24,704	concatenate[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 128)	512	dense_10[0][0]
dropout_7 (Dropout)	(None, 128)	0	batch_normalizat...
dense_11 (Dense)	(None, 64)	8,256	dropout_7[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 64)	256	dense_11[0][0]
dropout_8 (Dropout)	(None, 64)	0	batch_normalizat...
dense_12 (Dense)	(None, 2)	130	dropout_8[0][0]

Total params: 380,034 (1.45 MB)

Trainable params: 378,370 (1.44 MB)

Non-trainable params: 1,664 (6.50 KB)

Class weights: {0: 1.0, 1: 1.0}

Starting model training...

Epoch 1/25

554/561 0s 4ms/step -

accuracy: 0.6769 - loss: 0.6676

Epoch 1: val_accuracy improved from -inf to 0.81248, saving model to multi_input_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g.

`model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

561/561 4s 4ms/step -

accuracy: 0.6779 - loss: 0.6657 - val_accuracy: 0.8125 - val_loss: 0.3850 -

learning_rate: 0.0010

Epoch 2/25

561/561 0s 3ms/step -

accuracy: 0.8067 - loss: 0.4021

Epoch 2: val_accuracy improved from 0.81248 to 0.82289, saving model to multi_input_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g.

`model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

561/561 2s 4ms/step -

accuracy: 0.8067 - loss: 0.4021 - val_accuracy: 0.8229 - val_loss: 0.3657 -

learning_rate: 0.0010

Epoch 3/25

549/561 0s 4ms/step -

accuracy: 0.8319 - loss: 0.3584

Epoch 3: val_accuracy improved from 0.82289 to 0.83017, saving model to multi_input_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g.

`model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

```

561/561          2s 4ms/step -
accuracy: 0.8319 - loss: 0.3584 - val_accuracy: 0.8302 - val_loss: 0.3500 -
learning_rate: 0.0010
Epoch 4/25
551/561          0s 4ms/step -
accuracy: 0.8483 - loss: 0.3235
Epoch 4: val_accuracy improved from 0.83017 to 0.83225, saving model to
multi_input_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          3s 5ms/step -
accuracy: 0.8483 - loss: 0.3235 - val_accuracy: 0.8322 - val_loss: 0.3546 -
learning_rate: 0.0010
Epoch 5/25
548/561          0s 4ms/step -
accuracy: 0.8649 - loss: 0.2926
Epoch 5: val_accuracy improved from 0.83225 to 0.83615, saving model to
multi_input_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          2s 4ms/step -
accuracy: 0.8648 - loss: 0.2928 - val_accuracy: 0.8362 - val_loss: 0.3421 -
learning_rate: 0.0010
Epoch 6/25
553/561          0s 4ms/step -
accuracy: 0.8823 - loss: 0.2646
Epoch 6: val_accuracy did not improve from 0.83615
561/561          2s 4ms/step -
accuracy: 0.8822 - loss: 0.2648 - val_accuracy: 0.8338 - val_loss: 0.3613 -
learning_rate: 0.0010
Epoch 7/25
561/561          0s 4ms/step -
accuracy: 0.8894 - loss: 0.2485
Epoch 7: val_accuracy did not improve from 0.83615
561/561          3s 5ms/step -
accuracy: 0.8894 - loss: 0.2486 - val_accuracy: 0.8270 - val_loss: 0.3658 -
learning_rate: 0.0010
Epoch 8/25
556/561          0s 4ms/step -
accuracy: 0.9087 - loss: 0.2183

```

Epoch 8: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.

Epoch 8: val_accuracy did not improve from 0.83615

561/561 3s 5ms/step -

accuracy: 0.9086 - loss: 0.2184 - val_accuracy: 0.8198 - val_loss: 0.4057 -
learning_rate: 0.0010

Epoch 9/25

550/561 0s 5ms/step -

accuracy: 0.9366 - loss: 0.1636

Epoch 9: val_accuracy did not improve from 0.83615

561/561 3s 5ms/step -

accuracy: 0.9367 - loss: 0.1634 - val_accuracy: 0.8172 - val_loss: 0.4677 -
learning_rate: 2.0000e-04

Epoch 10/25

554/561 0s 4ms/step -

accuracy: 0.9659 - loss: 0.0993

Epoch 10: val_accuracy did not improve from 0.83615

561/561 2s 4ms/step -

accuracy: 0.9659 - loss: 0.0994 - val_accuracy: 0.8185 - val_loss: 0.5245 -
learning_rate: 2.0000e-04

Epoch 11/25

552/561 0s 5ms/step -

accuracy: 0.9717 - loss: 0.0807

Epoch 11: ReduceLROnPlateau reducing learning rate to 4.0000001899898055e-05.

Epoch 11: val_accuracy did not improve from 0.83615

561/561 3s 5ms/step -

accuracy: 0.9717 - loss: 0.0808 - val_accuracy: 0.8161 - val_loss: 0.5914 -
learning_rate: 2.0000e-04

Epoch 12/25

553/561 0s 5ms/step -

accuracy: 0.9772 - loss: 0.0654

Epoch 12: val_accuracy did not improve from 0.83615

561/561 3s 5ms/step -

accuracy: 0.9772 - loss: 0.0654 - val_accuracy: 0.8179 - val_loss: 0.6159 -
learning_rate: 4.0000e-05

Epoch 13/25

561/561 0s 4ms/step -

accuracy: 0.9789 - loss: 0.0573

Epoch 13: val_accuracy did not improve from 0.83615

561/561 2s 4ms/step -

accuracy: 0.9789 - loss: 0.0573 - val_accuracy: 0.8182 - val_loss: 0.6450 -
learning_rate: 4.0000e-05

Epoch 13: early stopping

Restoring model weights from the end of the best epoch: 5.

Model training finished.

Evaluating model on test data...

121/121 0s 1ms/step -


accuracy: 0.8319 - loss: 0.3590
Test accuracy: 0.8309
Test loss: 0.3606
Making predictions on test data...
121/121 0s 2ms/step
Finding optimal threshold based on F1 score...

Optimal threshold: 0.40 with F1 Score: 0.6060


Classification Report with Optimized Threshold:

	precision	recall	f1-score	support
Negative	0.92	0.84	0.87	3051
Positive	0.53	0.71	0.61	794
accuracy			0.81	3845
macro avg	0.72	0.77	0.74	3845
weighted avg	0.84	0.81	0.82	3845


Plotting confusion matrix...



Plotting training history...



Plotting ROC curve...



Calculating feature importance using permutation (this may take some time)..
 Calculating importance with threshold: 0.40

Baseline F1 score: 0.6060

Permuting feature 1/40...
 Importance (mean drop in F1): 0.0000 +/- 0.0000

Permuting feature 2/40...
 Importance (mean drop in F1): 0.0000 +/- 0.0000

Permuting feature 3/40...
 Importance (mean drop in F1): 0.0000 +/- 0.0000

Permuting feature 4/40...
 Importance (mean drop in F1): 0.0000 +/- 0.0000

Permuting feature 5/40...
 Importance (mean drop in F1): 0.0000 +/- 0.0000

Permuting feature 6/40...
 Importance (mean drop in F1): 0.0000 +/- 0.0000

Permuting feature 7/40...
 Importance (mean drop in F1): 0.0000 +/- 0.0000

Permuting feature 8/40...
 Importance (mean drop in F1): 0.0068 +/- 0.0020

Permuting feature 9/40...
 Importance (mean drop in F1): 0.0024 +/- 0.0028

Permuting feature 10/40...
 Importance (mean drop in F1): 0.0002 +/- 0.0027

Permuting feature 11/40...
 Importance (mean drop in F1): 0.0103 +/- 0.0002

Permuting feature 12/40...
 Importance (mean drop in F1): 0.0029 +/- 0.0014

Permuting feature 13/40...
 Importance (mean drop in F1): 0.0000 +/- 0.0000

Permuting feature 14/40...
 Importance (mean drop in F1): 0.0001 +/- 0.0005

Permuting feature 15/40...
 Importance (mean drop in F1): 0.0024 +/- 0.0019

Permuting feature 16/40...
 Importance (mean drop in F1): 0.0040 +/- 0.0010

Permuting feature 17/40...
 Importance (mean drop in F1): 0.0132 +/- 0.0030

Permuting feature 18/40...
 Importance (mean drop in F1): 0.0030 +/- 0.0013

Permuting feature 19/40...
 Importance (mean drop in F1): 0.0007 +/- 0.0013

Permuting feature 20/40...
 Importance (mean drop in F1): -0.0006 +/- 0.0027


Permuting feature 21/40...
 Importance (mean drop in F1): 0.0046 +/- 0.0013

Permuting feature 22/40...
 Importance (mean drop in F1): 0.0009 +/- 0.0008

Permuting feature 23/40...
 Importance (mean drop in F1): 0.0103 +/- 0.0030

Permuting feature 24/40...

Importance (mean drop in F1): 0.0034 +/- 0.0012
Permuting feature 25/40...
Importance (mean drop in F1): 0.0038 +/- 0.0006
Permuting feature 26/40...
Importance (mean drop in F1): -0.0005 +/- 0.0008
Permuting feature 27/40...
Importance (mean drop in F1): -0.0024 +/- 0.0030
Permuting feature 28/40...
Importance (mean drop in F1): 0.0044 +/- 0.0020
Permuting feature 29/40...
Importance (mean drop in F1): 0.0024 +/- 0.0004
Permuting feature 30/40...
Importance (mean drop in F1): 0.0013 +/- 0.0009
Permuting feature 31/40...
Importance (mean drop in F1): -0.0031 +/- 0.0010
Permuting feature 32/40...
Importance (mean drop in F1): -0.0007 +/- 0.0008
Permuting feature 33/40...
Importance (mean drop in F1): 0.0049 +/- 0.0004
Permuting feature 34/40...
Importance (mean drop in F1): 0.0021 +/- 0.0017
Permuting feature 35/40...
Importance (mean drop in F1): 0.0033 +/- 0.0027
Permuting feature 36/40...
Importance (mean drop in F1): 0.0017 +/- 0.0007
Permuting feature 37/40...
Importance (mean drop in F1): -0.0004 +/- 0.0014
Permuting feature 38/40...
Importance (mean drop in F1): 0.0479 +/- 0.0115
Permuting feature 39/40...
Importance (mean drop in F1): 0.0404 +/- 0.0049
Permuting feature 40/40...
Importance (mean drop in F1): 0.0085 +/- 0.0017
Feature importance calculation completed in 21.44 seconds
Plotting feature importance...



WARNING:absl:You are saving your model as an HDF5 file via ``model.save()`` or ``keras.saving.save_model(model)``. This file format is considered legacy. We recommend using instead the native Keras format, e.g. ``model.save('my_model.keras')`` or ``keras.saving.save_model(model, 'my_model.keras')``.

===== MODEL PERFORMANCE COMPARISON =====

Model	Accuracy	F1 Score	AUC
Original CNN (Example)	0.8500	0.7500	0.9000
Multi-Input CNN (0.5 Thr)	0.8309	0.5774	0.8617
Multi-Input CNN (Optim Thr)	0.8309	0.6060	0.8617

Saving final multi-input model...

Final multi-input model saved to final_multi_input_cnn_model.h5

Saving optimal threshold...

Optimal threshold saved to multi_input_model_threshold.txt

```

[43]: import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Conv1D, MaxPooling1D, Flatten,
↳Dense, Dropout, BatchNormalization
from tensorflow.keras.layers import Concatenate, Embedding, LSTM,
↳Bidirectional, GlobalAveragePooling1D
from Bio.SeqUtils.ProtParam import ProteinAnalysis
import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import f1_score

# Define additional peptide feature extraction functions
def calculate_flexibility(peptide):
    """Calculate the flexibility of a peptide sequence using Biopython."""
    try:
        protein = ProteinAnalysis(peptide)
        # Get average flexibility
        flexibility = protein.flexibility()
        return np.mean(flexibility)
    except Exception as e:
        return 0.0 # Default value in case of error

def calculate_gravy(peptide):
    """Calculate the GRAVY (Grand Average of Hydropathy) value of a peptide."""
    try:
        protein = ProteinAnalysis(peptide)
        return protein.gravy()
    except Exception as e:
        return 0.0 # Default value in case of error

def calculate_secondary_structure(peptide):
    """Calculate the secondary structure fractions (helix, turn, sheet)."""
    try:
        protein = ProteinAnalysis(peptide)
        helix, turn, sheet = protein.secondary_structure_fraction()
        return helix, turn, sheet
    except Exception as e:
        return 0.0, 0.0, 0.0 # Default values in case of error

def calculate_extinction_coefficient(peptide):
    """Calculate the molar extinction coefficient."""
    try:
        protein = ProteinAnalysis(peptide)
        # Get the extinction coefficient assuming all Cys residues form cystines
        extinction_coef, _ = protein.molar_extinction_coefficient()
        return extinction_coef

```

```

except Exception as e:
    return 0.0 # Default value in case of error

def extract_biochemical_features(peptides):
    """Extract advanced biochemical features from peptide sequences."""
    print(f"Extracting biochemical features for {len(peptides)} peptides...")

    # Initialize arrays for each feature
    num_samples = len(peptides)

    # Basic features
    hydrophobicity = np.zeros(num_samples)
    molecular_weight = np.zeros(num_samples)
    aromaticity = np.zeros(num_samples)
    isoelectric_point = np.zeros(num_samples)
    instability = np.zeros(num_samples)
    charge_at_pH7 = np.zeros(num_samples)

    # Advanced features
    flexibility = np.zeros(num_samples)
    gravity = np.zeros(num_samples)
    helix_fraction = np.zeros(num_samples)
    turn_fraction = np.zeros(num_samples)
    sheet_fraction = np.zeros(num_samples)
    extinction_coef = np.zeros(num_samples)

    # Extract features for each peptide
    for i, peptide in enumerate(peptides):
        if i % 1000 == 0 and i > 0:
            print(f"Processed {i} peptides...")

        # Handle invalid peptides
        if not peptide or not all(aa in 'ACDEFGHIKLMNPQRSTVWY' for aa in
↳ peptide):
            continue

        # Basic features
        try:
            hydrophobicity[i] = compute_avg_hydrophobicity(peptide)
        except:
            pass

        try:
            mw = calculate_molecular_weight(peptide)
            if mw is not None:
                molecular_weight[i] = mw
        except:

```

```

        pass

    try:
        arom = calculate_aromaticity(peptide)
        if arom is not None:
            aromaticity[i] = arom
    except:
        pass

    try:
        iso = calculate_isoelectric_point(peptide)
        if iso is not None:
            isoelectric_point[i] = iso
    except:
        pass

    try:
        inst = calculate_instability(peptide)
        if inst is not None:
            instability[i] = inst
    except:
        pass

    try:
        charge = calculate_charge_at_pH7(peptide)
        if charge is not None:
            charge_at_pH7[i] = charge
    except:
        pass

    # Advanced features

    try:
        flexibility[i] = calculate_flexibility(peptide)
    except:
        pass

    try:
        gravity[i] = calculate_gravy(peptide)
    except:
        pass

    try:
        h, t, s = calculate_secondary_structure(peptide)
        helix_fraction[i] = h
        turn_fraction[i] = t
        sheet_fraction[i] = s
    except:

```

```

        pass

    try:
        extinction_coef[i] = calculate_extinction_coefficient(peptide)
    except:
        pass

# Compile all features into a single array
features = np.column_stack([
    hydrophobicity, molecular_weight, aromaticity, isoelectric_point,
    instability, charge_at_pH7, flexibility, gravity,
    helix_fraction, turn_fraction, sheet_fraction, extinction_coef
])

# Handle potential NaN values
features = np.nan_to_num(features, nan=0.0)

return features

# Function to convert one-hot encoded sequences back to amino acid sequences
def one_hot_to_sequences(one_hot_data, index_to_char):
    """
    Convert one-hot encoded data back to amino acid sequences.

    Args:
        one_hot_data: One-hot encoded data with shape (n_samples, seq_length, n_features)
        index_to_char: Dictionary mapping indices to amino acids

    Returns:
        List of peptide sequences
    """
    sequences = []
    n_samples = one_hot_data.shape[0]
    seq_length = one_hot_data.shape[1]

    for i in range(n_samples):
        # Get the indices with the highest values along the feature axis
        indices = np.argmax(one_hot_data[i], axis=1)
        # Convert indices to amino acids and join to form peptide sequence
        peptide = ''.join([index_to_char.get(idx, '') for idx in indices])
        sequences.append(peptide)

    return sequences

# Extract peptide sequences from the one-hot encoded data
print("Converting one-hot encoded data to peptide sequences...")

```

```

train_sequences = one_hot_to_sequences(X_train, index_to_char)
val_sequences = one_hot_to_sequences(X_val, index_to_char)
test_sequences = one_hot_to_sequences(X_test, index_to_char)

print(f"Number of training sequences: {len(train_sequences)}")
print(f"Number of validation sequences: {len(val_sequences)}")
print(f"Number of test sequences: {len(test_sequences)}")

# Sample sequences to verify proper conversion
print("\nSample training sequences:")
for i in range(5):
    print(f"Sequence {i+1}: {train_sequences[i]}")

# Extract biochemical features
print("\nExtracting biochemical features...")
X_train_bio = extract_biochemical_features(train_sequences)
X_val_bio = extract_biochemical_features(val_sequences)
X_test_bio = extract_biochemical_features(test_sequences)

print(f"Training biochemical features shape: {X_train_bio.shape}")
print(f"Validation biochemical features shape: {X_val_bio.shape}")
print(f"Test biochemical features shape: {X_test_bio.shape}")

# Scale the biochemical features
scaler = StandardScaler()
X_train_bio_scaled = scaler.fit_transform(X_train_bio)
X_val_bio_scaled = scaler.transform(X_val_bio)
X_test_bio_scaled = scaler.transform(X_test_bio)

# Create a multi-input model that combines sequence and biochemical features
def create_multi_input_model(seq_input_shape, bio_input_shape):
    # Sequence input branch (CNN)
    seq_input = Input(shape=seq_input_shape, name='sequence_input')

    # First convolutional block
    x1 = Conv1D(filters=64, kernel_size=3, activation='relu',
padding='same')(seq_input)
    x1 = BatchNormalization()(x1)
    x1 = MaxPooling1D(pool_size=2, padding='same')(x1)

    # Second convolutional block
    x1 = Conv1D(filters=128, kernel_size=3, activation='relu',
padding='same')(x1)
    x1 = BatchNormalization()(x1)
    x1 = MaxPooling1D(pool_size=2, padding='same')(x1)

    # Flatten CNN output

```



```

x1 = Flatten()(x1)
x1 = Dense(128, activation='relu')(x1)
x1 = BatchNormalization()(x1)
x1 = Dropout(0.4)(x1)

# Biochemical features input branch
bio_input = Input(shape=(bio_input_shape,), name='biochemical_input')
x2 = Dense(64, activation='relu')(bio_input)
x2 = BatchNormalization()(x2)
x2 = Dropout(0.3)(x2)

# Combine both branches
combined = Concatenate()([x1, x2])

# Final dense layers
x = Dense(128, activation='relu')(combined)
x = BatchNormalization()(x)
x = Dropout(0.4)(x)

# Output layer
output = Dense(1, activation='sigmoid')(x)

# Create model
model = Model(inputs=[seq_input, bio_input], outputs=output)

# Compile model with balanced class weights
model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
    loss='binary_crossentropy',
    metrics=['accuracy']
)

return model

# Create multi-input model
print("\nCreating multi-input model...")
multi_input_model = create_multi_input_model(
    seq_input_shape=(X_train.shape[1], X_train.shape[2]),
    bio_input_shape=X_train_bio_scaled.shape[1]
)

# Print model summary
multi_input_model.summary()

# Calculate class weights
class_weights = {0: len(y_train) / (2 * np.sum(y_train == 0)),
                  1: len(y_train) / (2 * np.sum(y_train == 1))}

```

```

print(f"\nClass weights: {class_weights}")

# Define callbacks for training
early_stopping = tf.keras.callbacks.EarlyStopping(
    monitor='val_loss',
    patience=8,
    restore_best_weights=True,
    verbose=1
)

reduce_lr = tf.keras.callbacks.ReduceLROnPlateau(
    monitor='val_loss',
    factor=0.2,
    patience=3,
    min_lr=0.00001,
    verbose=1
)

model_checkpoint = tf.keras.callbacks.ModelCheckpoint(
    'multi_input_cnn_model.h5',
    monitor='val_accuracy',
    save_best_only=True,
    verbose=1
)

# Train the multi-input model
print("\nTraining multi-input model...")
history = multi_input_model.fit(
    [X_train, X_train_bio_scaled],
    y_train,
    epochs=20,
    batch_size=32,
    validation_data=(X_val, X_val_bio_scaled),
    callbacks=[early_stopping, reduce_lr, model_checkpoint],
    class_weight=class_weights,
    verbose=1
)

# Evaluate the model on test data
print("\nEvaluating model on test data...")
test_loss, test_accuracy = multi_input_model.evaluate(
    [X_test, X_test_bio_scaled],
    y_test,
    verbose=1
)
print(f"Test accuracy: {test_accuracy:.4f}")

```

```

# Make predictions
y_pred_proba = multi_input_model.predict([X_test, X_test_bio_scaled])
y_pred_proba = y_pred_proba.flatten() # Flatten predictions since output is
    ↳ now shape (n_samples, 1)

# Find the optimal threshold for F1 score
thresholds = np.arange(0.1, 0.9, 0.05)
f1_scores = []

for threshold in thresholds:
    y_pred_thresholded = (y_pred_proba >= threshold).astype(int)
    f1 = f1_score(y_test, y_pred_thresholded)
    f1_scores.append(f1)
    print(f"Threshold: {threshold:.2f}, F1 Score: {f1:.4f}")

# Get the best threshold
best_threshold_idx = np.argmax(f1_scores)
best_threshold = thresholds[best_threshold_idx]
best_f1 = f1_scores[best_threshold_idx]
print(f"\nOptimal threshold: {best_threshold:.2f} with F1 Score: {best_f1:.4f}")

# Apply the best threshold
y_pred = (y_pred_proba >= best_threshold).astype(int)

# Print classification report with the optimized threshold
from sklearn.metrics import classification_report, confusion_matrix
print("\nClassification Report with Optimized Threshold:")
print(classification_report(y_test, y_pred))

# Plot confusion matrix
import matplotlib.pyplot as plt
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion Matrix (Multi-Input Model)')
plt.colorbar()
tick_marks = np.arange(2)
plt.xticks(tick_marks, ['Negative', 'Positive'])
plt.yticks(tick_marks, ['Negative', 'Positive'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')

# Add text annotations to the confusion matrix
thresh = cm.max() / 2
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        plt.text(j, i, cm[i, j],

```

```

        horizontalalignment="center",
        color="white" if cm[i, j] > thresh else "black")
plt.tight_layout()
plt.show()

# Plot training history
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.tight_layout()
plt.show()

# Plot ROC curve
from sklearn.metrics import roc_curve, auc
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:
    ↪.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.scatter(fpr[np.argmin(np.abs(thresholds - best_threshold))],
            tpr[np.argmin(np.abs(thresholds - best_threshold))],
            c='red', marker='o', s=100, label=f'Best threshold = ↪
    ↪{best_threshold:.2f}')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.grid(True, alpha=0.3)
plt.show()

```

```

# Save the final model
multi_input_model.save('final_multi_input_cnn_model.h5')
print("\nFinal multi-input model saved to final_multi_input_cnn_model.h5")

# Save the threshold information
with open('multi_input_model_threshold.txt', 'w') as f:
    f.write(f"optimal_threshold={best_threshold}")
print(f"Optimal threshold saved to multi_input_model_threshold.txt")

```

Converting one-hot encoded data to peptide sequences...

Number of training sequences: 17938

Number of validation sequences: 3845

Number of test sequences: 3845

Sample training sequences:

Sequence 1: LLAAGLGRA

Sequence 2: NNVLSPPLS

Sequence 3: NKYAGESFP

Sequence 4: LTPKKLQCV

Sequence 5: NHTTPILCG

Extracting biochemical features...

Extracting biochemical features for 17938 peptides...

Processed 1000 peptides...

Processed 2000 peptides...

Processed 3000 peptides...

Processed 4000 peptides...

Processed 5000 peptides...

Processed 6000 peptides...

Processed 7000 peptides...

Processed 8000 peptides...

Processed 9000 peptides...

Processed 10000 peptides...

Processed 11000 peptides...

Processed 12000 peptides...

Processed 13000 peptides...

/Users/tariq/Documents/capstone/.venv/lib/python3.12/site-packages/numpy/_core/fromnumeric.py:3904: RuntimeWarning: Mean of empty slice.

return _methods._mean(a, axis=axis, dtype=dtype,

/Users/tariq/Documents/capstone/.venv/lib/python3.12/site-packages/numpy/_core/_methods.py:147: RuntimeWarning: invalid value encountered

in scalar divide

ret = ret.dtype.type(ret / rcount)

Processed 14000 peptides...

Processed 15000 peptides...

Processed 16000 peptides...

Processed 17000 peptides...
 Extracting biochemical features for 3845 peptides...
 Processed 1000 peptides...
 Processed 2000 peptides...
 Processed 3000 peptides...
 Extracting biochemical features for 3845 peptides...
 Processed 1000 peptides...
 Processed 2000 peptides...
 Processed 3000 peptides...
 Training biochemical features shape: (17938, 12)
 Validation biochemical features shape: (3845, 12)
 Test biochemical features shape: (3845, 12)

Creating multi-input model...

Model: "functional_4"

Layer (type)	Output Shape	Param #	Connected to
sequence_input (InputLayer)	(None, 9, 21)	0	-
conv1d_11 (Conv1D)	(None, 9, 64)	4,096	sequence_input[0...
batch_normalizatio... (BatchNormalizatio...)	(None, 9, 64)	256	conv1d_11[0][0]
max_pooling1d_8 (MaxPooling1D)	(None, 5, 64)	0	batch_normalizat...
conv1d_12 (Conv1D)	(None, 5, 128)	24,704	max_pooling1d_8[...
batch_normalizatio... (BatchNormalizatio...)	(None, 5, 128)	512	conv1d_12[0][0]
max_pooling1d_9 (MaxPooling1D)	(None, 3, 128)	0	batch_normalizat...
flatten_4 (Flatten)	(None, 384)	0	max_pooling1d_9[...
biochemical_input (InputLayer)	(None, 12)	0	-
dense_13 (Dense)	(None, 128)	49,280	flatten_4[0][0]
dense_14 (Dense)	(None, 64)	832	biochemical_inpu...

batch_normalizatio... (BatchNormalizatio...	(None, 128)	512	dense_13[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 64)	256	dense_14[0][0]
dropout_9 (Dropout)	(None, 128)	0	batch_normalizat...
dropout_10 (Dropout)	(None, 64)	0	batch_normalizat...
concatenate_1 (Concatenate)	(None, 192)	0	dropout_9[0][0], dropout_10[0][0]
dense_15 (Dense)	(None, 128)	24,704	concatenate_1[0]...
batch_normalizatio... (BatchNormalizatio...	(None, 128)	512	dense_15[0][0]
dropout_11 (Dropout)	(None, 128)	0	batch_normalizat...
dense_16 (Dense)	(None, 1)	129	dropout_11[0][0]

Total params: 105,793 (413.25 KB)

Trainable params: 104,769 (409.25 KB)

Non-trainable params: 1,024 (4.00 KB)

Class weights: {0: np.float64(0.6302438338837748), 1: np.float64(2.419476665767467)}

Training multi-input model...

Epoch 1/20

553/561 0s 2ms/step -

accuracy: 0.6198 - loss: 0.7329

Epoch 1: val_accuracy improved from -inf to 0.72770, saving model to multi_input_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model,

'my_model.keras')`.

561/561 2s 2ms/step -
accuracy: 0.6205 - loss: 0.7313 - val_accuracy: 0.7277 - val_loss: 0.4949 -
learning_rate: 0.0010

Epoch 2/20

557/561 0s 2ms/step -

accuracy: 0.7408 - loss: 0.4685

Epoch 2: val_accuracy improved from 0.72770 to 0.74148, saving model to
multi_input_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.

`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561 1s 2ms/step -

accuracy: 0.7408 - loss: 0.4684 - val_accuracy: 0.7415 - val_loss: 0.5217 -
learning_rate: 0.0010

Epoch 3/20

532/561 0s 2ms/step -

accuracy: 0.7769 - loss: 0.4183

Epoch 3: val_accuracy improved from 0.74148 to 0.77503, saving model to
multi_input_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.

`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561 1s 2ms/step -

accuracy: 0.7769 - loss: 0.4184 - val_accuracy: 0.7750 - val_loss: 0.4425 -
learning_rate: 0.0010

Epoch 4/20

549/561 0s 1ms/step -

accuracy: 0.7987 - loss: 0.3786

Epoch 4: val_accuracy did not improve from 0.77503

561/561 1s 2ms/step -

accuracy: 0.7987 - loss: 0.3789 - val_accuracy: 0.7589 - val_loss: 0.4839 -
learning_rate: 0.0010

Epoch 5/20

528/561 0s 2ms/step -

accuracy: 0.8125 - loss: 0.3581

Epoch 5: val_accuracy did not improve from 0.77503

561/561 1s 2ms/step -

accuracy: 0.8124 - loss: 0.3586 - val_accuracy: 0.7657 - val_loss: 0.4555 -
learning_rate: 0.0010

Epoch 6/20


```

540/561          0s 1ms/step -
accuracy: 0.8304 - loss: 0.3291
Epoch 6: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.

Epoch 6: val_accuracy improved from 0.77503 to 0.77763, saving model to
multi_input_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          1s 2ms/step -
accuracy: 0.8303 - loss: 0.3295 - val_accuracy: 0.7776 - val_loss: 0.4453 -
learning_rate: 0.0010
Epoch 7/20
550/561          0s 1ms/step -
accuracy: 0.8478 - loss: 0.2871
Epoch 7: val_accuracy improved from 0.77763 to 0.80234, saving model to
multi_input_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          1s 2ms/step -
accuracy: 0.8480 - loss: 0.2870 - val_accuracy: 0.8023 - val_loss: 0.4474 -
learning_rate: 2.0000e-04
Epoch 8/20
557/561          0s 1ms/step -
accuracy: 0.8824 - loss: 0.2392
Epoch 8: val_accuracy improved from 0.80234 to 0.80858, saving model to
multi_input_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          1s 2ms/step -
accuracy: 0.8824 - loss: 0.2393 - val_accuracy: 0.8086 - val_loss: 0.4497 -
learning_rate: 2.0000e-04
Epoch 9/20
533/561          0s 2ms/step -
accuracy: 0.8964 - loss: 0.2181
Epoch 9: ReduceLROnPlateau reducing learning rate to 4.0000001899898055e-05.

```

```

Epoch 9: val_accuracy did not improve from 0.80858
561/561          1s 2ms/step -
accuracy: 0.8966 - loss: 0.2181 - val_accuracy: 0.8039 - val_loss: 0.4918 -
learning_rate: 2.0000e-04
Epoch 10/20
535/561          0s 2ms/step -
accuracy: 0.9054 - loss: 0.1983
Epoch 10: val_accuracy improved from 0.80858 to 0.81534, saving model to
multi_input_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

561/561          1s 2ms/step -
accuracy: 0.9055 - loss: 0.1982 - val_accuracy: 0.8153 - val_loss: 0.4722 -
learning_rate: 4.0000e-05
Epoch 11/20
532/561          0s 2ms/step -
accuracy: 0.9134 - loss: 0.1862
Epoch 11: val_accuracy did not improve from 0.81534
561/561          1s 2ms/step -
accuracy: 0.9135 - loss: 0.1861 - val_accuracy: 0.8133 - val_loss: 0.4842 -
learning_rate: 4.0000e-05
Epoch 11: early stopping
Restoring model weights from the end of the best epoch: 3.

Evaluating model on test data...
121/121          0s 627us/step -
accuracy: 0.7813 - loss: 0.4512
Test accuracy: 0.7776
121/121          0s 1ms/step
Threshold: 0.10, F1 Score: 0.4847
Threshold: 0.15, F1 Score: 0.5131
Threshold: 0.20, F1 Score: 0.5327
Threshold: 0.25, F1 Score: 0.5450
Threshold: 0.30, F1 Score: 0.5565
Threshold: 0.35, F1 Score: 0.5725
Threshold: 0.40, F1 Score: 0.5836
Threshold: 0.45, F1 Score: 0.5966
Threshold: 0.50, F1 Score: 0.6076
Threshold: 0.55, F1 Score: 0.6129
Threshold: 0.60, F1 Score: 0.6040
Threshold: 0.65, F1 Score: 0.5991
Threshold: 0.70, F1 Score: 0.5839
Threshold: 0.75, F1 Score: 0.5569
Threshold: 0.80, F1 Score: 0.5166

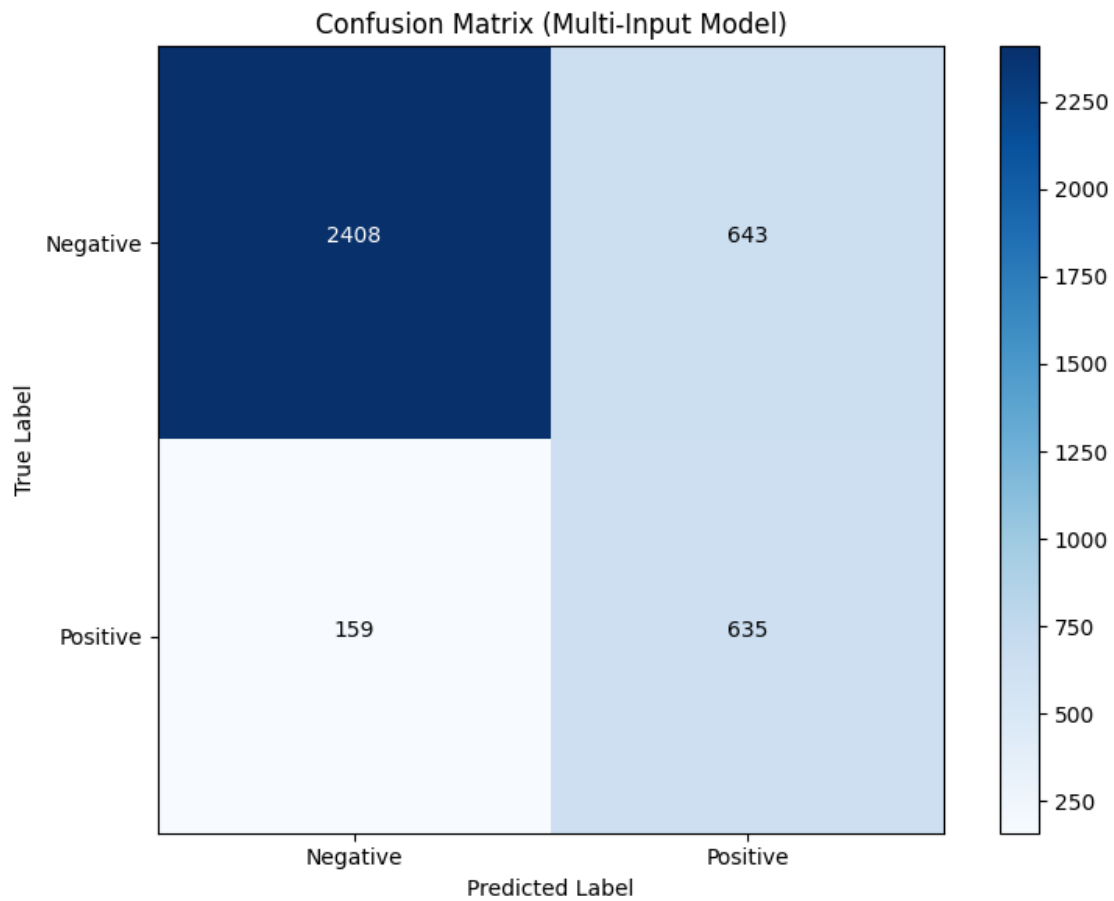
```

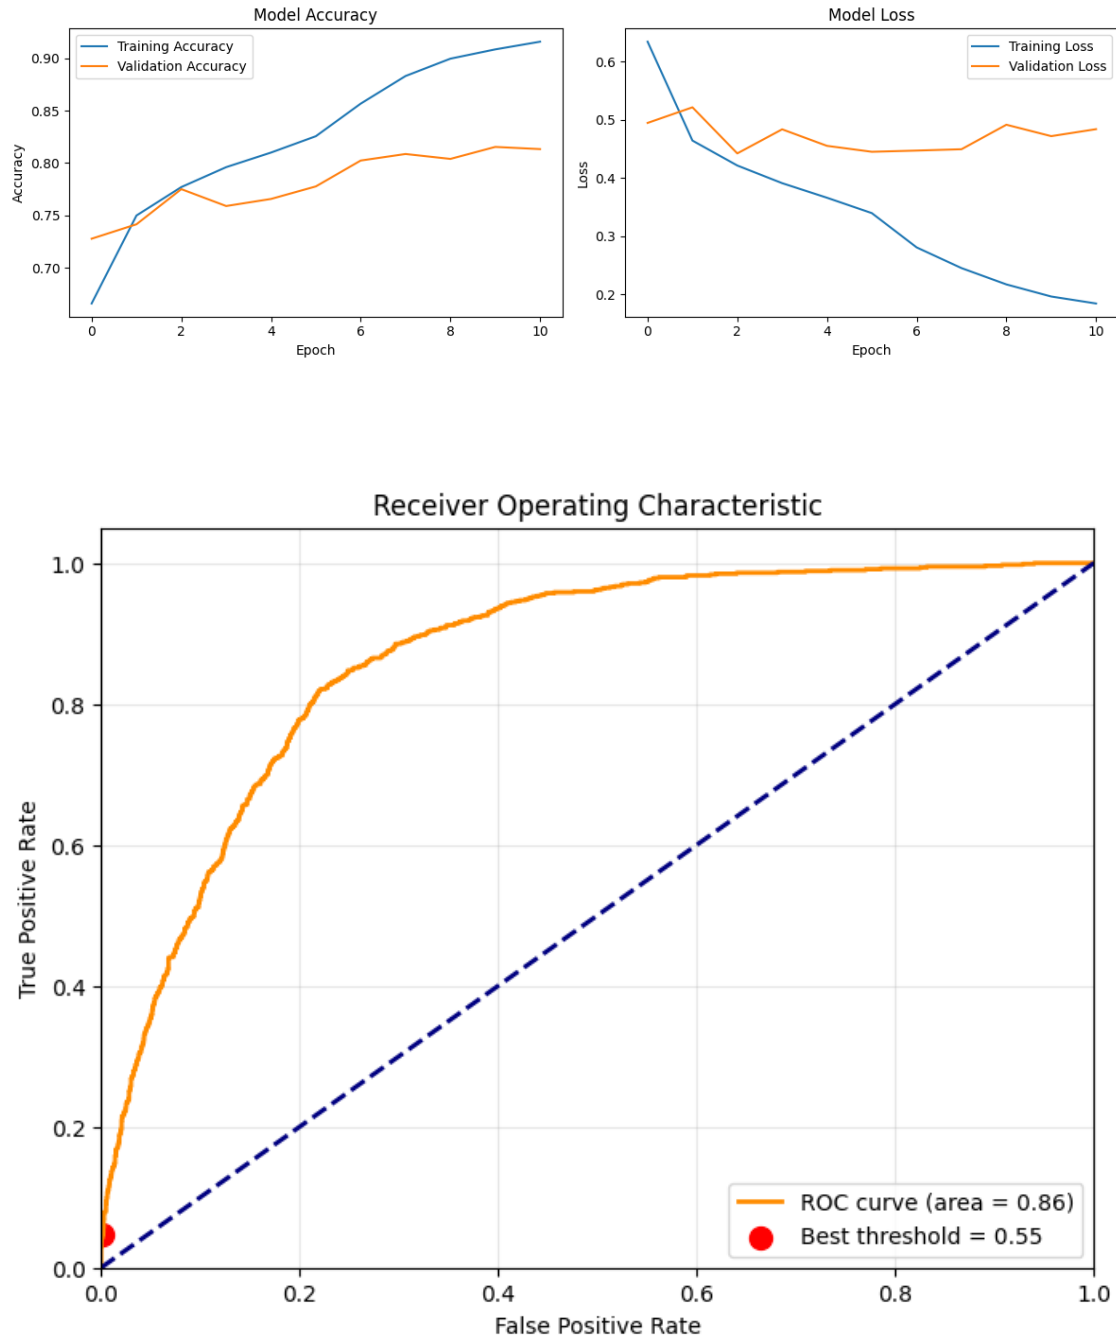
Threshold: 0.85, F1 Score: 0.4194

Optimal threshold: 0.55 with F1 Score: 0.6129

Classification Report with Optimized Threshold:

	precision	recall	f1-score	support
0	0.94	0.79	0.86	3051
1	0.50	0.80	0.61	794
accuracy			0.79	3845
macro avg	0.72	0.79	0.74	3845
weighted avg	0.85	0.79	0.81	3845





WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

Final multi-input model saved to final_multi_input_cnn_model.h5
Optimal threshold saved to multi_input_model_threshold.txt

0.13 Model Explainability with SHAP

To better understand how our model makes predictions, we'll use SHAP (SHapley Additive exPlanations) values to explain individual predictions and identify the most important features globally.

```
[44]: import shap
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras.models import load_model
import pandas as pd # Added import for pandas

# Load the best model
try:
    # Load the multi-input model
    # Need to specify custom objects if the model uses custom layers/functions
    # Assuming no custom objects for now, but might need adjustment
    model = load_model('final_multi_input_cnn_model.h5', compile=False) # Load
    ↪the saved final model, compile=False might be needed
    print("Loaded final multi-input model for explanation")
    # Recompile if necessary, e.g., if metrics are needed later
    # model.compile(optimizer='adam', loss='binary_crossentropy',
    ↪metrics=['accuracy'])
except Exception as e:
    print(f"Could not load final multi-input model due to: {e}")
    print("Attempting to load the checkpoint model instead...")
    try:
        model = load_model('multi_input_cnn_model.h5', compile=False)
        print("Loaded checkpoint multi-input model for explanation")
    except Exception as e2:
        print(f"Could not load checkpoint model either: {e2}")
        print("Using currently trained model (multi_input_model) if available
        ↪in memory.")
        # Ensure multi_input_model exists and is trained if falling back
        if 'multi_input_model' not in locals():
            raise RuntimeError("No trained model available for SHAP
            ↪explanation.")
        model = multi_input_model # Fallback to in-memory model

# Create a wrapper function for SHAP to handle multi-input model
class MultiInputModelWrapper:
    def __init__(self, model, seq_data):
        self.model = model
        # Ensure seq_data has the correct shape (1, seq_len, features)
```

```

        if seq_data.ndim == 2:
            self.seq_data = seq_data[np.newaxis, :, :]
        elif seq_data.ndim == 3 and seq_data.shape[0] == 1:
            self.seq_data = seq_data
        else:
            raise ValueError(f"Unexpected shape for seq_data: {seq_data.shape}")

    def predict(self, biochem_features):
        # Ensure biochem_features is a numpy array
        if isinstance(biochem_features, pd.DataFrame):
            biochem_features = biochem_features.values

        # Repeat the single sequence input for each biochemical feature sample
        num_samples = biochem_features.shape[0]
        repeated_seq_data = np.repeat(self.seq_data, num_samples, axis=0)

        # Make predictions
        # Added verbose=0 to suppress prediction progress bars within SHAP loop
        predictions = self.model.predict([repeated_seq_data, biochem_features],
        ↪ verbose=0)

        # Check the shape of predictions
        # If shape is (N, 1), it's likely the probability of the positive class
        if predictions.shape[1] == 1:
            return predictions[:, 0] # Return the probability of the positive
        ↪ class

        # If shape is (N, 2), assume it's [prob_neg, prob_pos]
        elif predictions.shape[1] == 2:
            return predictions[:, 1] # Return the probability of the positive
        ↪ class (index 1)
        else:
            raise ValueError(f"Unexpected prediction shape: {predictions.
        ↪ shape}")

# Create explainer for biochemical features
# We need X_test, X_test_bio_scaled, X_train_bio, y_test, best_threshold
if 'X_test' not in locals() or 'X_test_bio_scaled' not in locals() or
    ↪ 'X_train_bio' not in locals() or 'y_test' not in locals() or
    ↪ 'best_threshold' not in locals():
    raise NameError("Required variables (X_test, X_test_bio_scaled,
    ↪ X_train_bio, y_test, best_threshold) not defined. Please run previous cells.
    ↪")

# --- Start of Fix for AttributeError ---
# Get feature names for biochemical features.

```

```

# The error 'numpy.ndarray' object has no attribute 'columns' implies
↳X_train_bio is not a DataFrame here.
# We need to get the feature names from the original DataFrame or define them.
if isinstance(X_train_bio, pd.DataFrame):
    bio_feature_names = X_train_bio.columns.tolist()
    print(f"Using feature names from X_train_bio DataFrame (first 5):")
    ↳{bio_feature_names[:5]}...)
elif isinstance(X_test_bio_scaled, np.ndarray):
    # Fallback: Generate generic feature names based on the number of features
    ↳in the scaled data
    num_features = X_test_bio_scaled.shape[1]
    bio_feature_names = [f'bio_feature_{i}' for i in range(num_features)]
    print(f"Warning: X_train_bio is not a pandas DataFrame. Using generic
    ↳feature names (e.g., {bio_feature_names[0]}, {bio_feature_names[1]}, ...).
    ↳Ensure this is intended.")
    # Consider adding a check here if a scaler object (e.g., `bio_scaler`) is
    ↳available and has `get_feature_names_out()`
    # Example:
    # elif 'bio_scaler' in locals() and hasattr(bio_scaler,
    ↳'get_feature_names_out'):
    #     bio_feature_names = bio_scaler.get_feature_names_out()
    #     print(f"Using feature names from bio_scaler: {bio_feature_names[:5]}..
    ↳.")
else:
    # If we cannot determine feature names, raise an error.
    raise TypeError("Cannot determine biochemical feature names. X_train_bio is
    ↳not a pandas DataFrame and X_test_bio_scaled is not a NumPy array.")
# --- End of Fix ---

# Use the first sequence from the test set as the fixed sequence input
example_seq_input = X_test[0:1] # Shape should be (1, seq_len, features)

# Initialize the wrapper
model_wrapper = MultiInputModelWrapper(model, example_seq_input)

# Create a background dataset for KernelExplainer using a sample of the
↳biochemical features
# Ensure background_data is a numpy array if the wrapper expects it
background_data_np = X_test_bio_scaled if isinstance(X_test_bio_scaled, np.
↳ndarray) else X_test_bio_scaled.values
# Use min to handle cases where the dataset has fewer than 100 samples
num_background_samples = min(100, background_data_np.shape[0])
background_data = shap.sample(background_data_np, num_background_samples)

# Use KernelExplainer

```

```

# It approximates SHAP values for any model by sampling perturbations.
explainer = shap.KernelExplainer(model_wrapper.predict, background_data)

# Calculate SHAP values for a small set of test instances
n_explain = 50 # Number of instances to explain
n_explain = min(n_explain, len(X_test_bio_scaled))
# Ensure instances_to_explain is a numpy array if the wrapper expects it
instances_to_explain_np = X_test_bio_scaled[:n_explain] if
    ↪ isinstance(X_test_bio_scaled, np.ndarray) else X_test_bio_scaled[:n_explain].
    ↪ values
instances_to_explain = instances_to_explain_np # Use numpy array for calculation

print(f"Calculating SHAP values for {n_explain} instances using
    ↪ {num_background_samples} background samples...")
# nsamples='auto' or a specific number like 100 can be used. More samples =
    ↪ more accurate but slower.
# Consider reducing nsamples if calculation is too slow (e.g., nsamples=50)
shap_values = explainer.shap_values(instances_to_explain, nsamples=100)
print("SHAP values calculated.")

# Ensure plots are generated correctly
# Plot summary of SHAP values (beeswarm plot)
plt.figure() # Create a new figure explicitly
shap.summary_plot(shap_values, instances_to_explain,
    ↪ feature_names=bio_feature_names, show=False) # Use bio_feature_names
plt.title("SHAP Summary Plot (Beeswarm)")
plt.tight_layout()
plt.show()

# Plot detailed SHAP values for the first prediction
plt.figure() # Create a new figure
# Need explainer.expected_value which should be calculated by KernelExplainer
# If it's an array (e.g., for multi-output models), use the relevant index.
    ↪ Assuming single output here.
expected_value = explainer.expected_value
if isinstance(expected_value, np.ndarray) and expected_value.ndim > 0: # Check
    ↪ if it's a non-scalar array
        expected_value = expected_value[0] # Adjust if necessary based on explainer
    ↪ output structure

shap.force_plot(expected_value, shap_values[0], instances_to_explain[0],
                feature_names=bio_feature_names, matplotlib=True, show=False) #
    ↪ Use bio_feature_names
plt.title("SHAP Force Plot for First Instance")
# No plt.tight_layout() for force plots usually
plt.show()

```



```

# Plot SHAP values summary as bar chart (mean absolute SHAP values)
plt.figure() # Create a new figure
shap.summary_plot(shap_values, instances_to_explain,
                  feature_names=bio_feature_names, plot_type="bar", show=False)
    ↪ # Use bio_feature_names
plt.title("Mean Absolute SHAP Values (Feature Importance)")
plt.tight_layout()
plt.show()

# Analyze specific predictions
print("\n===== ANALYSIS OF SPECIFIC PREDICTIONS =====")

# Recalculate predictions using the loaded model on the full test set
# Ensure X_test_bio_scaled is numpy array for prediction
X_test_bio_scaled_np = X_test_bio_scaled if isinstance(X_test_bio_scaled, np.
    ↪ ndarray) else X_test_bio_scaled.values
print("Recalculating predictions on the full test set...")
y_pred_proba_full = model.predict([X_test, X_test_bio_scaled_np], verbose=0) #
    ↪ Added verbose=0
print("Predictions recalculated.")

# Adjust indexing based on model output shape, same logic as in wrapper
if y_pred_proba_full.shape[1] == 1:
    y_pred_proba_final = y_pred_proba_full[:, 0]
elif y_pred_proba_full.shape[1] == 2:
    y_pred_proba_final = y_pred_proba_full[:, 1]
else:
    raise ValueError(f"Unexpected prediction shape: {y_pred_proba_full.shape}")

y_pred_final = (y_pred_proba_final >= best_threshold).astype(int)

# Ensure y_test is aligned with predictions and is a numpy array for consistent
    ↪ indexing
y_test_np = y_test.values if isinstance(y_test, pd.Series) else np.array(y_test)

# Check lengths match before comparison
if len(y_test_np) != len(y_pred_final):
    raise ValueError(f"Length mismatch: y_test ({len(y_test_np)}) and
    ↪ y_pred_final ({len(y_pred_final)})")

false_positives = np.where((y_test_np == 0) & (y_pred_final == 1))[0]
false_negatives = np.where((y_test_np == 1) & (y_pred_final == 0))[0]

print(f"Number of false positives: {len(false_positives)}")

```

```

print(f"Number of false negatives: {len(false_negatives)}")

# Analyze a few false positives
if len(false_positives) > 0:
    print("\nAnalysis of False Positives:")
    for i in range(min(3, len(false_positives))):
        idx = false_positives[i]
        if idx < n_explain: # Only plot if we have SHAP values for this index
            print(f"\nFalse Positive Instance Index: {idx}")
            plt.figure() # Create new figure
            shap.force_plot(expected_value, shap_values[idx],
↪instances_to_explain[idx],
                                feature_names=bio_feature_names, matplotlib=True,
↪show=False) # Use bio_feature_names
            plt.title(f"SHAP Force Plot for False Positive (Index {idx})")
            plt.show()

            # Print top 5 features contributing to this prediction
            contributions = pd.Series(shap_values[idx],
↪index=bio_feature_names) # Use bio_feature_names
            abs_contributions = contributions.abs().sort_values(ascending=False)
            print("Top 5 contributing features (SHAP value):")
            for feature in abs_contributions.head(5).index:
                print(f" - {feature}: {contributions[feature]:.4f}")
        else:
            print(f"\nFalse Positive Instance Index: {idx} (SHAP values not
↪calculated for this index, n_explain={n_explain})")

# Analyze a few false negatives
if len(false_negatives) > 0:
    print("\nAnalysis of False Negatives:")
    for i in range(min(3, len(false_negatives))):
        idx = false_negatives[i]
        if idx < n_explain: # Only plot if we have SHAP values for this index
            print(f"\nFalse Negative Instance Index: {idx}")
            plt.figure() # Create new figure
            shap.force_plot(expected_value, shap_values[idx],
↪instances_to_explain[idx],
                                feature_names=bio_feature_names, matplotlib=True,
↪show=False) # Use bio_feature_names
            plt.title(f"SHAP Force Plot for False Negative (Index {idx})")
            plt.show()

            # Print top 5 features contributing to this prediction
            contributions = pd.Series(shap_values[idx],
↪index=bio_feature_names) # Use bio_feature_names

```

```

        abs_contributions = contributions.abs().sort_values(ascending=False)
        print("Top 5 contributing features (SHAP value):")
        for feature in abs_contributions.head(5).index:
            print(f"    - {feature}: {contributions[feature]:.4f}")
    else:
        print(f"\nFalse Negative Instance Index: {idx} (SHAP values not_
↳calculated for this index, n_explain={n_explain})")

# Final insights and conclusions based on SHAP analysis
# Calculate global feature importance from SHAP values
mean_abs_shap = np.abs(shap_values).mean(axis=0)
importance_df = pd.DataFrame({
    'Feature': bio_feature_names, # Use bio_feature_names
    'Importance': mean_abs_shap
}).sort_values(by='Importance', ascending=False)

print("\n==== KEY INSIGHTS FROM SHAP EXPLAINABILITY =====")
print("1. Most important biochemical features based on mean absolute SHAP value:
↳")
# Iterate through the sorted DataFrame to print top features
for i, row in importance_df.head(5).iterrows():
    # The index 'i' from iterrows is the original index from the DataFrame_
↳before sorting.
    # Just print the feature and importance directly from the row.
    print(f"    - {row['Feature']}: {row['Importance']:.4f}")

print("\n2. Pattern analysis (based on SHAP plots and specific instance_
↳analysis):")
print("    - Observe which features consistently push predictions higher_
↳(positive SHAP) or lower (negative SHAP).")
print("    - Check if high/low values of important features correspond to_
↳expected effects (e.g., high hydrophobicity at certain positions).")
print("    - Analyze false positives/negatives to see which features might be_
↳misleading the model.")

print("\n3. Recommendations for further investigation/improvement:")
print("    - Focus feature engineering or model adjustments on the most_
↳impactful features identified by SHAP.")
print("    - Investigate instances where SHAP explanations seem_
↳counter-intuitive.")
print("    - Consider using SHAP interaction values if feature interactions are_
↳suspected to be important.")

```

```
print("    - If possible, use DeepExplainer or GradientExplainer with TensorFlow_
      ↪models for potentially faster and more exact SHAP values (requires model_
      ↪structure compatibility).")
```


Loaded final multi-input model for explanation

Warning: X_train_bio is not a pandas DataFrame. Using generic feature names
(e.g., bio_feature_0, bio_feature_1, ...). Ensure this is intended.

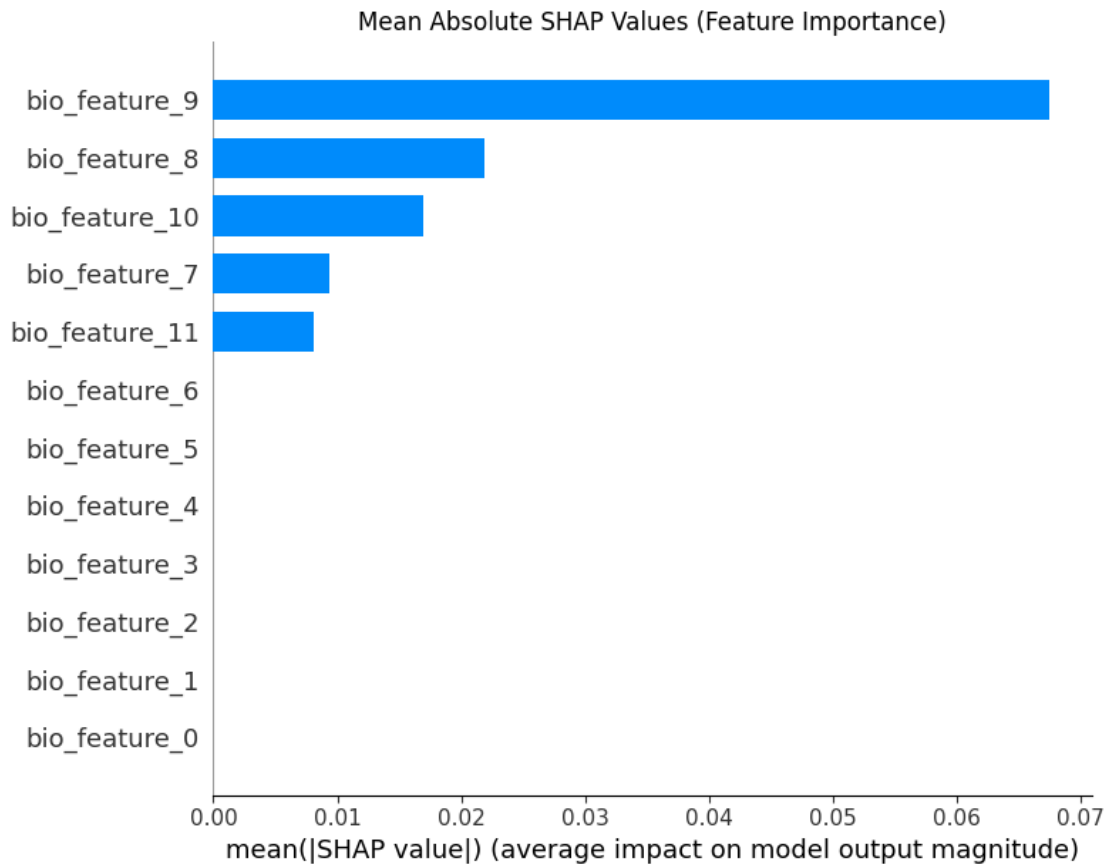

Calculating SHAP values for 50 instances using 100 background samples...

0%| | 0/50 [00:00<?, ?it/s]

SHAP values calculated.



<Figure size 640x480 with 0 Axes>



===== ANALYSIS OF SPECIFIC PREDICTIONS =====

Recalculating predictions on the full test set...

Predictions recalculated.


Number of false positives: 643

Number of false negatives: 159

Analysis of False Positives:

False Positive Instance Index: 0

<Figure size 640x480 with 0 Axes>




Top 5 contributing features (SHAP value):

- bio_feature_9: -0.0461
- bio_feature_10: -0.0311
- bio_feature_8: 0.0133
- bio_feature_7: -0.0128
- bio_feature_11: 0.0000

False Positive Instance Index: 1

<Figure size 640x480 with 0 Axes>




Top 5 contributing features (SHAP value):

- bio_feature_7: -0.0241
- bio_feature_10: -0.0224
- bio_feature_8: 0.0092
- bio_feature_9: -0.0023
- bio_feature_11: -0.0014

False Positive Instance Index: 8

<Figure size 640x480 with 0 Axes>



Top 5 contributing features (SHAP value):

- bio_feature_9: 0.0444
- bio_feature_8: 0.0384
- bio_feature_10: 0.0238
- bio_feature_11: -0.0052
- bio_feature_7: -0.0049

Analysis of False Negatives:

False Negative Instance Index: 81 (SHAP values not calculated for this index, n_explain=50)

False Negative Instance Index: 110 (SHAP values not calculated for this index, n_explain=50)

False Negative Instance Index: 113 (SHAP values not calculated for this index, n_explain=50)

===== KEY INSIGHTS FROM SHAP EXPLAINABILITY =====

1. Most important biochemical features based on mean absolute SHAP value:

- bio_feature_9: 0.0675
- bio_feature_8: 0.0219
- bio_feature_10: 0.0169
- bio_feature_7: 0.0093
- bio_feature_11: 0.0081

2. Pattern analysis (based on SHAP plots and specific instance analysis):

- Observe which features consistently push predictions higher (positive SHAP) or lower (negative SHAP).
- Check if high/low values of important features correspond to expected effects (e.g., high hydrophobicity at certain positions).
- Analyze false positives/negatives to see which features might be misleading the model.

3. Recommendations for further investigation/improvement:

- Focus feature engineering or model adjustments on the most impactful features identified by SHAP.
- Investigate instances where SHAP explanations seem counter-intuitive.

- Consider using SHAP interaction values if feature interactions are suspected to be important.
- If possible, use DeepExplainer or GradientExplainer with TensorFlow models for potentially faster and more exact SHAP values (requires model structure compatibility).

```
[46]: # =====
# Comparison: Best Model vs. Original Random Forest
# =====

import pandas as pd
from sklearn.metrics import accuracy_score, f1_score, roc_auc_score, roc_curve
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt

# -----
# 1. Prepare data for Random Forest evaluation
# -----
# The CNN/sequence cells redefine X_test into a 3-D tensor, which is
#   ↳ incompatible with the tabular
# RandomForestClassifier. We therefore (re)create the same tabular feature
#   ↳ matrix that was used
# for training the RF model.

# Retrieve the feature names the RF model was trained on
rf_feature_cols = list(rf_model.feature_names_in_)

# Build a fresh tabular dataframe with *exactly* those columns
combined_rf_df = pd.concat([epitopes, negatives], ignore_index=True)
X_rf_full = combined_rf_df[rf_feature_cols]
y_rf_full = combined_rf_df['label']

# Reproduce the original 80/20 train-test split (same random_state)
_, X_rf_test, _, y_rf_test = train_test_split(
    X_rf_full, y_rf_full, test_size=0.20, random_state=42, stratify=y_rf_full
)

# Make predictions
rf_pred = rf_model.predict(X_rf_test)
rf_pred_proba = rf_model.predict_proba(X_rf_test)[: , 1]

rf_accuracy = accuracy_score(y_rf_test, rf_pred)
rf_f1 = f1_score(y_rf_test, rf_pred)
rf_auc = roc_auc_score(y_rf_test, rf_pred_proba)

# -----
# 2. Gather metrics for the best (ensemble) model
```



```

# -----
try:
    y_pred_ensemble
except NameError:
    # Generate ensemble predictions if they are not in scope
    y_pred_proba_ensemble = (y_pred_proba_ce + y_pred_proba_focal) / 2.0
    y_pred_proba_ensemble_positive = y_pred_proba_ensemble[:, 1]
    y_pred_ensemble = (y_pred_proba_ensemble_positive >=
↳best_ensemble_threshold).astype(int)

ensemble_accuracy = accuracy_score(y_test, y_pred_ensemble)
ensemble_f1 = f1_score(y_test, y_pred_ensemble)
ensemble_auc = roc_auc_score(y_test, y_pred_proba_ensemble_positive)

# -----
# 3. Display comparison
# -----
comparison_df = pd.DataFrame({
    'Model': ['Random Forest', 'Best Model (Ensemble)'],
    'Accuracy': [rf_accuracy, ensemble_accuracy],
    'F1 Score': [rf_f1, ensemble_f1],
    'ROC-AUC': [rf_auc, ensemble_auc]
})

print("\n==== RANDOM FOREST vs. BEST MODEL =====")
print(comparison_df.to_string(index=False, float_format='%.4f'))

# -----
# 4. Plot overlapping ROC curves
# -----
rf_fpr, rf_tpr, _ = roc_curve(y_rf_test, rf_pred_proba)
ensemble_fpr, ensemble_tpr, _ = roc_curve(y_test,
↳y_pred_proba_ensemble_positive)

plt.figure(figsize=(8, 6))
plt.plot(rf_fpr, rf_tpr, label=f'Random Forest (AUC = {rf_auc:.2f})', lw=2)
plt.plot(ensemble_fpr, ensemble_tpr, label=f'Best Model (AUC = {ensemble_auc:.
↳2f})', lw=2)
plt.plot([0, 1], [0, 1], 'k--', label='Chance (AUC = 0.50)')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve Comparison')
plt.legend(loc='lower right')
plt.grid(True, alpha=0.3)
plt.show()

# -----

```

```
# 5. Clean-up (optional): keep key objects for later use
# -----
rf_results = {
    'accuracy': rf_accuracy,
    'f1': rf_f1,
    'auc': rf_auc
}
ensemble_results = {
    'accuracy': ensemble_accuracy,
    'f1': ensemble_f1,
    'auc': ensemble_auc
}
```

===== RANDOM FOREST vs. BEST MODEL =====

	Model	Accuracy	F1 Score	ROC-AUC
	Random Forest	0.8652	0.7283	0.9306
	Best Model (Ensemble)	0.8057	0.6024	0.8637

