

# Cancer T-Cell Epitope Classification

Identifying key target antigens for cancer immunotherapy

Tariq Alagha

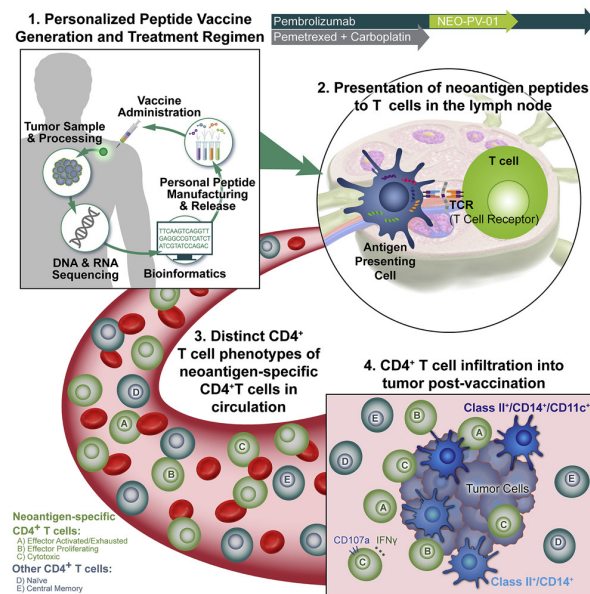


Figure 1: Awad et al. (2022)

## Introduction

At the heart of immune defense are tiny molecular “flags” called epitopes. These short sequences of amino acids, like “ADVEFCLSL”, sit on larger proteins and tell immune cells whether something is a threat. When an immune cell recognizes an epitope on a virus or a cancer cell, it can launch a protective attack.

Being able to reliably identify these epitopes in creating new vaccines and cancer treatments is crucial. Think of epitopes as the precise handshake between the immune system and a threat. Finding the right ones means smarter therapies can be designed. Unlike older treatments

like chemotherapy that harm healthy cells, therapies targeting specific epitopes can attack diseases with more accuracy. This promises better results for patients, with fewer debilitating side effects.

This project is all about teaching computers to do this vital identification work. Machine learning models are being built and tested that can look at an amino acid sequence and its properties and decide if it's an epitope or not. By making this process faster and more accurate, the discovery of new vaccine and immunotherapy candidates can be sped up, ultimately leading to more effective and kinder treatments.

## Data

### Dataset

To train these models, the main source is the [Immune Epitope Database](#)(IEDB), the largest public library of knowledge about how the immune system sees and reacts to epitopes on different molecules. It tells which sequences are known to be recognized by T-cells or antibodies.

For this project, the focus has been put on epitopes found on human cancer cells that have been proven in experiments to activate T-cell immune defenses. For each epitope, its unique amino acid sequence and the specific MHC allele it interacts with will be used.

Importantly, to train a good model, it doesn't just need to learn what an epitope is; it also needs to learn what it isn't. Many epitope records in the IEDB link to the full protein they come from. These full proteins are used to carefully select sequences that are not currently known epitopes. This provides a set of epitope and non-epitope examples.

## Libraries and packages

```
# Importing libraries
import pandas as pd
from collections import Counter
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import Bio
from Bio import SeqIO
from Bio.SeqUtils.ProtParam import ProteinAnalysis
from io import StringIO
import logomaker
import requests
from scipy import stats
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
    accuracy_score, auc, average_precision_score, classification_report,
    confusion_matrix, precision_recall_curve, roc_auc_score, roc_curve
)
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.utils.class_weight import compute_class_weight
import tensorflow as tf
from tensorflow.keras import metrics
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint, ReduceLROnPlateau
from tensorflow.keras.layers import (
    Conv1D, MaxPooling1D, Flatten, Dense, Dropout, Input, BatchNormalization
)
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.regularizers import l2
from sklearn.metrics import (
    classification_report, confusion_matrix, roc_curve, auc,
    precision_recall_curve, average_precision_score, accuracy_score
)
import matplotlib.pyplot as plt
```

## Preprocessing

Retrieving the data from IEDB was as simple as doing a search and clicking export. Using the requests python library, the full antigen sequence was downloaded and appended to the epitope dataset. Next, simple formatting was done to standardize the column names. Finally, the epitope dataset was merged with the assays dataset and filtered to include the following columns:

```
epitopes = pd.read_csv(r'/Users/tariq/Documents/capstone/data/epitope_table_export_174027958')
assays = pd.read_csv(r'/Users/tariq/Documents/capstone/data/tcell_table_export_1740279970.csv')

def fetch_full_sequence(url):
    if pd.isna(url): # Check if the URL is not NaN
        url = f'{url}.fasta'
    try:
        response = requests.get(url)
        if response.status_code == 200:
            fasta_io = StringIO(response.text)
            records = list(SeqIO.parse(fasta_io, "fasta"))
            if records: # Check if there are any records
                return str(records[0].seq)
            else:
                print("No records found in the FASTA file.")
    except requests.exceptions.RequestException as e:
        print(f"Request failed: {e}")
    return None

#epitopes['Full Sequence'] = epitopes['Epitope - Molecule Parent IRI'].apply(fetch_full_sequence)
epitopes = pd.read_csv(r'/Users/tariq/Documents/capstone/data/epitope_full_seq.csv')

# make all column names snake case
epitopes.columns = epitopes.columns.str.lower()
assays.columns = assays.columns.str.lower()

# remove spaces from column names
epitopes.columns = epitopes.columns.str.replace(' ', '')
epitopes.columns = epitopes.columns.str.replace('-', '_')
epitopes.columns = epitopes.columns.str.replace(' ', '_')

assays.columns = assays.columns.str.replace(' ', '')
assays.columns = assays.columns.str.replace('-', '_')
assays.columns = assays.columns.str.replace(' ', '_')
```

```

epitopes = epitopes.filter(['epitope_name', 'fullsequence'])
assays = assays.filter(['epitope_name', 'epitope_moluculeparent', 'host_name', 'host_mhcpres

# map mhc name and class from the assays dataframe to a new column in the epitopes dataframe
mhc = assays.filter(['epitope_name', 'mhcrestriction_name', 'mhcrestriction_class'])
mhc = mhc.drop_duplicates(subset=['epitope_name'])
epitopes = epitopes.merge(mhc, on='epitope_name', how='left')

epitopes.head()

```

	epitope_name	fullsequence	mhcrestric
0	AAGIGILTV	MPREDAHFIYGYPKKGHGHSYTTAEAAAGIGILTVILGVLLLIGCW...	HLA-A2
1	AAGIGILTVI	MPREDAHFIYGYPKKGHGHSYTTAEAAAGIGILTVILGVLLLIGCW...	HLA-A*0
2	ACDPHSGHFV	NaN	HLA-A2
3	ADLVGFLLK	MSLEQRSLHCKPEEALAEAAQAEALGLVCVQAATSSSSPLVLGTLEEV...	HLA-A*1
4	ADVEFCLSL	MLLAVLYCLLWSFQTSAGHFPRACVSSKNLMEKECCPPWSGDRSPC...	HLA-B*4

## MHC Allele Distribution

It's important to understand the distribution of MHC alleles associated with the epitopes in our dataset, as MHC molecules are responsible for presenting these peptides to T-cells. A skewed distribution could influence later analysis and create bias towards more represented alleles.

```

# Assuming 'epitopes' DataFrame with 'mhcrestriction_name' column is available
# from the previous preprocessing cell.

# Count MHC allele frequencies, dropping any NaNs first
mhc_counts = epitopes['mhcrestriction_name'].dropna().value_counts()

# Select top N for visualization
N = 20 # Show the top 20 alleles
top_mhc_counts = mhc_counts.head(N)

# Plotting
plt.figure(figsize=(12, 8))
sns.barplot(x=top_mhc_counts.values, y=top_mhc_counts.index, palette='viridis')
plt.title(f'Top {N} Most Frequent MHC Alleles in Epitope Dataset')
plt.xlabel('Frequency (Number of Epitopes)')
plt.ylabel('MHC Allele')

```

```
plt.tight_layout()
plt.show()

# Optionally, print some stats for context
# print(f"Total unique alleles found: {len(mhc_counts)}")
# print("Top 5 allele counts:\n", top_mhc_counts.head())
```




Figure 2: Top 20 Most Frequent MHC Alleles in Epitope Dataset

The plot reveals a skew in the MHC allele distribution within the initial epitope dataset. The allele **HLA-A\*02:01** is much more common compared to all others. This skew towards HLA-A\*02:01, is critical to acknowledge. It implies that later analyses and models might be heavily influenced by, or perform best on, peptides presented by this specific allele.

## Negative Sample Generation

```

def generate_negatives(row):
    epitope = row["epitope_name"]
    full_seq = row["fullsequence"]
    mhc = row["mhcrestriction_name"]

    # Handle missing or empty sequences
    if pd.isnull(full_seq) or full_seq == "":
        return []

    epitope = str(epitope)
    full_seq = str(full_seq)
    ep_len = len(epitope)

    negatives = []
    for i in range(len(full_seq) - ep_len + 1):
        window = full_seq[i:i+ep_len]
        if window != epitope:
            negatives.append({"peptide": window, "mhc": mhc})
    return negatives

# Apply the function to each row
negatives = pd.DataFrame()
negatives['negatives'] = epitopes.apply(generate_negatives, axis=1)
negatives = negatives[["negatives"]].explode("negatives").reset_index(drop=True)
negatives.dropna(subset=["negatives"], inplace=True)

# Remove duplicates
print(f"Shape before removing duplicates: {negatives.shape}")
negatives = negatives.drop_duplicates(subset=['negatives'])
print(f"Shape after removing duplicates: {negatives.shape}")

# Check for any remaining NaN values
print(f"Number of NaN values in negatives: {negatives['negatives'].isna().sum()}")

# Extract peptide and mhc into separate columns
negatives['peptide'] = negatives['negatives'].apply(lambda x: x['peptide'])
negatives['mhc'] = negatives['negatives'].apply(lambda x: x['mhc'])

```

Although the IEDB database provided a substantial amount of epitopes, in order to draw visual comparisons and create models to classify epitopes, samples of non-epitope peptides are

needed. These can be generated by shuffling and sampling amino acid sequences from the full antigen sequences of the epitopes, ensuring that the sampled sequences did not overlap with the epitope sequences.

There are pros and cons to this methodology. As opposed to generating completely random sequences of amino acids — sampling from larger sequences allows for natural patterns and physiochemical motifs to be retained. That is not to say the performance of statistical models or qualitative analysis will be better. Random sequences are more likely to be highly irregular, or even biologically implausible. Sampling from the full antigen sequences eliminates this potential bias.

Conversely, it is possible for a randomly sampled peptide to be an epitope that has not been tested yet, or simply isn't in the subset of data used for this analysis — resulting in an increase in the number of false negatives in our data.

## Feature Engineering

```
# Kyte-Doolittle hydrophobicity scale
kyte_doolittle = {
    'I': 4.5, 'V': 4.2, 'L': 3.8, 'F': 2.8, 'C': 2.5,
    'M': 1.9, 'A': 1.8, 'G': -0.4, 'T': -0.7, 'S': -0.8,
    'W': -0.9, 'Y': -1.3, 'P': -1.6, 'H': -3.2, 'E': -3.5,
    'Q': -3.5, 'D': -3.5, 'N': -3.5, 'K': -3.9, 'R': -4.5
}

def compute_avg_hydrophobicity(peptide):
    # Get hydrophobicity scores for each amino acid; default to 0 if missing
    scores = [kyte_doolittle.get(aa, 0) for aa in peptide]
    return sum(scores) / len(scores) if scores else 0

# Apply the function to the 'peptide' column to create a new column 'avg_hydro'
epitopes['epitope_avg_hydro'] = epitopes['epitope_name'].apply(compute_avg_hydrophobicity)
# Import the molecular_weight function from Bio.SeqUtils

def calculate_molecular_weight(peptide):
    """Calculate the molecular weight of a peptide sequence using Biopython."""
    try:
        # ProteinAnalysis only works with standard amino acids
        protein = ProteinAnalysis(peptide)
        return protein.molecular_weight()
    except Exception as e:
```



```

        # Handle peptides with non-standard amino acids
        return None

# Apply the function to calculate molecular weight for each epitope
epitopes['molecular_weight'] = epitopes['epitope_name'].apply(calculate_molecular_weight)
def calculate_aromaticity(peptide):
    """Calculate the aromaticity of a peptide sequence using Biopython."""
    try:
        # ProteinAnalysis only works with standard amino acids
        protein = ProteinAnalysis(peptide)
        return protein.aromaticity()
    except Exception as e:
        # Handle peptides with non-standard amino acids
        return None

# Apply the function to calculate molecular weight for each epitope
epitopes['aromaticity'] = epitopes['epitope_name'].apply(calculate_aromaticity)
def calculate_isoelectric_point(peptide):
    """Calculate the isoelectric point of a peptide sequence using Biopython."""
    try:
        # ProteinAnalysis only works with standard amino acids
        protein = ProteinAnalysis(peptide)
        return protein.isoelectric_point()
    except Exception as e:
        # Handle peptides with non-standard amino acids
        return None

# Apply the function to calculate molecular weight for each epitope
epitopes['isoelectric_point'] = epitopes['epitope_name'].apply(calculate_isoelectric_point)
def calculate_instability(peptide):
    """Calculate the instability of a peptide sequence using Biopython."""
    try:
        # ProteinAnalysis only works with standard amino acids
        protein = ProteinAnalysis(peptide)
        return protein.instability_index()
    except Exception as e:
        # Handle peptides with non-standard amino acids
        return None

# Apply the function to calculate molecular weight for each epitope
epitopes['instability'] = epitopes['epitope_name'].apply(calculate_instability)
def calculate_charge_at_pH7(peptide):

```

```

"""Calculate the charge of a peptide sequence at pH 7 using Biopython."""
try:
    # ProteinAnalysis only works with standard amino acids
    protein = ProteinAnalysis(peptide)
    return protein.charge_at_pH(7)
except Exception as e:
    # Handle peptides with non-standard amino acids
    return None

# Apply the function to calculate molecular weight for each epitope
epitopes['charge_at_pH7'] = epitopes['epitope_name'].apply(calculate_charge_at_pH7)

# Calculate features on the peptide column
negatives['peptide_length'] = negatives['peptide'].apply(len)
negatives['peptide_avg_hydro'] = negatives['peptide'].apply(compute_avg_hydrophobicity)
negatives['molecular_weight'] = negatives['peptide'].apply(calculate_molecular_weight)
negatives['aromaticity'] = negatives['peptide'].apply(calculate_aromaticity)
negatives['isoelectric_point'] = negatives['peptide'].apply(calculate_isoelectric_point)
negatives['instability'] = negatives['peptide'].apply(calculate_instability)
negatives['charge_at_pH7'] = negatives['peptide'].apply(calculate_charge_at_pH7)

negatives.drop('negatives', axis=1, inplace=True)

```

The protein analysis tool from the BioPython package allows for some quick feature engineering on most given peptides. For this analysis, the relevant features would be hydrophobicity, molecular weight, aromaticity, isoelectric point, instability, and the charge at pH7. Publications on epitope classification hold binding affinity — the ability for a peptide to bind to the body's MHC complex — to be a strong predictor. The BioPython package does not come with any functionality for binding affinity prediction but IEDB provides a tool called **netMHCpan**, which is the leading binding affinity prediction algorithm.

The IEDB website offers a GUI for using netMHCpan to predict binding affinities. However, it is only possible to run predictions on 100 peptides at a time and this analysis is examining many more than that. NetMHCpan can be downloaded and installed as a command line tool allowing more flexibility using python. Given an amino acid sequence and an MHC allele specification, netMHCpan returns a binding affinity score. This score ranges from 0 to 1, where higher values indicate a stronger likelihood of binding.

Subsequent analysis filters for 9-mer peptides, a common length for MHC Class I epitopes, for which binding prediction tools like netMHCpan are well-suited. After feature engineering, we have:

```
# Generated using the standalone netMHCpan tool with relevant MHC alleles for each peptide
epitopes = pd.read_csv("/Users/tariq/Documents/capstone/data/ninemer_epitopes.csv")
epitopes = epitopes.drop(columns=['fullsequence', 'mhcrestriction_name', 'mhcrestriction_classification'])
epitopes = epitopes.rename(columns={'epitope_name': 'peptide', 'epitope_avg_hydro': 'peptide_avg_hydro'})
epitopes_BA_pred = pd.read_csv("/Users/tariq/Documents/capstone/data/ninemer_epitopes_BA_pred.csv")

negatives = pd.read_csv("/Users/tariq/Documents/capstone/data/ninemer_negatives_trimmed.csv")
negatives = negatives.drop(columns=['mhc', 'peptide_length'])
negatives = negatives.rename(columns={'peptide': 'peptide'})
negatives = negatives.drop_duplicates(subset=['peptide'])
negatives_BA_pred = pd.read_csv("/Users/tariq/Documents/capstone/data/ninemer_negatives_BA_pred.csv")
negatives_BA_pred = negatives_BA_pred.drop_duplicates(subset=['peptide'])

# Merge the 'Score_BA' column from epitopes_BA_pred into the epitopes dataframe
epitopes = pd.merge(epitopes, epitopes_BA_pred[['peptide', 'Score_BA']], on='peptide', how='left')

negatives = pd.merge(negatives, negatives_BA_pred[['peptide', 'Score_BA']], on='peptide', how='left')

epitopes.head()
```

	peptide	peptide_avg_hydro	molecular_weight	aromaticity	isoelectric_point	instability	charge_at_pH7
0	AAGIGILTV	2.122222	813.9814	0.000000	5.570017	11.422222	-1.000000
1	ADVEFCLSL	1.233333	996.1348	0.111111	4.050028	20.855556	-1.000000
2	AFLPWHLRF	0.533333	1186.4061	0.333333	9.800371	53.400000	0.000000
3	ALAETSYVK	0.155556	981.1004	0.111111	6.045191	5.688889	-1.000000
4	ALDVYNGLL	0.966667	977.1115	0.111111	4.050028	-16.188889	-1.000000

## What Distinguishes an epitope from any other peptide?

### Statistical Comparison

```
# Compare numeric features between epitopes and negatives datasets
numeric_features = ['peptide_avg_hydro', 'molecular_weight', 'aromaticity',
                    'isoelectric_point', 'instability', 'charge_at_pH7', 'Score_BA']

# Create a figure with subplots for each numeric feature
fig, axes = plt.subplots(len(numeric_features), 1, figsize=(12, 5*len(numeric_features)))
#fig.tight_layout(pad=5.0)
```

```

# Plot density plots for each feature
for i, feature in enumerate(numeric_features):
    ax = axes[i]

    # Create density plot for Epitopes
    sns.kdeplot(epitopes[feature].dropna(), ax=ax, label='Epitopes', fill=True)

    # Create density plot for Negatives
    sns.kdeplot(negatives[feature].dropna(), ax=ax, label='Negatives', fill=True)

    # Add feature statistics
    epitope_mean = epitopes[feature].mean()
    negative_mean = negatives[feature].mean()

    ax.set_title(f'{feature} Density Plot Comparison')
    ax.text(0.02, 0.95, f'Epitopes mean: {epitope_mean:.4f}', transform=ax.transAxes)
    ax.text(0.02, 0.90, f'Negatives mean: {negative_mean:.4f}', transform=ax.transAxes)
    ax.legend()

    # Add p-value from t-test
    t_stat, p_value = stats.ttest_ind(
        epitopes[feature].dropna(),
        negatives[feature].dropna(),
        equal_var=False # Welch's t-test (doesn't assume equal variances)
    )
    #ax.text(0.02, 0.85, f'p-value: {p_value:.4e}', transform=ax.transAxes)

plt.suptitle('Comparison of Numeric Features Between Epitopes and Negatives', fontsize=16)
plt.show()

```

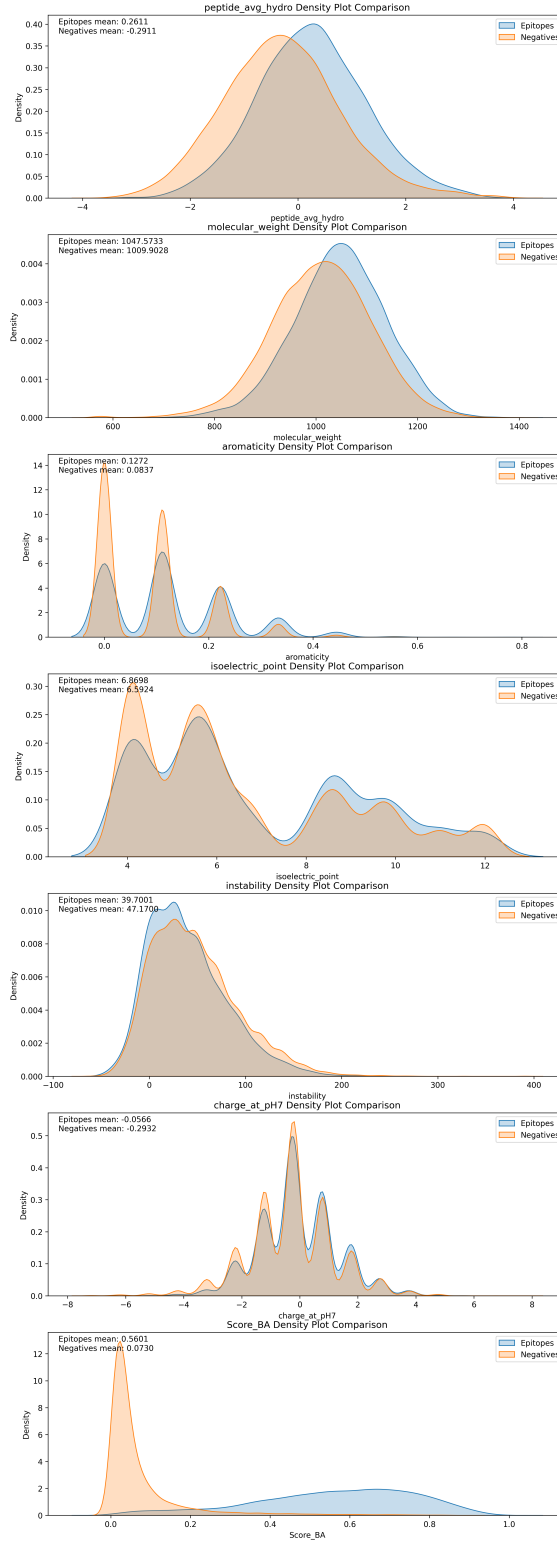


Figure 3: Density Plot Comparison of Numeric Features between Epitopes and Negatives

A boxplot comparison of the numerical variables reveals hardly significant differences between the epitope and non-epitope peptides. The clear outlier being the predicted binding affinity score.

```
# plot Score_BA for epitopes and negatives overlaid on the same plot
plt.figure(figsize=(10, 6))

# Use density instead of raw counts to normalize the histograms
plt.hist(epitopes['Score_BA'], bins=20, alpha=0.5, color='blue', edgecolor='black',
        label='Epitopes', density=True)
plt.hist(negatives['Score_BA'], bins=20, alpha=0.5, color='red', edgecolor='black',
        label='Negatives', density=True)

# Alternative approach: use log scale for y-axis
plt.yscale('log')

plt.xlabel('Binding Affinity')
plt.ylabel('Density (log scale)')
plt.title('Normalized Histogram of Binding Affinity for Epitopes vs Negatives')
plt.legend(prop={'size': 14}) # Increased legend font size
plt.tight_layout()
plt.show()
```

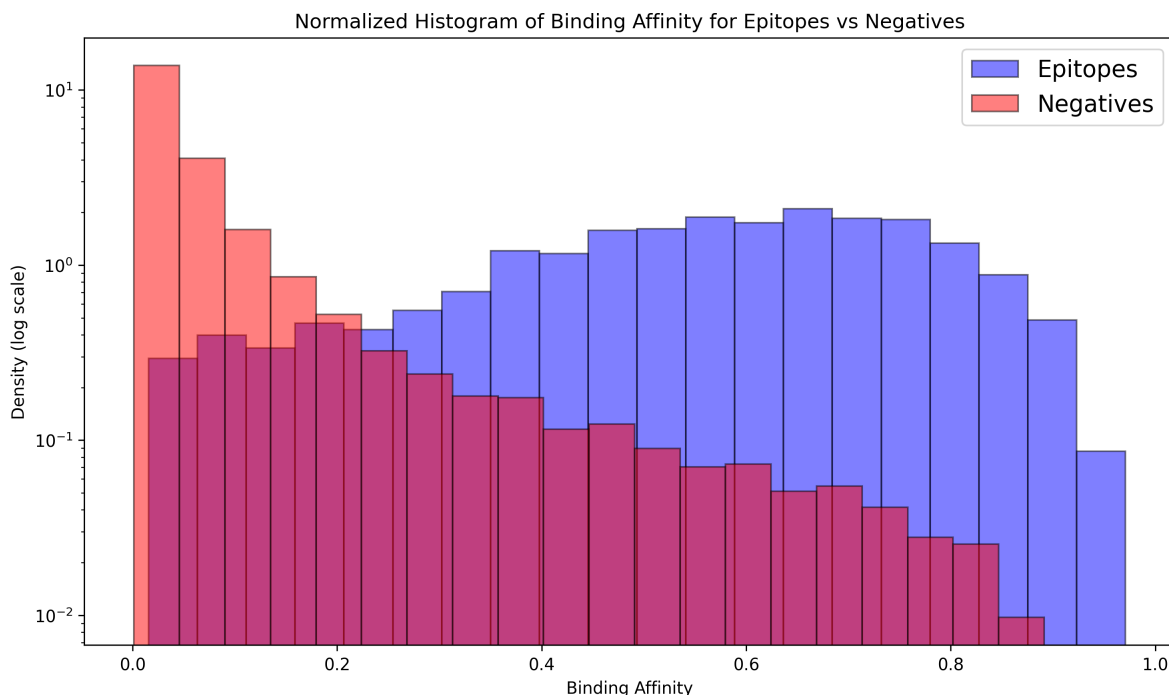


Figure 4: Normalized Histogram of Binding Affinity Scores

Upon further inspection of the difference in predicted binding affinity score, we see the non-epitope peptides exhibit a right-skewed distribution with a mean of 0.07, and the epitopes show a broad, moderate-variance spread with a much higher mean of 0.56.

## Sequence Motifs

While the predicted binding affinity score is a strong predictor of epitope classification, it is not the only feature that distinguishes an epitope from a non-epitope. To further understand the differences between the two classes, we can look for patterns in the amino acid sequences.

One way to explore these patterns is to examine the frequency of short amino acid motifs, such as tripeptides. By comparing the most frequent tripeptides in known epitopes versus non-epitope sequences, we might identify motifs that are enriched in one class or the other.

```
# Create temporary copies of epitopes and negatives, then add 'label' column for this analysis.
# Assumes 'epitopes' and 'negatives' DataFrames (without 'label' yet) are available from prior steps.
temp_epitopes_for_motifs = epitopes.copy()
temp_negatives_for_motifs = negatives.copy()
```

```

temp_epitopes_for_motifs['label'] = 1
temp_negatives_for_motifs['label'] = 0

# Now proceed with the lines you uncommented, using these temporary DataFrames
epitopes_filtered = temp_epitopes_for_motifs[['peptide', 'label']].copy()
epitopes_filtered.rename(columns={'peptide': 'sequence'}, inplace=True)
negatives_filtered = temp_negatives_for_motifs[['peptide', 'label']].copy()
negatives_filtered.rename(columns={'peptide': 'sequence'}, inplace=True)

# Use a specific name for this combined_data to avoid conflict
combined_data_for_tripeptides = pd.concat([epitopes_filtered, negatives_filtered], ignore_index=True)

def get_tripeptides(sequence):
    """Extracts all overlapping tripeptides from a sequence."""
    return [sequence[i:i+3] for i in range(len(sequence) - 2)]

# Separate epitope and non-epitope sequences
epitope_sequences = combined_data_for_tripeptides[combined_data_for_tripeptides['label'] == 1]
non_epitope_sequences = combined_data_for_tripeptides[combined_data_for_tripeptides['label'] == 0]

# Get all tripeptides for epitopes
all_epitope_tripeptides = []
for seq in epitope_sequences:
    all_epitope_tripeptides.extend(get_tripeptides(seq))

# Get all tripeptides for non-epitopes
all_non_epitope_tripeptides = []
for seq in non_epitope_sequences:
    all_non_epitope_tripeptides.extend(get_tripeptides(seq))

# Count frequencies
epitope_tripeptide_counts = Counter(all_epitope_tripeptides)
non_epitope_tripeptide_counts = Counter(all_non_epitope_tripeptides)

# Get top N most common tripeptides
N = 15
top_epitope_tripeptides = epitope_tripeptide_counts.most_common(N)
top_non_epitope_tripeptides = non_epitope_tripeptide_counts.most_common(N)

# Plotting
fig, axes = plt.subplots(2, 1, figsize=(12, 10))

```



```

# Epitope tripeptides
if top_epitope_tripeptides:
    peptides, counts = zip(*top_epitope_tripeptides)
    axes[0].bar(peptides, counts, color='skyblue')
    axes[0].set_title(f'Top {N} Most Frequent Tripeptides in Epitopes')
    axes[0].set_ylabel('Frequency')
    axes[0].tick_params(axis='x', rotation=45)
else:
    axes[0].text(0.5, 0.5, 'No tripeptides found for epitopes', horizontalalignment='center')

# Non-epitope tripeptides
if top_non_epitope_tripeptides:
    peptides, counts = zip(*top_non_epitope_tripeptides)
    axes[1].bar(peptides, counts, color='lightcoral')
    axes[1].set_title(f'Top {N} Most Frequent Tripeptides in Non-Epitopes')
    axes[1].set_ylabel('Frequency')
    axes[1].set_xlabel('Tripeptide')
    axes[1].tick_params(axis='x', rotation=45)
else:
    axes[1].text(0.5, 0.5, 'No tripeptides found for non-epitopes', horizontalalignment='center')

plt.tight_layout()
plt.show()

```




Figure 5: Top 15 Most Frequent Tripeptides in Epitope vs. Non-Epitope Sequences

These plots show the most common tripeptide sequences found within the epitope and non-epitope datasets. Comparing these can help identify if certain short amino acid patterns are more prevalent in one group over the other, potentially hinting at structural or functional differences recognized by the immune system or affecting MHC binding.

Interpreting the tripeptide frequency plots:

- **Epitopes:** The tripeptide LLL is dominant, with Leucine-rich sequences being prevalent.
- **Non-Epitopes:** SSS is the most frequent. While LLL is also common, this group shows more diversity with prominent Serine, Proline (PPP), and Glycine-based motifs (e.g., GGS).

The most frequent tripeptides differ significantly between epitopes and non-epitopes. Epitopes favor Leucine-based motifs, while non-epitopes have a broader range with SSS leading. This suggests that these short sequence patterns contribute to distinguishing the two.

Another powerful way to visualize conserved patterns in a set of sequences is by generating sequence logos. Each position in the logo consists of a stack of letters, where the height of each letter indicates its frequency at that position.

```
# Assuming 'combined_data_for_tripeptides' DataFrame with 'sequence' and 'label' columns is a
# from the previous tripeptide analysis cell.

# Standard 20 amino acids
amino_acids = sorted(list("ACDEFGHIKLMNPQRSTVWY"))
sequence_length = 9 # Assuming all sequences are 9-mers

epitope_sequences_for_logo = combined_data_for_tripeptides[combined_data_for_tripeptides['label'] == 'epitope']
non_epitope_sequences_for_logo = combined_data_for_tripeptides[combined_data_for_tripeptides['label'] == 'non-epitope']

def create_ppm_from_sequences(sequences, amino_acids_list, seq_len):
    if not sequences:
        return pd.DataFrame(0.0, index=amino_acids_list, columns=range(seq_len))

    pfm = pd.DataFrame(0, index=amino_acids_list, columns=range(seq_len))
    for seq in sequences:
        if len(seq) == seq_len: # Ensure sequence has expected length
            for i, char in enumerate(seq):
                if char in amino_acids_list: # Ensure character is a standard amino acid
                    pfm.loc[char, i] += 1

    # Convert PFM to PPM (Position Probability Matrix)
    ppm = pfm.div(len(sequences), axis='columns')
    return ppm

# Create PPM for epitopes
ppm_epitopes = create_ppm_from_sequences(epitope_sequences_for_logo, amino_acids, sequence_length)

# Create PPM for non-epitopes
ppm_non_epitopes = create_ppm_from_sequences(non_epitope_sequences_for_logo, amino_acids, sequence_length)

# Generate and display sequence logos
fig, axes = plt.subplots(2, 1, figsize=(12, 6))

# Epitope Logo
```

```

if not ppm_epitopes.empty and ppm_epitopes.sum().sum() > 0:
    # Transpose the PPM DataFrame for logomaker
    logomaker.Logo(ppm_epitopes.T, ax=axes[0], font_name='Arial Rounded MT Bold')
    axes[0].set_title('Sequence Logo for Epitopes')
    axes[0].set_ylabel('Bits') # Typically, height is in bits of information
else:
    axes[0].text(0.5, 0.5, 'No data or empty PPM for epitope logo', horizontalalignment='center')

# Non-Epitope Logo
if not ppm_non_epitopes.empty and ppm_non_epitopes.sum().sum() > 0:
    # Transpose the PPM DataFrame for logomaker
    logomaker.Logo(ppm_non_epitopes.T, ax=axes[1], font_name='Arial Rounded MT Bold')
    axes[1].set_title('Sequence Logo for Non-Epitopes')
    axes[1].set_xlabel('Position')
    axes[1].set_ylabel('Bits')
else:
    axes[1].text(0.5, 0.5, 'No data or empty PPM for non-epitope logo', horizontalalignment='center')

plt.tight_layout()
plt.show()

```

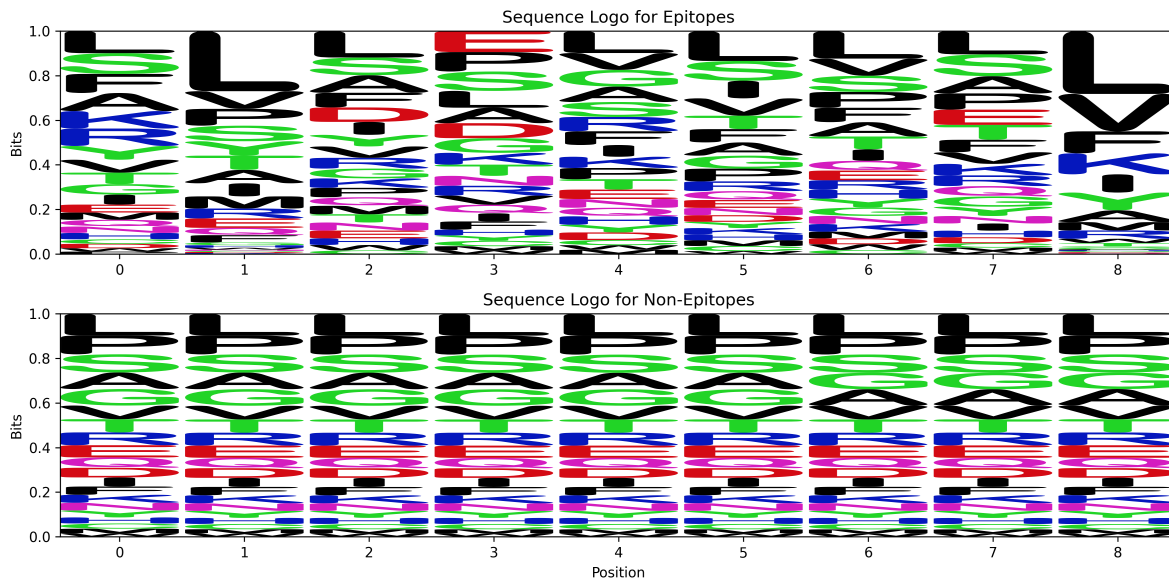


Figure 6: Sequence Logos for Epitope vs. Non-Epitope Sequences

These sequence logos visually represent the amino acid frequencies at each of the 9 positions

for epitopes and non-epitopes. Bigger letters indicate a higher frequency of that amino acid at that specific position.

Interpreting the sequence logos:

- **Epitope Logo:** While showing variability, there's a noticeable preference for Leucine at the position 8 and also a tendency for L at position 1. Other positions show a mix of amino acids with generally lower conservation, though residues like Serine, Lysine, and Valine appear at various spots.
- **Non-Epitope Logo:** This logo generally shows more diversity across all positions. While Leucine and Serine are common, no single amino acid dominates at most positions. The position 8 preference for L seen in epitopes is less pronounced here.

## Can we predict?

### Model Selection: A Baseline with State-of-the-Art Binding Prediction

To establish a baseline, we first develop a Random Forest classifier. This model incorporates a unique feature: predicted binding affinity scores (**Score\_BA**) derived from **netMHCpan**, a state-of-the-art algorithm for MHC binding prediction. By including this, our baseline leverages existing sophisticated domain knowledge. The full feature set includes:

- Average Hydrophobicity
- Molecular Weight
- Aromaticity
- Isoelectric Point
- Instability
- Charge at pH7

Performance will be evaluated based on accuracy, precision, and recall.

### Preprocessing

Prior to training, labels are assigned to the epitopes and non-epitopes as 1 or 0 respectively. The two samples are then concatenated, scaled, and shuffled. Finally, the data is split into training and testing sets with an 80/20 ratio.

```

# Add label column to epitopes dataframe (positive class = 1)
epitopes['label'] = 1

# Add label column to negatives dataframe (negative class = 0)
negatives['label'] = 0

# Combine the positive and negative examples
combined_data = pd.concat([epitopes, negatives], ignore_index=True)

# Shuffle the combined dataset
combined_data = combined_data.sample(frac=1, random_state=42).reset_index(drop=True)

# Define features and target
X = combined_data.drop(columns=['peptide', 'label'])
y = combined_data['label']

# Identify numerical columns to scale (exclude one-hot encoded amino acid columns)
numerical_cols = ['peptide_avg_hydro', 'molecular_weight', 'aromaticity', 'isoelectric_point']

# Split the data into training and testing sets (80% train, 20% test)

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Scale numerical features using StandardScaler
scaler = StandardScaler()
X_train[numerical_cols] = scaler.fit_transform(X_train[numerical_cols])
X_test[numerical_cols] = scaler.transform(X_test[numerical_cols])

# Print the shapes to verify the split
print(f"Training set: {X_train.shape[0]} samples")
print(f"Testing set: {X_test.shape[0]} samples")
print(f"Positive samples in training: {sum(y_train == 1)}")
print(f"Negative samples in training: {sum(y_train == 0)}")
print(f"Positive samples in testing: {sum(y_test == 1)}")
print(f"Negative samples in testing: {sum(y_test == 0)}")

Training set: 20502 samples
Testing set: 5126 samples
Positive samples in training: 4236
Negative samples in training: 16266
Positive samples in testing: 1059
Negative samples in testing: 4067

```

Figure 7

## Training + Evaluation

The random forest classifier is fit to the training data and evaluated on the testing data.

```
# Initialize the Random Forest Classifier
rf_model = RandomForestClassifier(
    n_estimators=100, # Number of trees
    max_depth=None,   # Maximum depth of trees
    min_samples_split=2,
    min_samples_leaf=1,
    random_state=42
)

# Train the model
rf_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = rf_model.predict(X_test)
y_pred_proba = rf_model.predict_proba(X_test)[:, 1] # Probability estimates for positive class

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)

# Plot Confusion Matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Predicted Negative', 'Predicted Positive'],
            yticklabels=['Actual Negative', 'Actual Positive'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix - Random Forest')
plt.show()

# Evaluate the model
print("\nRandom Forest Model Classification Report:")
print(classification_report(y_test, y_pred))
```




Figure 8: Confusion Matrix for Random Forest with Binding Affinity

Random Forest Model Classification Report:				
	precision	recall	f1-score	support
0	0.93	0.96	0.95	4067
1	0.82	0.74	0.78	1059
accuracy			0.91	5126
macro avg	0.88	0.85	0.86	5126
weighted avg	0.91	0.91	0.91	5126

The results show strong performance from the random forest classifier, with an overall accuracy



of 91% and recall of 75% for the positive class. This strong performance is due to the powerful Score\_BA predictor from netMHCpan.

```
# Feature importance
feature_importance = pd.DataFrame({
    'Feature': X_train.columns,
    'Importance': rf_model.feature_importances_
})
feature_importance = feature_importance.sort_values('Importance', ascending=True)

# Plot top 15 features
plt.figure(figsize=(10, 6))
top_features = feature_importance.head(15)
plt.barh(np.arange(len(top_features)), top_features['Importance'], align='center')
plt.yticks(np.arange(len(top_features)), top_features['Feature'])
plt.xlabel('Importance')
plt.title('Feature Importance - Random Forest')
plt.tight_layout()
plt.show()
```




Figure 9: Feature Importance for Random Forest with Binding Affinity

## Quantifying the Impact of Binding Affinity Prediction

To highlight the significant contribution of the `netMHCpan` prediction of binding affinity, we next evaluate the Random Forest model without the `Score_BA` feature. This helps quantify the performance drop when relying solely on other physiochemical properties without this advanced binding prediction.

```
# drop the Score_BA column
X_train = X_train.drop(columns=['Score_BA'])
X_test = X_test.drop(columns=['Score_BA'])

# Initialize the Random Forest Classifier
rf_model = RandomForestClassifier(
    n_estimators=100, # Number of trees
    max_depth=None, # Maximum depth of trees
    min_samples_split=2,
    min_samples_leaf=1,
    random_state=42
)

# Train the model
rf_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = rf_model.predict(X_test)
y_pred_proba = rf_model.predict_proba(X_test)[:, 1] # Probability estimates for positive class

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)

# Plot Confusion Matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Predicted Negative', 'Predicted Positive'],
            yticklabels=['Actual Negative', 'Actual Positive'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix - Random Forest')
plt.show()

# Evaluate the model
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```




Figure 10: Confusion Matrix for Random Forest without Binding Affinity

Classification Report:					
	precision	recall	f1-score	support	
0	0.82	0.95	0.88	4067	
1	0.47	0.18	0.26	1059	
accuracy			0.79	5126	
macro avg	0.64	0.57	0.57	5126	
weighted avg	0.75	0.79	0.75	5126	

The accuracy only drops from 91% to 79%. However, this is misleading when considering the class imbalance of the data. The ratio of negative samples to positive is roughly 4:1, respectively. So, predicting the majority class — non-epitope — almost everytime would result in the majority of the testing data being correctly predicted and labeled.

A better performance metric to compare between models would be the model's recall rate on the positive class. How many of the epitopes in the testing data were correctly predicted and labeled? The same model, when predicted binding affinity was included as a predictor, produced a 74% recall rate while the current model has a much lower 18% recall rate. This dramatic drop in recall for epitopes clearly demonstrates the model's dependence on the `netMHCpan` binding affinity scores for identifying true positives.

```
# Feature importance
feature_importance = pd.DataFrame({
    'Feature': X_train.columns,
    'Importance': rf_model.feature_importances_
})
feature_importance = feature_importance.sort_values('Importance', ascending=True)

# Plot top 15 features
plt.figure(figsize=(10, 6))
top_features = feature_importance.head(15)
plt.barh(np.arange(len(top_features)), top_features['Importance'], align='center')
plt.yticks(np.arange(len(top_features)), top_features['Feature'])
plt.xlabel('Importance')
plt.title('Feature Importance - Random Forest')
plt.tight_layout()
plt.show()
```




Figure 11: Feature Importance for Random Forest without Binding Affinity

### A Different Approach: Learning Directly from Sequence with CNNs

Given the Random Forest model's significant reliance on pre-computed binding affinity from **netMHCpan**, we explore an alternative strategy: a Convolutional Neural Network (CNN). The key value of the CNN in this context is its ability to learn predictive patterns directly from the raw amino acid sequences themselves. Unlike the RF model which uses engineered features and external predictions like **Score\_BA**, the CNN can potentially uncover complex sequence motifs, positional tendencies, and other relevant features that may not be fully captured by binding affinity predictions alone. This approach allows the model to discover new features without feature engineering or reliance on separate binding prediction tools.

To start, the data will be filtered to only include the amino acid sequence and respective label.

```
# Filter and combine
epitopes_filtered = epitopes[['peptide', 'label']].copy()
epitopes_filtered.rename(columns={'peptide': 'sequence'}, inplace=True)
negatives_filtered = negatives[['peptide', 'label']].copy()
negatives_filtered.rename(columns={'peptide': 'sequence'}, inplace=True)
```

```
combined_data = pd.concat([epitopes_filtered, negatives_filtered], ignore_index=True)

# Shuffle the validated data
combined_data = combined_data.sample(frac=1, random_state=42).reset_index(drop=True)
combined_data.head()
```

	sequence	label
0	IENIWSPEG	1
1	PYQVPFVQA	0
2	FPEGLDPSA	0
3	LLFTDQHGL	1
4	DDRESWPSV	0

To train a neural network, the sequence must be represented in a numerical format. Assigning each amino acid a unique integer value up to 20, the sequences are converted to a list of integers. Then, the resulting integer sequences are one-hot encoded into a 3-dimensional array of shape (25628, 9, 20). 25,628 peptides, 9 amino acids in each peptide, 20 unique amino acids. The data is split into training, validation, and testing sets with a 70/15/15 ratio.

```
# Extract sequences
all_sequences = combined_data['sequence'].tolist()

# Find unique characters (amino acids) across all sequences
unique_chars = sorted(list(set("".join(all_sequences))))

# Map characters to indices starting from 0 (no padding index needed)
char_to_index = {char: i for i, char in enumerate(unique_chars)}
index_to_char = {i: char for i, char in enumerate(unique_chars)}

num_chars = len(unique_chars) # Vocabulary size is just the number of unique chars

# Convert sequences to integer sequences
int_sequences = []
for seq in all_sequences:
    int_seq = [char_to_index[char] for char in seq]
    int_sequences.append(int_seq)

# One-hot encode the integer sequences
# Shape: (num_samples, sequence_length, num_unique_chars)
# Assuming all sequences have length 9 as implied by the shape (9, num_chars)
```

```

sequence_length = 9 # Explicitly define sequence length
X_onehot = np.zeros((len(int_sequences), sequence_length, num_chars), dtype=np.float32)

for i, seq in enumerate(int_sequences):
    # Ensure sequence length matches expected length before encoding
    if len(seq) == sequence_length:
        for j, char_idx in enumerate(seq): # j is position (0-8), char_idx is the integer in
            X_onehot[i, j, char_idx] = 1.0
    else:
        print(f"Warning: Sequence at index {i} has length {len(seq)}, expected {sequence_length}")

y = combined_data['label'].values

# --- Data Splitting (70/15/15) ---
print("--- Data Splitting ---")
# Split into temp (85%) and test (15%)
X_temp, X_test, y_temp, y_test = train_test_split(
    X_onehot, y, test_size=0.15, random_state=42, stratify=y
)
# Split temp into train (70% of total) and validation (15% of total)
val_split_ratio = 0.15 / 0.85 # Calculate split ratio for validation set
X_train, X_val, y_train, y_val = train_test_split(
    X_temp, y_temp, test_size=val_split_ratio, random_state=42, stratify=y_temp
)

print(f"Training set: {X_train.shape}")
print(f"Validation set: {X_val.shape}")
print(f"Testing set: {X_test.shape}")
print("-" * 30)

```

```

--- Data Splitting ---
Training set: (17938, 9, 20)
Validation set: (3845, 9, 20)
Testing set: (3845, 9, 20)
-----

```

This CNN model processes one-hot encoded sequences using three 1D convolutional blocks (64, 128, 256 filters with BatchNormalization, MaxPooling, L2 regularization). Two subsequent dense blocks (256, 128 units with BatchNormalization, Dropout, L2 regularization, ReLU) follow. A final softmax layer outputs class probabilities. The model is compiled with Adam, sparse\_categorical\_crossentropy loss, and accuracy.

```

def create_cnn_model(input_shape, num_classes=2):
    """Creates and compiles the CNN model."""
    inputs = Input(shape=input_shape)
    x = Conv1D(64, kernel_size=8, activation='relu', padding='same', kernel_regularizer=l2(0.001))(x)
    x = BatchNormalization()(x)
    x = MaxPooling1D(pool_size=2, padding='same')(x)

    x = Conv1D(128, kernel_size=8, activation='relu', padding='same', kernel_regularizer=l2(0.001))(x)
    x = BatchNormalization()(x)
    x = MaxPooling1D(pool_size=2, padding='same')(x)

    x = Conv1D(256, kernel_size=8, activation='relu', padding='same', kernel_regularizer=l2(0.001))(x)
    x = BatchNormalization()(x)

    x = Flatten()(x)

    x = Dense(256, activation='relu', kernel_regularizer=l2(0.001))(x)
    x = BatchNormalization()(x)
    x = Dropout(0.5)(x)
    x = Dense(128, activation='relu', kernel_regularizer=l2(0.001))(x)
    x = BatchNormalization()(x)
    x = Dropout(0.4)(x)

    outputs = Dense(num_classes, activation='softmax')(x)

    model = Model(inputs=inputs, outputs=outputs)

    optimizer = Adam(learning_rate=0.001)
    model.compile(optimizer=optimizer,
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

    return model

input_shape = (9, num_chars)
model = create_cnn_model(input_shape)

```

Before training the model, class weights are calculated to handle the data imbalance, effectively telling the model to “pay more attention” to samples from the minority class.

```

from sklearn.utils.class_weight import compute_class_weight

# Calculate class weights

```



```

class_weights = compute_class_weight('balanced', classes=np.unique(y_train), y=y_train)
class_weight_dict = dict(enumerate(class_weights))

# Define callbacks
early_stopping = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True, verbose=0)
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, patience=5, min_lr=1e-6, verbose=0)
model_checkpoint = ModelCheckpoint('best_cnn_model_len9.keras', monitor='val_accuracy', save_best_only=True, verbose=0)

# Train the model
history = model.fit(
    X_train, y_train,
    epochs=50,
    batch_size=32,
    validation_data=(X_val, y_val),
    callbacks=[early_stopping, reduce_lr, model_checkpoint],
    class_weight=class_weight_dict,
    verbose=0
)

model = tf.keras.models.load_model('best_cnn_model_len9.keras')

# Evaluate on test data
test_loss, test_accuracy = model.evaluate(X_test, y_test, verbose=0)

# Get prediction probabilities for the positive class
y_pred_proba = model.predict(X_test, verbose=0)
y_pred_proba_positive = y_pred_proba[:, 1]

# Apply fixed threshold of 0.5
y_pred = (y_pred_proba_positive >= 0.5).astype(int)

# Print classification report
print("\nCNN Classification Report:")

# Calculate and plot confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 5))
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title(f'Confusion Matrix')
plt.colorbar()
tick_marks = np.arange(2)
plt.xticks(tick_marks, ['Negative', 'Positive'])

```

```

plt.yticks(tick_marks, ['Negative', 'Positive'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
thresh = cm.max() / 2.
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        plt.text(j, i, format(cm[i, j], 'd'),
                  horizontalalignment="center",
                  color="white" if cm[i, j] > thresh else "black")
plt.tight_layout()
print(classification_report(y_test, y_pred))
plt.show()

```

#### CNN Classification Report:

	precision	recall	f1-score	support
0	0.90	0.89	0.90	3051
1	0.60	0.61	0.60	794
accuracy			0.83	3845
macro avg	0.75	0.75	0.75	3845
weighted avg	0.84	0.83	0.84	3845




Figure 12: Confusion Matrix for CNN Model

The Convolutional Neural Network produced a reasonable overall accuracy of 82%, primarily driven by its strong ability to correctly identify the majority, non-epitopes, with high precision and recall. However, its performance on the minority class, epitopes, was mixed; while managing to recall 65% of true epitopes, its precision was significantly lower at 56%, indicating that nearly half of its positive predictions were false positives.

The Convolutional Neural Network significantly outperforms the Random Forest classifier while relying only on the sequence data, and no additional features. The RF struggles with identifying the actual epitopes, shown by the low recall for Class 1, whereas the CNN, while still challenged with precision for Class 1, provides a better balance and significantly higher recall for the positive class, making it the more effective model in this comparison.

## What did the CNN learn?

To better understand what the CNN model learned, we can generate saliency maps. A saliency map highlights which input features (amino acids at specific positions) were most influential in the model's decision for a given input sample. This is done by calculating the gradient of the model's output with respect to the input features.

The following plots examine a single sample sequence, AASCFTASV, which the model misclassified. The true label is Non-epitope, but the model predicted it as an Epitope with high confidence. The first map shows which features drove the incorrect prediction, and the second map shows which features would have been important for the correct classification.

```
# Ensure the model is loaded or available from previous cells
# If not, load it:
model = tf.keras.models.load_model('best_cnn_model_len9.keras')

# Ensure X_test, y_test, and index_to_char are available from the 'sequence_encoding' cell

# Select a sample from the test set
sample_index = 5 # You can change this to inspect different samples
sample_input = X_test[sample_index:sample_index+1] # Keep batch dimension
sample_label = y_test[sample_index]
true_class_index = int(sample_label) # Ensure it's an integer for indexing

# Convert the one-hot encoded sample back to a sequence for display
sample_sequence_onehot = X_test[sample_index]
sample_amino_acid_indices = np.argmax(sample_sequence_onehot, axis=1)
sample_sequence_str = "".join([index_to_char[idx] for idx in sample_amino_acid_indices])

print(f"Sample Sequence: {sample_sequence_str}")
print(f"True Label: {'Epitope' if sample_label == 1 else 'Non-epitope'} (Class {sample_label})")

# Make a prediction to confirm
sample_pred_proba = model.predict(sample_input, verbose=0)
predicted_class = np.argmax(sample_pred_proba[0])
print(f"Predicted Label: {'Epitope' if predicted_class == 1 else 'Non-epitope'} (Class {predicted_class})")

# Convert sample_input to tf.Tensor for gradient taping
sample_input_tf = tf.convert_to_tensor(sample_input, dtype=tf.float32)

with tf.GradientTape() as tape:
    tape.watch(sample_input_tf)
    predictions_tensor = model(sample_input_tf) # Get model output for current sample
```

```

predicted_class_idx = tf.argmax(predictions_tensor, axis=1).numpy()[0] # Determine predicted class

# Using predicted class for saliency as per latest changes
output_neuron_to_explain = predictions_tensor[:, predicted_class_idx]

# Get the gradients of the output neuron with respect to the input
saliency_grads = tape.gradient(output_neuron_to_explain, sample_input_tf)

# saliency will have shape (1, 9, 20)
# We take the absolute values and then sum across the one-hot encoding dimension
# or take the max across the one-hot encoding for each position
saliency_map_per_position = np.sum(np.abs(saliency_grads[0]), axis=1)
# saliency_map_per_position = np.max(np.abs(saliency_grads[0]), axis=1)

# Normalize the saliency map for visualization
saliency_map_normalized = (saliency_map_per_position - np.min(saliency_map_per_position)) /

# Plotting the saliency map
plt.figure(figsize=(10, 4))
plt.bar(range(len(sample_sequence_str)), saliency_map_normalized, color='skyblue')
plt.xticks(range(len(sample_sequence_str)), list(sample_sequence_str))
plt.xlabel("Amino Acid Position")
plt.ylabel("Normalized Saliency")
plt.title(f"Saliency Map for Sequence: {sample_sequence_str} (Predicted: {'Epitope' if predicted_class_idx == true_class_index else 'Not Epitope'})")
plt.tight_layout()
plt.show()

# You can also try visualizing the gradients for the actual true class if different
if predicted_class != true_class_index:
    with tf.GradientTape() as tape_true:
        tape_true.watch(sample_input_tf)
        predictions_true = model(sample_input_tf)
        output_neuron_true_class = predictions_true[:, true_class_index]
        saliency_true = tape_true.gradient(output_neuron_true_class, sample_input_tf)
        saliency_map_per_position_true = np.sum(np.abs(saliency_true[0]), axis=1)
        saliency_map_normalized_true = (saliency_map_per_position_true - np.min(saliency_map_per_position_true)) /

    plt.figure(figsize=(10, 4))
    plt.bar(range(len(sample_sequence_str)), saliency_map_normalized_true, color='salmon')

```

```
plt.xticks(range(len(sample_sequence_str)), list(sample_sequence_str))
plt.xlabel("Amino Acid Position")
plt.ylabel("Normalized Saliency (for True Class)")
plt.title(f"Saliency Map for Sequence: {sample_sequence_str} (Influence on True Class: {
plt.tight_layout()
plt.show()
```

Sample Sequence: DLSPDGPRS

True Label: Non-epitope (Class 0)

Predicted Label: Non-epitope (Class 0) with probability 0.9954

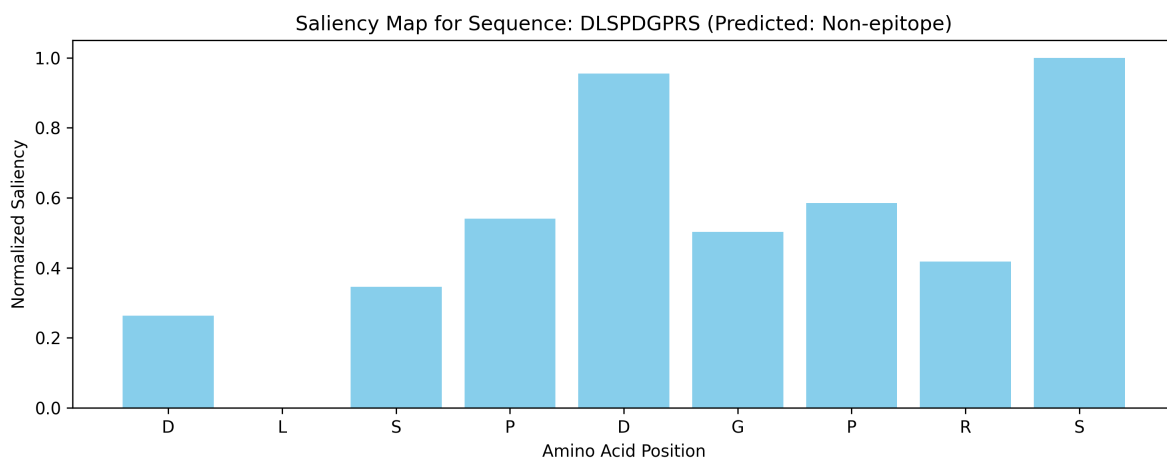


Figure 13: Example Saliency Maps for a Misclassified Sample (AASCFTASV)

For this specific misclassified sample (AASCFTASV):

- The model's incorrect Epitope prediction was strongly influenced by the C at position 3, A at position 6, and V at position 8.
- Interestingly, for the correct Non-epitope classification, the C at position 3 and V at position 8 also show high importance.

This suggests these positions are decision points for the model for this sequence with conflicting signals leading to the misclassification.

## What about the whole dataset?

To get a more general understanding of feature importance across the dataset, we can compute average saliency maps for different classes of samples. Below, we calculate and plot the average saliency map for true epitopes and true non-epitopes in the test set. The saliency is calculated with respect to the true class output for each sample. This helps to reveal general patterns the model has learned.

```
# Ensure model, X_test, y_test, sequence_length, and num_chars are available
# model = tf.keras.models.load_model('best_cnn_model_len9.keras') # If not already loaded

num_samples_to_process = len(X_test) # Or a smaller number for quicker testing, e.g., 100

# Accumulators for saliency maps and counts
saliency_accumulator_epitopes = np.zeros(sequence_length)
count_epitopes = 0
saliency_accumulator_non_epitopes = np.zeros(sequence_length)
count_non_epitopes = 0

for i in range(num_samples_to_process):
    sample_input = X_test[i:i+1] # Keep batch dimension
    true_label = y_test[i]
    true_class_idx = int(true_label)

    sample_input_tf = tf.convert_to_tensor(sample_input, dtype=tf.float32)

    with tf.GradientTape() as tape:
        tape.watch(sample_input_tf)
        predictions_tensor = model(sample_input_tf) # Get model output for current sample

        # Using true class for saliency
        output_neuron_to_explain = predictions_tensor[:, true_class_idx]

    saliency_grads = tape.gradient(output_neuron_to_explain, sample_input_tf)

    if saliency_grads is not None:
        # Using np.max as per your previous preference
        saliency_map_per_position = np.sum(np.abs(saliency_grads[0].numpy()), axis=1)
        # saliency_map_per_position = np.max(np.abs(saliency_grads[0].numpy()), axis=1)

        if true_label == 1: # Epitope
            saliency_accumulator_epitopes += saliency_map_per_position
            count_epitopes += 1
```

```

        else: # Non-epitope
            saliency_accumulator_non_epitopes += saliency_map_per_position
            count_non_epitopes += 1
    else:
        print(f"Warning: Gradients were None for sample {i}. Skipping.")

    if (i + 1) % 100 == 0:
        #print(f"Processed {i+1}/{num_samples_to_process} samples...")
        pass

#print("Aggregation complete.")

# Calculate average saliency maps
if count_epitopes > 0:
    avg_saliency_epitopes = saliency_accumulator_epitopes / count_epitopes
    avg_saliency_epitopes_normalized = (avg_saliency_epitopes - np.min(avg_saliency_epitopes))
else:
    avg_saliency_epitopes_normalized = None
    print("No epitope samples processed or found to calculate average saliency.")

if count_non_epitopes > 0:
    avg_saliency_non_epitopes = saliency_accumulator_non_epitopes / count_non_epitopes
    avg_saliency_non_epitopes_normalized = (avg_saliency_non_epitopes - np.min(avg_saliency_non_epitopes))
else:
    avg_saliency_non_epitopes_normalized = None
    print("No non-epitope samples processed or found to calculate average saliency.")

# Plotting the average saliency maps
fig, axs = plt.subplots(2, 1, figsize=(12, 8), sharex=True)

if avg_saliency_epitopes_normalized is not None:
    axs[0].bar(range(sequence_length), avg_saliency_epitopes_normalized, color='lightcoral')
    axs[0].set_ylabel("Normalized Saliency")
    axs[0].set_title("Average Saliency Map for True Epitopes")
    axs[0].set_xticks(range(sequence_length))
else:
    axs[0].text(0.5, 0.5, 'No data for epitopes', horizontalalignment='center', verticalalignment='center')

if avg_saliency_non_epitopes_normalized is not None:
    axs[1].bar(range(sequence_length), avg_saliency_non_epitopes_normalized, color='skyblue')
    axs[1].set_xlabel("Amino Acid Position")
    axs[1].set_ylabel("Normalized Saliency")

```



```

    axs[1].set_title("Average Saliency Map for True Non-Epitopes")
    axs[1].set_xticks(range(sequence_length))
else:
    axs[1].text(0.5, 0.5, 'No data for non-epitopes', horizontalalignment='center', verticalalignment='center')

plt.tight_layout()
plt.show()

```

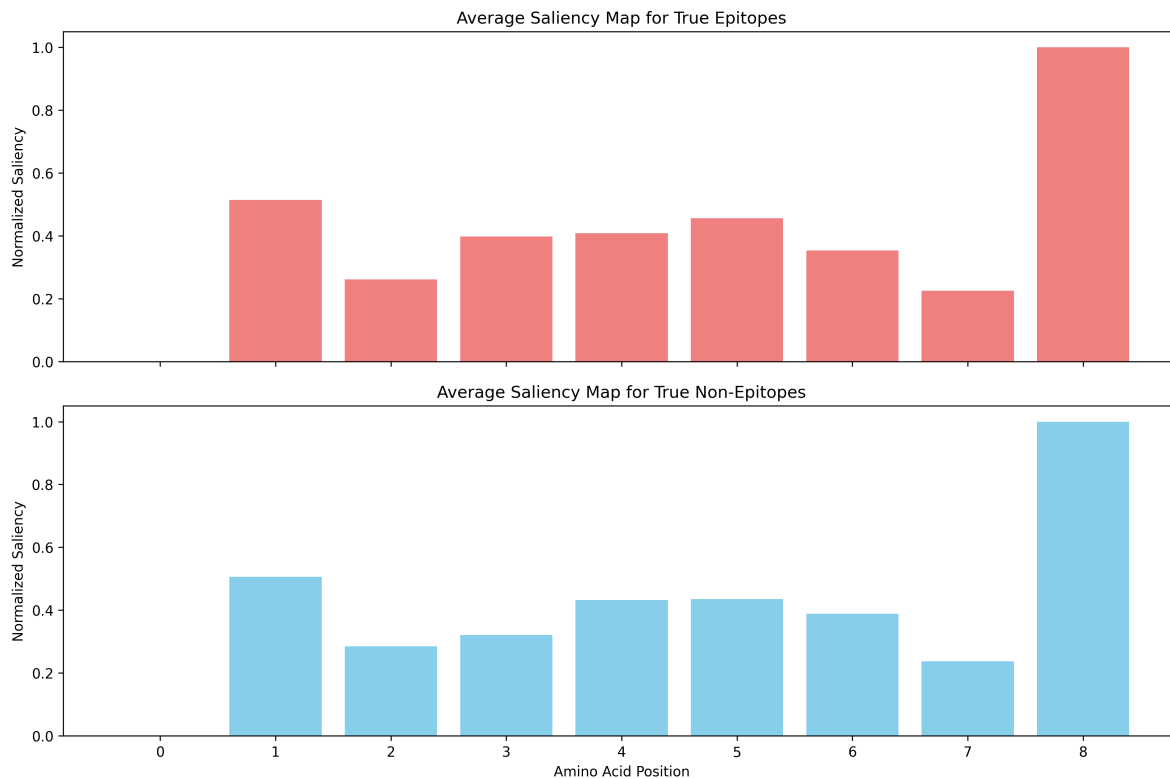


Figure 14: Average Saliency Maps for True Epitopes and Non-Epitopes

Interpreting the aggregated saliency maps:

- **Dominance of position 8:** For both true epitopes and true non-epitopes, the amino acid at position 8 shows by far the highest average saliency. This indicates the model heavily relies on the identity of this final residue to make its classification.
- **Importance of position 1 for non-epitopes:** For true non-epitopes, the amino acid at position 1 also shows notably high average saliency, suggesting its importance in identifying a sequence as not an epitope.

- **General pattern:** The model appears to have learned that the C-terminal residue is a primary determinant, with other positions like position 1 playing secondary, but still significant, roles.

## Conclusion

This project aimed to develop and evaluate computational models for classifying cancer T-cell epitopes, a necessary task for advancing personalized cancer treatment. By utilizing data from the Immune Epitope Database (IEDB) and employing feature engineering and machine learning techniques, we explored factors differentiating epitopes from non-epitope peptides derived from the same source proteins.

A Random Forest model incorporating predicted binding affinity scores from **netMHCpan** achieved high overall performance and reasonable recall for identifying true epitopes. However, removing this binding affinity feature caused a dramatic drop in the model’s ability to identify epitopes, indicating that standard physicochemical features alone, while descriptive, were insufficient for robust classification in this context.

Significantly, a Convolutional Neural Network (CNN) trained on one-hot encoded peptide sequences, without relying on external binding predictors or engineered features, demonstrated superior performance compared to the feature-based Random Forest without binding affinity. The CNN achieved balanced performance with better recall for the epitope class. This demonstrates the potential of deep learning models to capture patterns directly from sequence data.

## References

- Awad, Mark M., Ramaswamy Govindan, Kristen N. Balogh, David R. Spigel, Edward B. Garon, Meghan E. Bushway, Asaf Poran, et al. 2022. “Personalized Neoantigen Vaccine NEO-PV-01 with Chemotherapy and Anti-PD-1 as First-Line Treatment for Non-Squamous Non-Small Cell Lung Cancer.” *Cancer Cell* 40 (9): 1010–1026.e11. <https://doi.org/10.1016/j.ccell.2022.08.003>.
- “Immune Epitope Database (IEDB).” 2025. <https://www.iedb.org>.
- Reynisson, Birkir, Bruno Alvarez, Sinu Paul, Bjoern Peters, and Morten Nielsen. 2020. “NetMHCpan-4.1 and NetMHCIIpan-4.0: Improved Predictions of MHC Antigen Presentation by Concurrent Motif Deconvolution and Integration of MS MHC Eluted Ligand Data.” *Nucleic Acids Research* 48 (W1): W449–54. <https://doi.org/10.1093/nar/gkaa379>.
- The UniProt Consortium. 2025. “UniProt: The Universal Protein Knowledgebase.” <https://www.uniprot.org>.

Yarchoan, Mark, Burles A. Johnson, Eric R. Lutz, Daniel A. Laheru, and Elizabeth M. Jaffee. 2017. “Targeting Neoantigens to Augment Antitumour Immunity.” *Nature Reviews Cancer* 17 (4): 209–22. <https://doi.org/10.1038/nrc.2016.154>.